



**NOAA
FISHERIES**

PIFSC/ESD/ARP

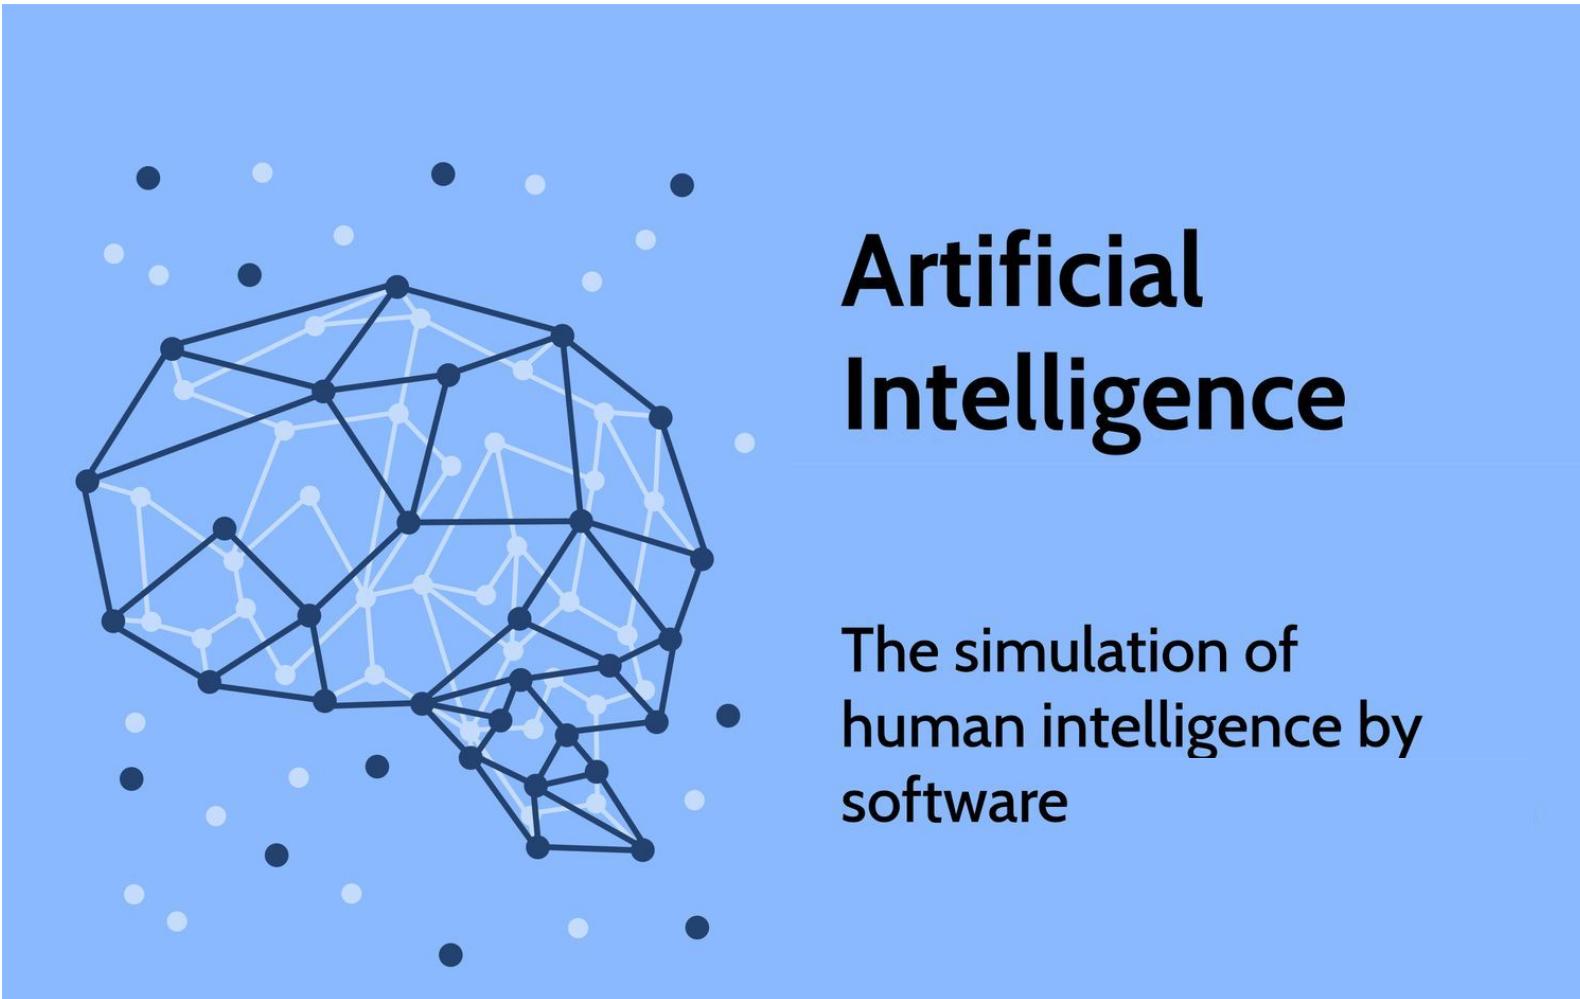
AI/ML Lunch & Learn: Object Detection

2025-01-08

michael.akridge@noaa.gov

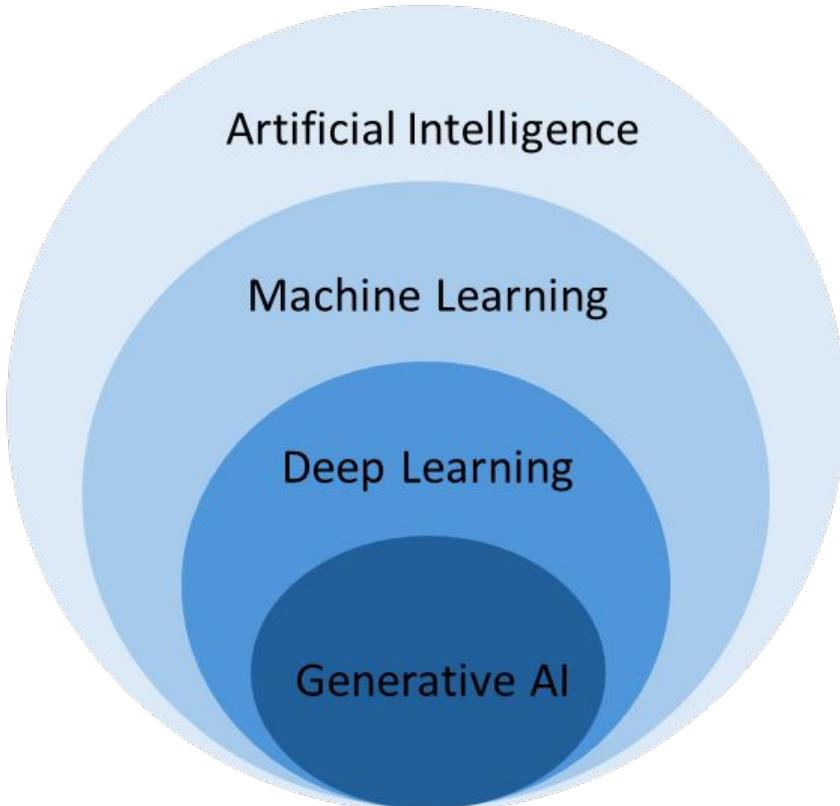
What is Object Detection?

AI



NOAA
FISHERIES

Object Detection: Deep learning powered AI



 Machine learning model

An algorithm that identifies patterns in a set of data and makes predictions over it.

What is AI? A complex pattern learner.

Artificial Intelligence (AI): The field of study making computer algorithms exhibit human-like intelligence and behaviors.

Machine Learning (ML): A subfield of AI using statistical methods to enable systems to learn from data.

Deep Learning (DL): A submethod of ML using deep neural networks to enable systems to learn from data.

Generative AI: A subset of models within DL that learn (from data) how to create brand new content in the form of text, images, etc.

Object Detection: Computer Vision

A CV field w/ a focus on identifying and locating specific objects in images

Key Components

1. Object Localization:

Identifying the position of objects in the form of bounding boxes (rectangles) or other techniques

2. Object Classification:

Assigning a class or label to each detected object.



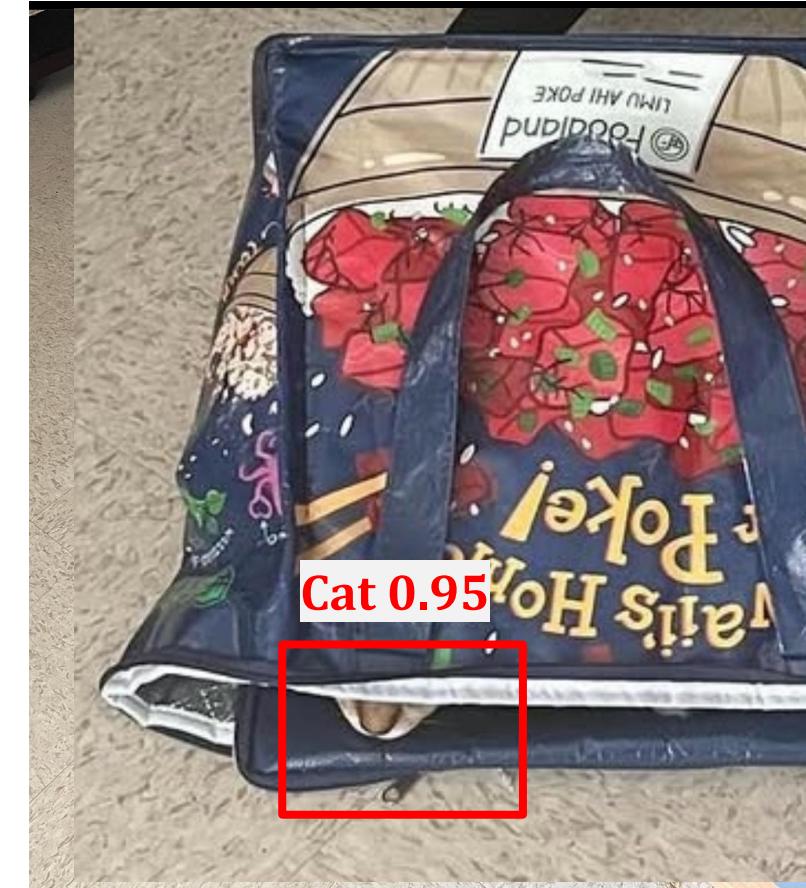
NOAA
FISHERIES

Is this a Cat?

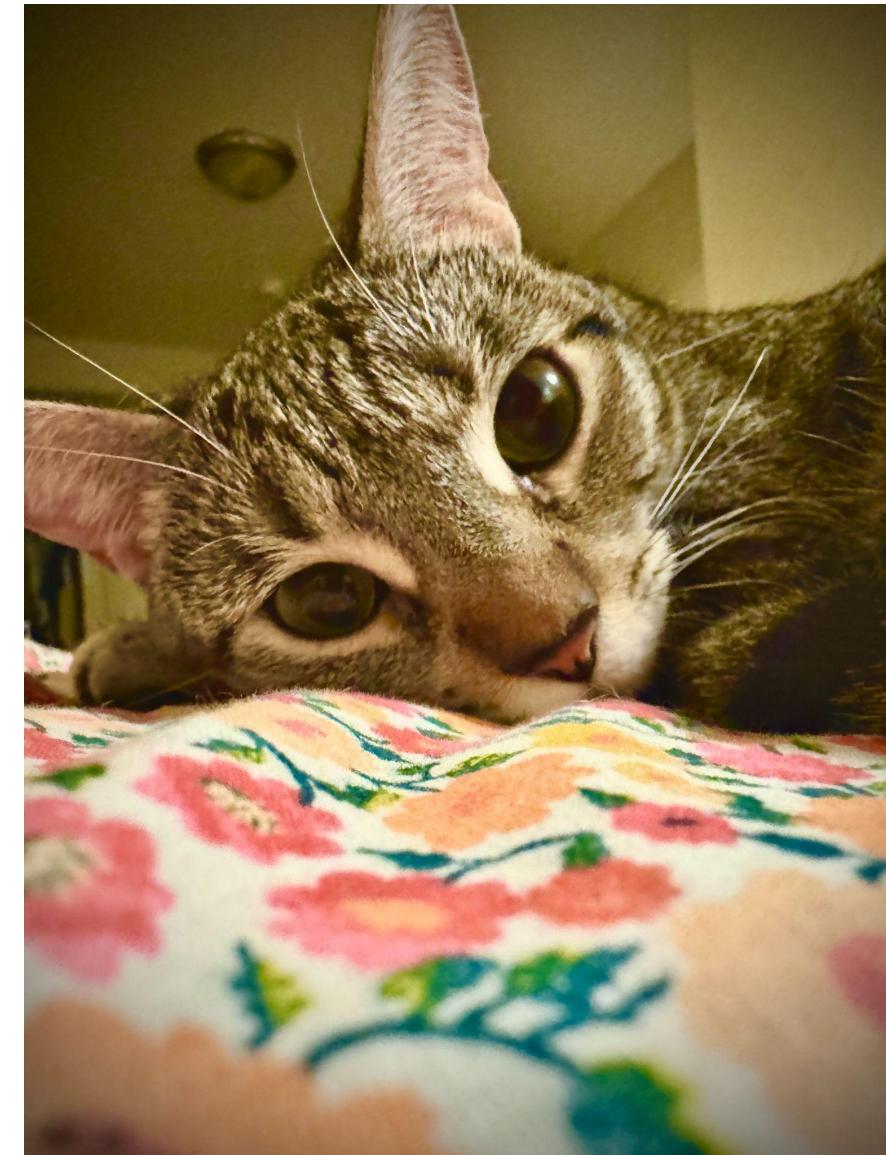


Image Classification

What is in the image & where is it?



Object Detection



NOAA
FISHERIES

Examples In Hawaii: Wildfire Detection System



Hawaiian
Electric

Hawaiian Electric deploys high resolution video cameras with artificial intelligence (AI) for early detection of wildfires

Strategically placed cameras will monitor 24/7 for ignitions

Release Date: 7/16/2024

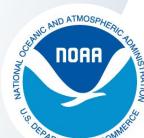
[Download PDF](#)

HONOLULU, July 16, 2024 – Hawaiian Electric has begun deploying a network of high-resolution video cameras using artificial intelligence (AI) technology to provide enhanced situational awareness and early detection of ignitions in elevated fire risk areas near company infrastructure. The public will also have access to the live feeds from any of the cameras.

Hawaiian Electric recently installed the first camera station in Lahaina and has plans to deploy a total of 78 stations in elevated fire risk areas on the five islands served by the company, with each location having two cameras to provide a full 360-degree view. The camera feeds will be monitored 24 hours a day, seven days a week. The \$14 million project is the latest step in Hawaiian Electric's ongoing effort to reduce the risk of wildfires associated with company equipment.

"We are continuing to take action to address the growing risks from wildfires across our service territory using a variety of technologies and methods," said Jim Alberts, Hawaiian Electric senior vice president and chief operations officer. "Installing publicly viewable AI-assisted video cameras in elevated fire risk areas will enable the company, fire agencies, and emergency operations centers the ability to identify potential wildfires early and respond quickly."

Hawaiian Electric signed a five-year contract with California-based ALERTWest. ALERTWest will install and maintain the camera stations as well as provide around-the-clock monitoring for potential ignitions by experienced wildfire safety professionals. Approximately 50% of the project costs will be covered by federal funds allocated under the Federal Infrastructure Investment and Jobs Act (IIJA) estimated at \$90 million in grant funding covering various costs related to Hawaiian Electric's resiliency and wildfire mitigation work. Hawaiian Electric also will be able to achieve cost savings by leveraging its existing telecom network to provide communications support for the project.



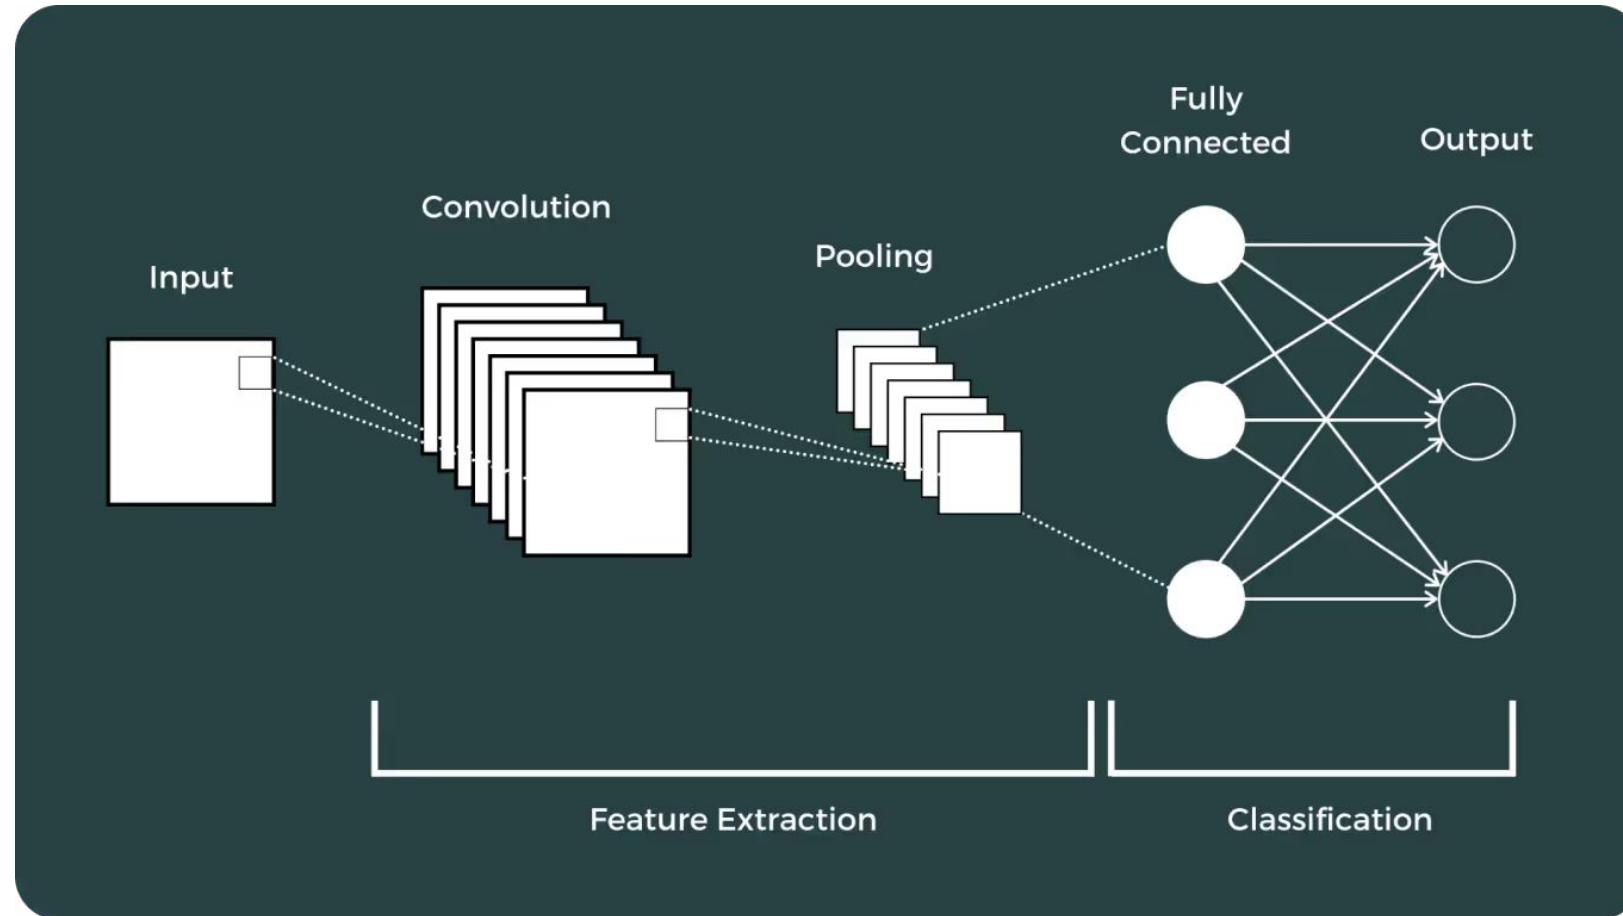
NOAA
FISHERIES

Examples In Hawaii: Wildfire Detection System



How does this all work?

Convolutional Neural Networks (CNNs)



“Convolutional refers to a process that involves intertwining two sources of information.”

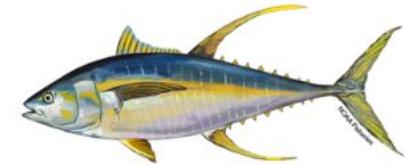
A bunch of Data



+ **Bunch
of Math** =

Classifier

Pacific Yellowfin Tuna
Thunnus albacares



Also Known As
Ahi, Kihada

Quick Facts

REGION Pacific Islands, West Coast



FISWATCH
U.S. SEAFOOD FACTS



NOAA
FISHERIES



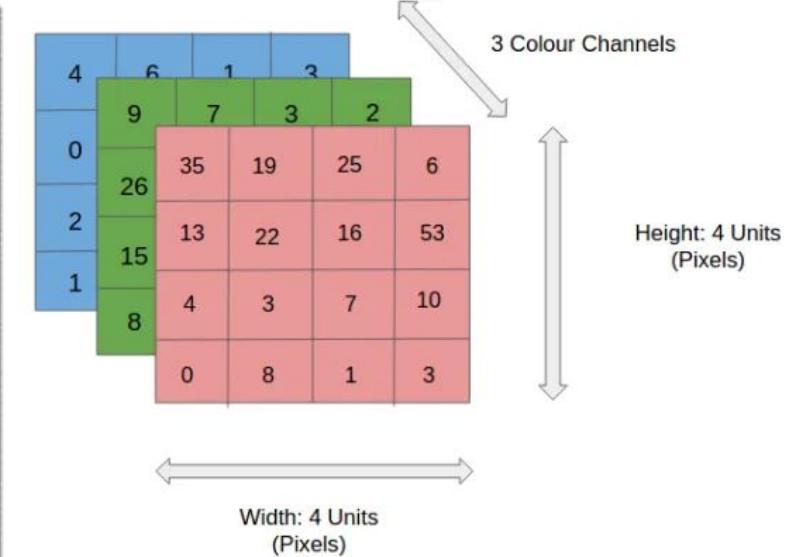
08	02	22	97	38	15	00	40	00	75	04	05	07	78	52	12	50	77	91	08
49	49	99	40	17	81	18	57	60	87	17	40	98	43	69	48	04	56	62	00
81	49	31	73	55	79	14	29	93	71	40	67	51	48	30	03	49	13	36	65
52	70	95	23	04	60	11	42	62	21	68	56	01	32	56	71	37	02	36	91
22	31	16	71	51	62	03	89	41	92	36	54	22	40	40	28	66	33	13	80
24	47	30	60	99	03	45	02	44	75	33	53	78	36	84	20	35	17	12	50
32	98	81	28	64	23	67	10	26	38	40	67	59	54	70	66	18	38	64	70
67	26	20	68	02	62	12	20	95	63	94	39	63	08	40	91	66	49	94	21
24	55	58	05	66	73	99	26	97	17	78	78	96	03	14	88	34	89	63	72
21	36	23	09	75	00	76	44	20	45	35	14	00	61	33	97	34	31	33	95
78	17	53	28	22	75	31	67	15	94	03	80	04	62	16	14	09	53	56	92
16	39	05	42	96	35	31	47	55	58	88	24	00	17	54	24	36	29	85	57
86	56	00	48	35	71	89	07	05	44	44	37	44	60	21	58	51	54	17	58
19	80	81	68	05	94	47	69	28	73	92	13	86	52	17	77	04	89	55	40
04	52	08	83	97	35	99	16	07	97	57	32	16	26	26	79	33	27	98	66
03	94	68	87	57	62	20	72	03	46	33	67	46	55	12	32	63	93	53	69
04	42	16	73	38	25	39	11	24	94	72	18	08	46	29	32	40	62	76	36
20	69	36	41	72	30	23	88	31	69	99	69	82	67	59	85	74	04	36	16
20	73	35	29	78	31	90	01	74	31	49	71	48	56	51	16	23	57	05	54
01	70	54	71	83	51	54	69	16	92	33	48	61	43	52	01	89	1	47	48

What the computer sees

image classification

82% cat
15% dog
2% hat
1% mug

<https://cs231n.github.io/>
<https://github.com/cs231n/cs231n.github.io>
<https://online.stanford.edu/courses/soe-ymls-machine-learning-specialization>
<https://online.stanford.edu/programs/artificial-intelligence-graduate-certificate>



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

4		

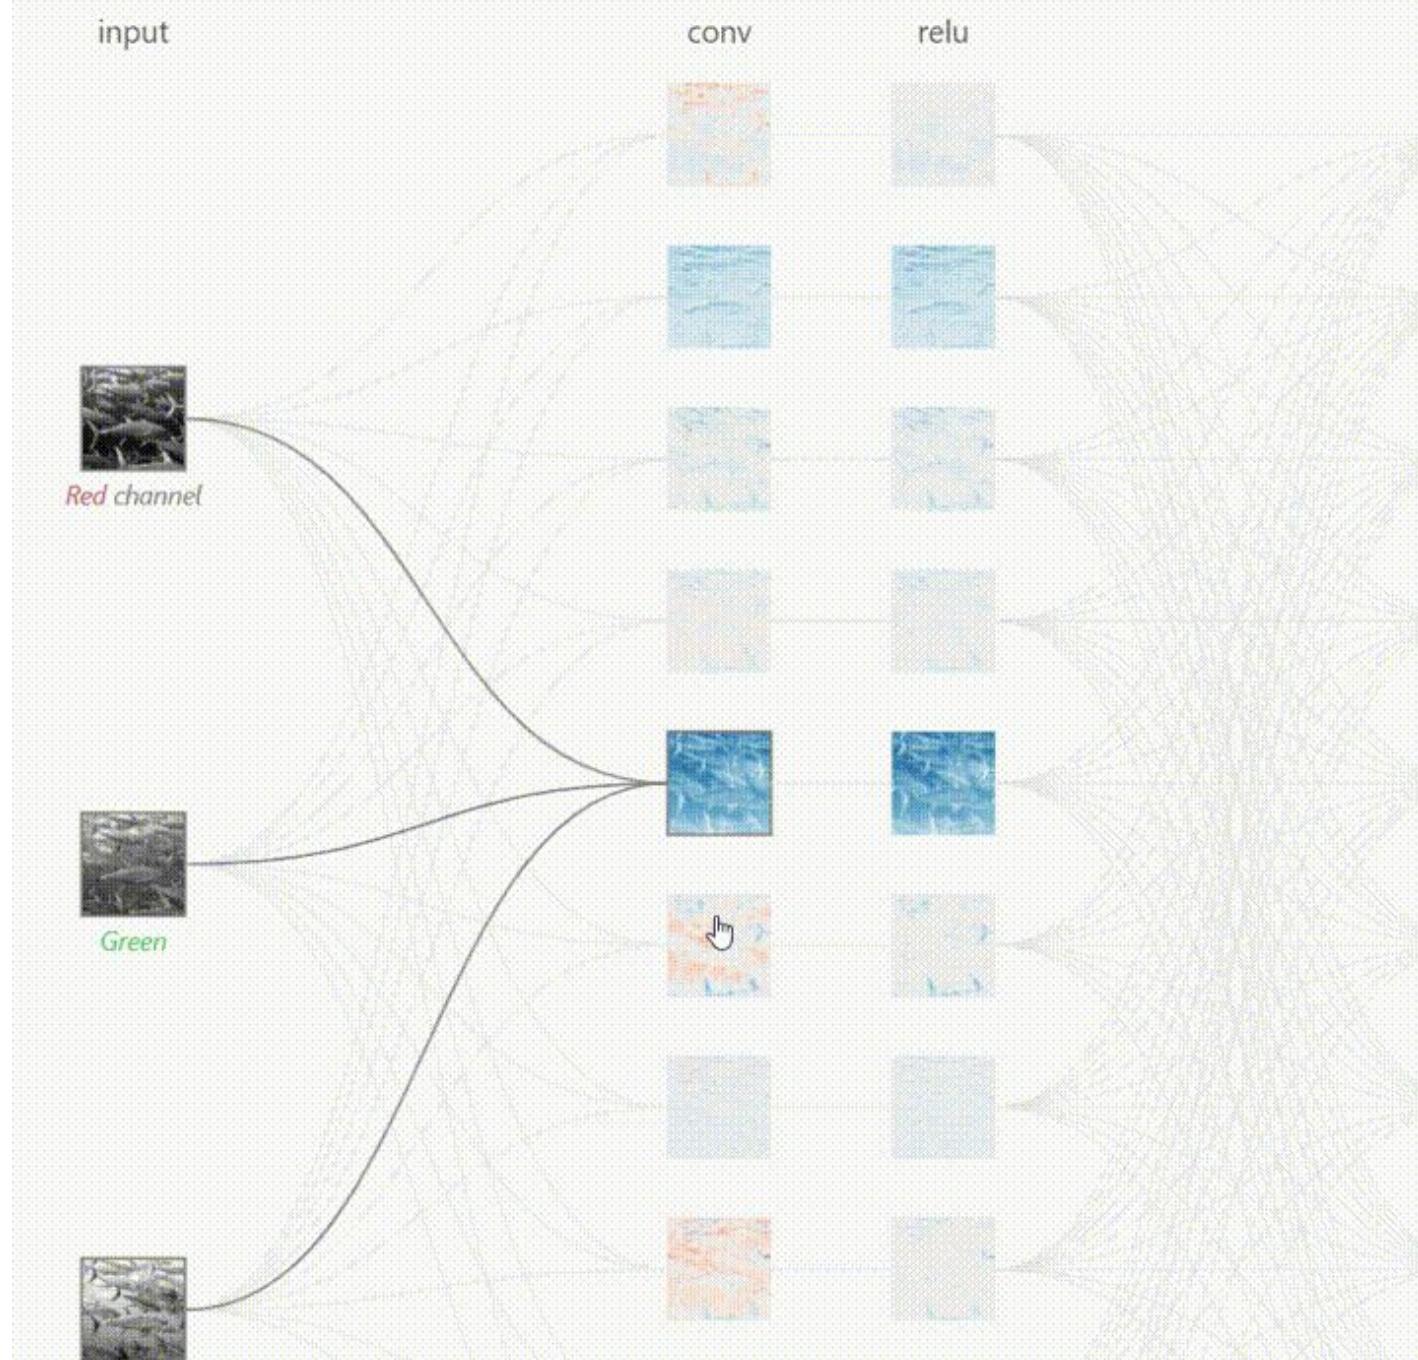
Image

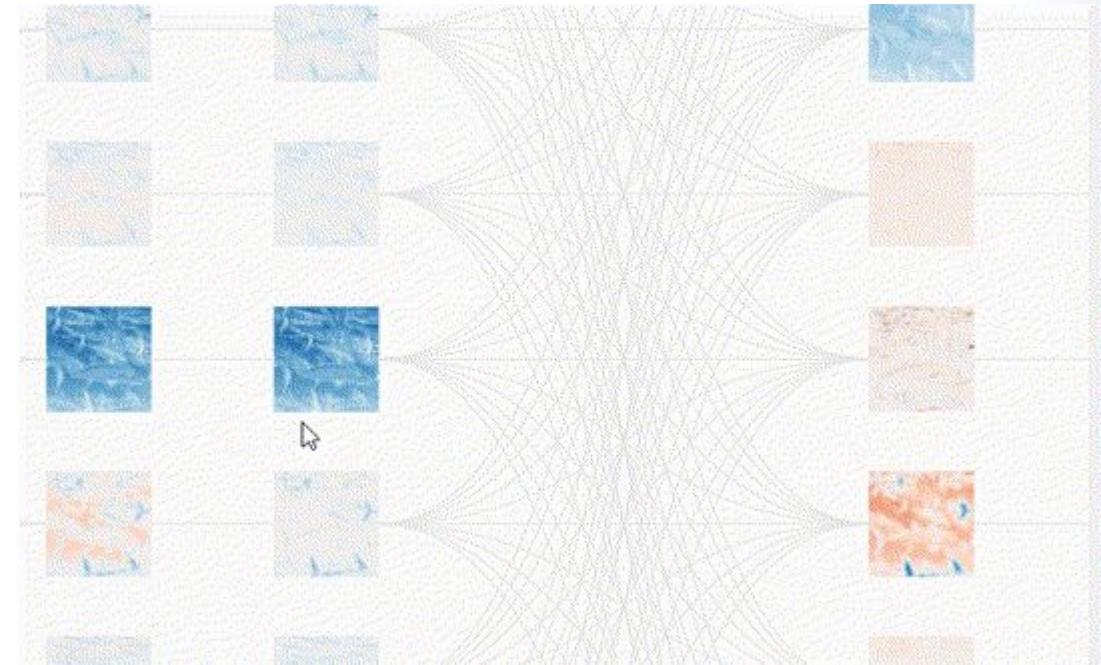




The convolutional layer is the first layer of a CNN and is where most of the computations take place.

It uses a filter or kernel, which is a small matrix of weights, to move across an input image to detect specific features.





Learn More here:

<https://poloclub.github.io/cnn-explainer/>



NOAA
FISHERIES



NOAA
FISHERIES

Common Types Object Detectors

Common Types of Object Detectors

(Two or Multi-Stage Detectors)

Region-Based Object Detectors

- R-CNN
- Fast R-CNN
- Faster R-CNN

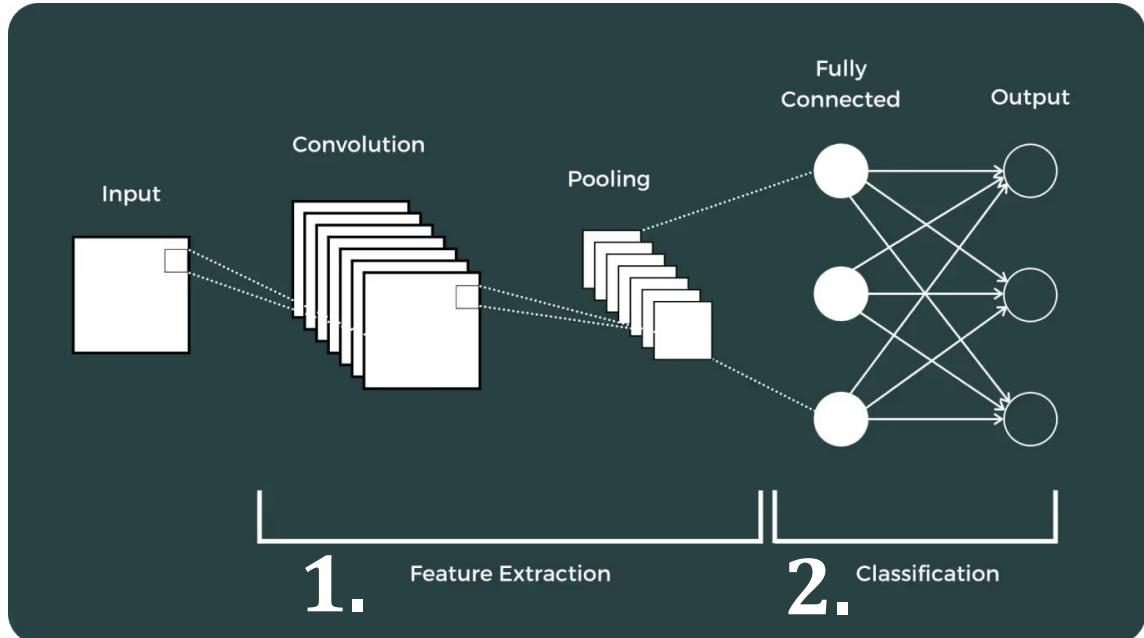
(One Stage Detectors)

Single Shot Object Detectors

- YOLO (You Only Look Once)
- SSD (Single Shot MultiBox Detector)

Common Types of Object Detectors

Region-Based Object Detectors (Two or Multi-Stage Detectors)

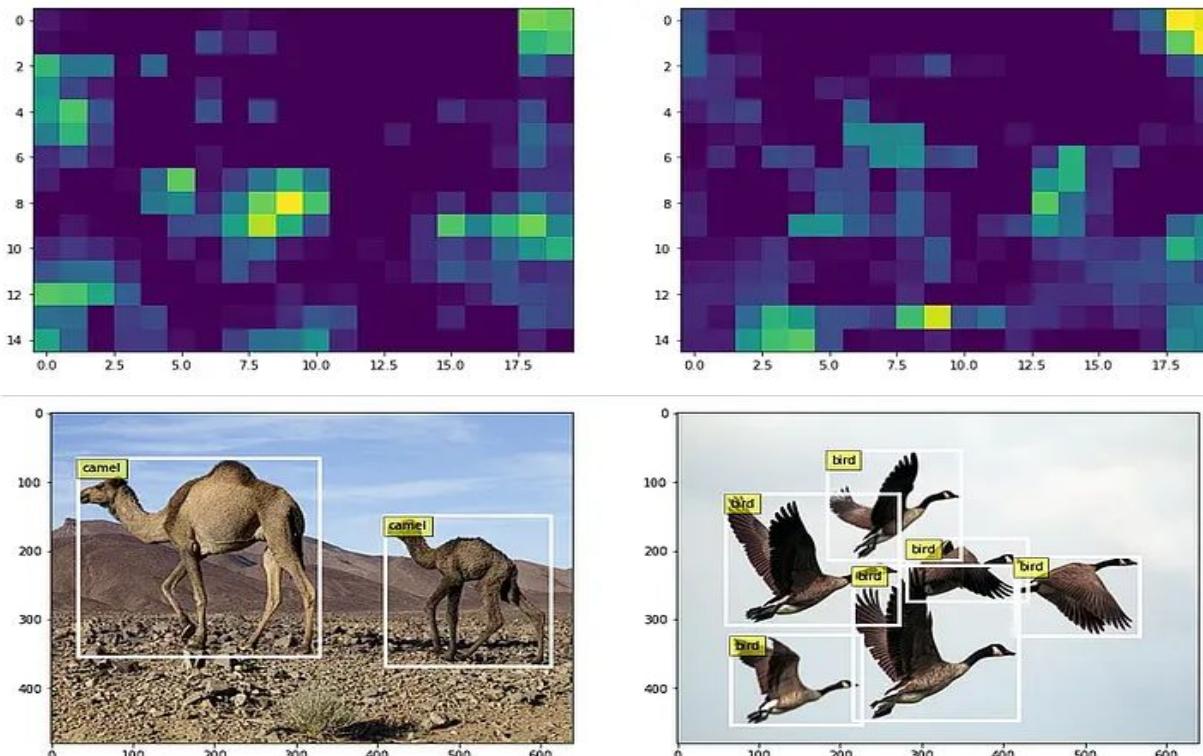


Single Shot Object Detectors (One Stage Detectors)

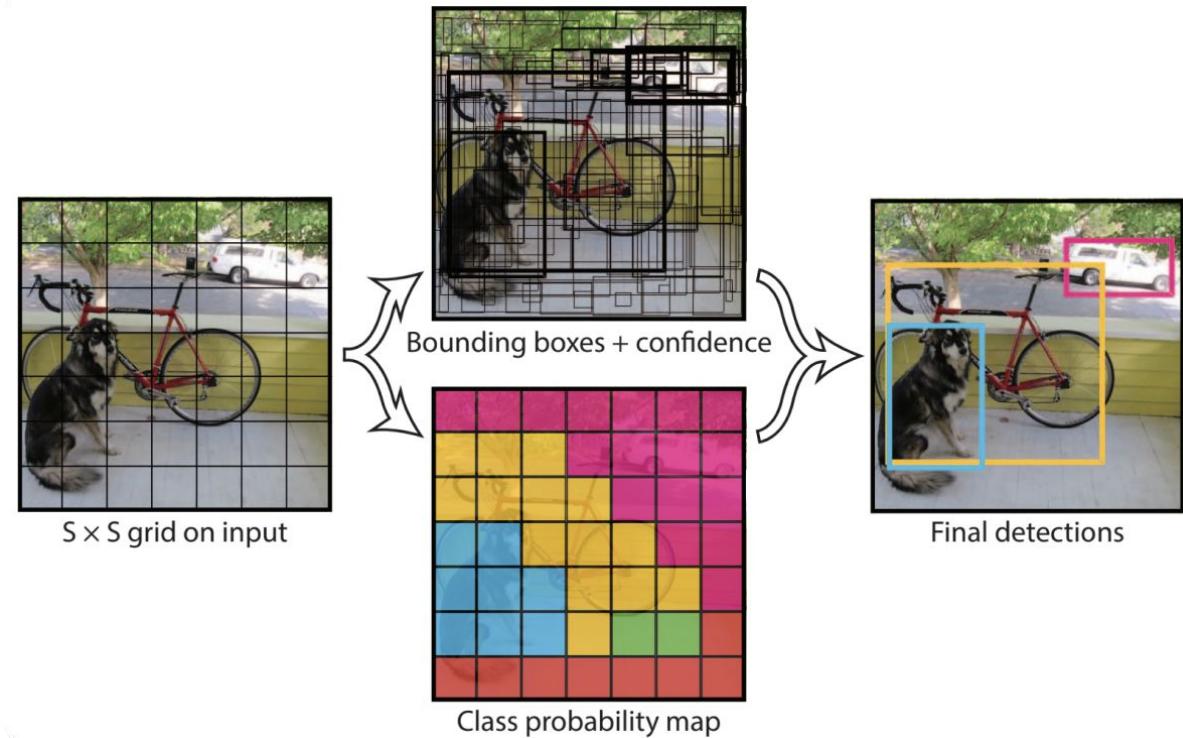


Common Types of Object Detectors

Region-Based Object Detectors (Two or Multi-Stage Detectors)



Single Shot Object Detectors (One Stage Detectors)



How does training an Object Detection model work?

Object Detection: Several Different Options

Training Method	Description	Use Case	Pros	Cons
From Scratch	Train from random initialization on new data	Custom models, large datasets	Full control	Data and compute-intensive
Fine-tuning	Adapt a foundational (pre-trained) model to a specific task	Domain-specific tasks	Efficient and high-performance	Risk of overfitting
Zero-shot Learning	Leverage pre-trained models without training	No labeled data available	No task-specific data needed	limited performance on edge-case/complex data
Transfer Learning	Use a model trained for a related task as a starting point	Related but different tasks	Speeds up training	Requires task/domain compatibility

What in an Object Detection Model?

Object Detection Model

The model is like a blueprint.

- It's a set of instructions that a computer can follow to identify objects in an image or video.

Model Weights:

- Model weights are the stored patterns(rules) the model has learned during training.

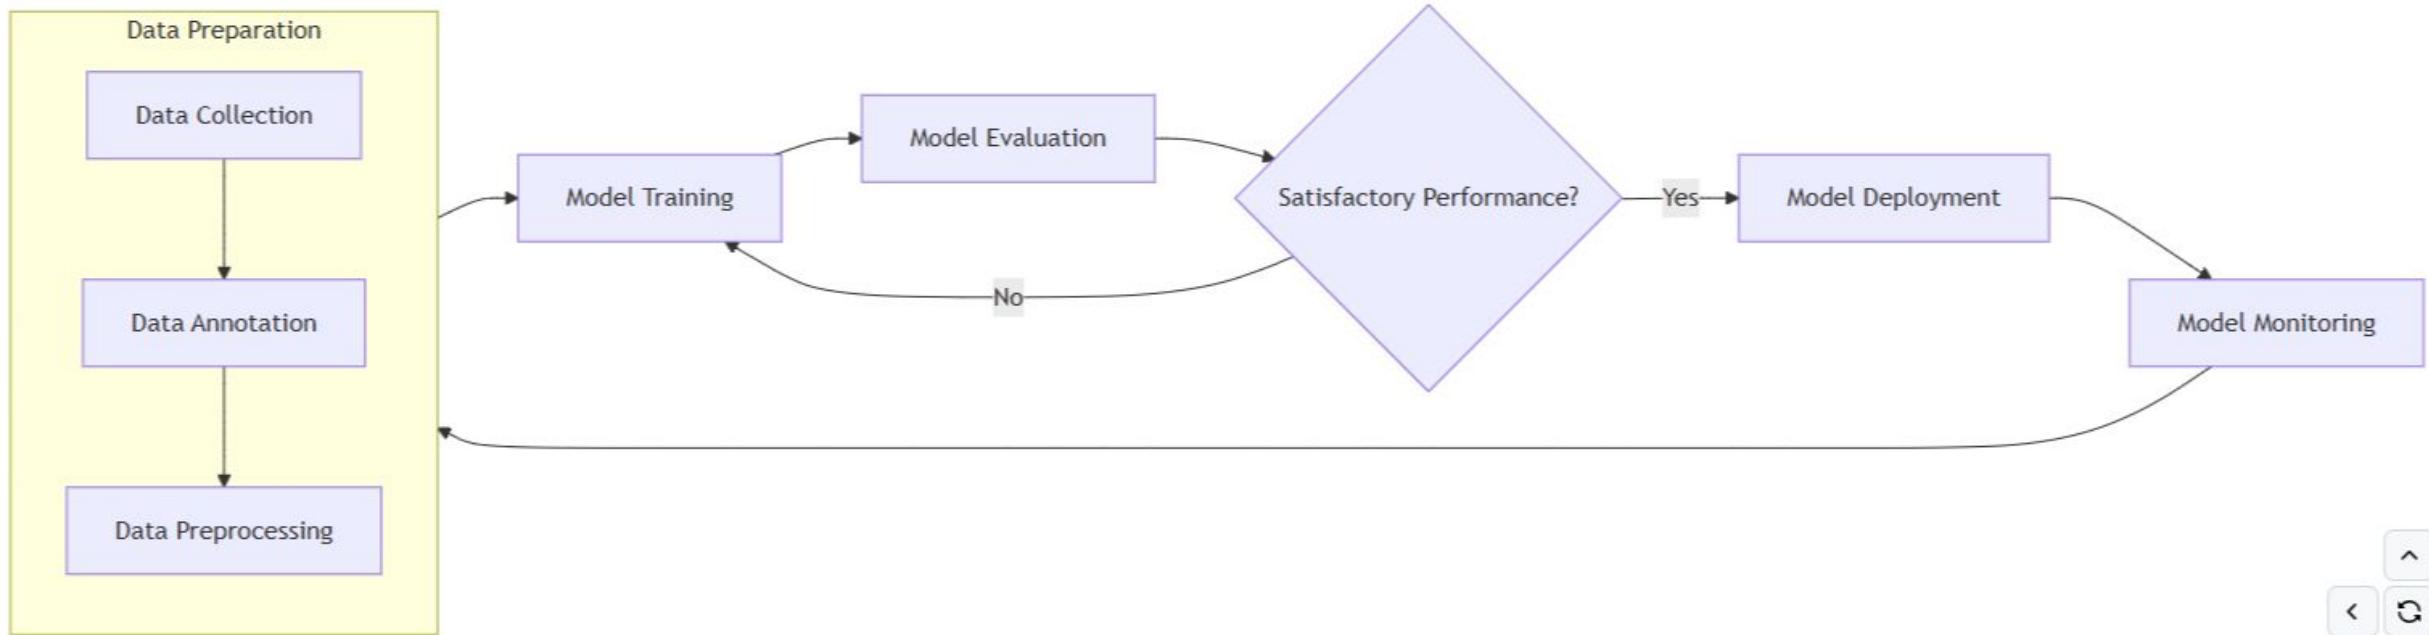
Common Model files

- .pt = model architecture & model weights.
- .pth = just model weights

How to get started?

Getting a Project Started

- Create Project Plan Summary
 - Define Scope
 - Create a class list



80% Data Management 20% Model Building

“In machine learning, 80% of your time is spent preparing and cleaning the data, and only 20% of the time is spent on the actual model building.”

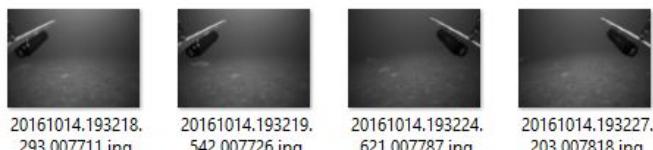
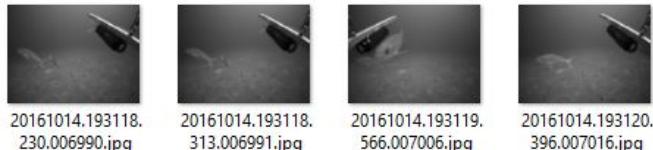
- Andrew Ng (former Stanford AI Director/Professor)

AI Ready Dataset:

What Does it Look Like?

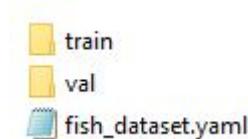
AI-Ready Training Dataset

1. Images



+ 2. Annotations (labels) = Training Dataset

- 📄 20161014.192648.087.003747.txt
- 📄 20161014.192742.649.004402.txt
- 📄 20161014.192742.980.004406.txt
- 📄 20161014.192742.982.004406.txt
- 📄 20161014.192743.065.004407.txt
- 📄 20161014.192743.146.004408.txt
- 📄 20161014.192743.147.004408.txt
- 📄 20161014.192743.311.004410.txt
- 📄 20161014.192743.315.004410.txt
- 📄 20161014.192743.396.004411.txt
- 📄 20161014.192743.399.004411.txt
- 📄 20161014.192743.480.004412.txt
- 📄 20161014.192743.482.004412.txt

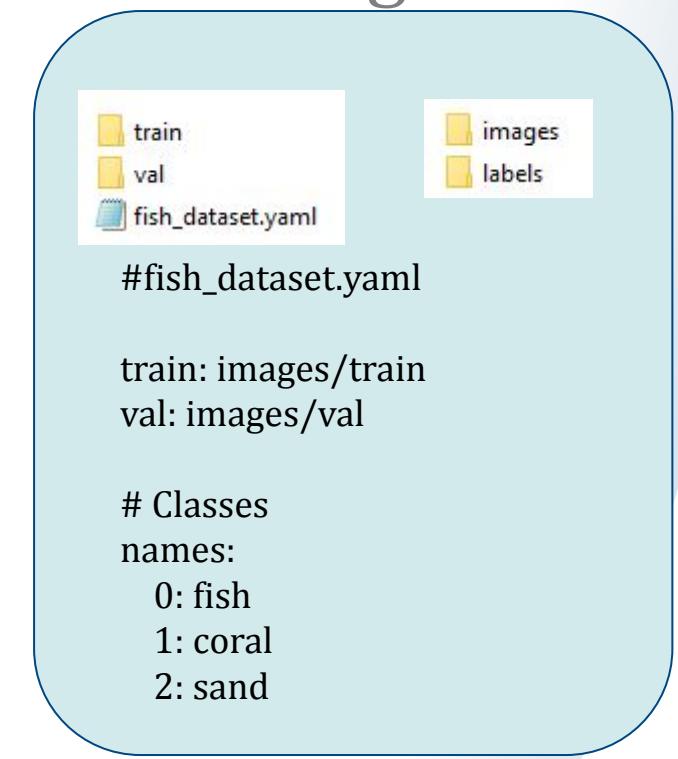
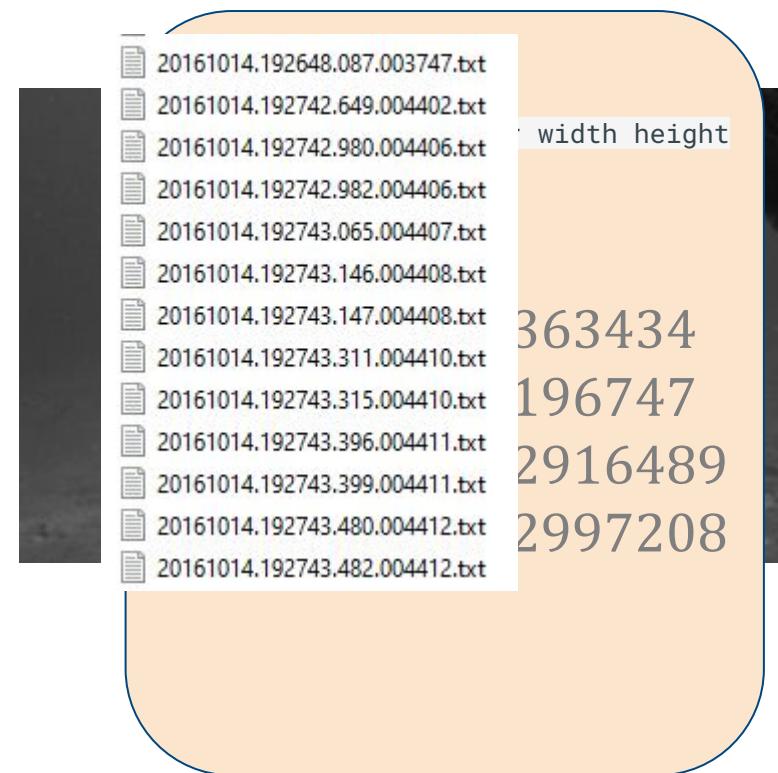


AI-Ready Training Dataset

1. Images



+ 2. Annotations (labels) = Training Dataset



NOAA
FISHERIES

Is your Data AI Ready?

Check out the ESIP AI-Ready Data Checklist

Link:

<https://github.com/ESIPFed/data-readiness/blob/main/checklist-published/ai-ready-data-checklist-v.1.0.md>

About the checklist

The checklist is developed using the 2019 draft readiness matrix developed by the Office of Science and Technology Policy Subcommittee on Open Science as a basis. The checklist has been improved based on further research and user feedback. Definitions for some concepts are listed at the end of this document. This checklist is developed through a collaboration of ESIP Data Readiness Cluster members include representatives from NOAA, NASA, USGS, and other organizations. The checklist will be updated periodically to reflect community feedback.

How to use the checklist

Ideally for AI-ready assessment, a dataset should be defined as the minimum measurable bundle (i.e., a physical parameter/variable of observational datasets or model simulations). The assessment at this scale will enable better integration of data from different sources for research and development. However, it can be an intensive process for manual assessment without automation. Therefore, we recommend current assessments be done on the data file level. If the dataset has different versions, the checklist should be applied to each dataset type (e.g. raw, derived).

How to cite this checklist:

ESIP Data Readiness Cluster (2023): Checklist to Examine AI-readiness for Open Environmental Datasets v.1.0. ESIP. Online resource.
<https://doi.org/10.6084/m9.figshare.19983722.v1>

History

Version: 1.0;
Last updated: June 21, 2023.

General Information

- Link to the dataset landing page (*questions below could be automatically filled from the landing page in the future*)
 - Name of the dataset
 - The current version of the dataset
 - Point of contact for the dataset
 - When was the dataset originally published?
- Is this raw data or a derived/processed data product? *Raw/Derived*
- Is this observational data, simulation/model output, or synthetic data? *Observed/Modeled/Synthetic*
- Is the data single-source or aggregated from several sources? *Single-source/Aggregated*

Data Quality

- Questions on timeliness:
 - Will the dataset be updated? *Yes, it will be updated. / No, it will not be updated*
 - If the data will be updated, how often will it be updated? [choose from]
 - Updated when new data are added.
 - Choose the update frequency that best describes the dataset: *near-real-time (irregularly) / hourly or sub-hourly / daily or sub-daily / weekly / monthly / yearly / longer than a year*
 - Will there be different stages of the update (e.g., updated with preliminary data first and replaced by a later update of the full report)? *Yes / No / Not applicable*

ESIP Data Readiness Cluster

This is the repository for the Data Readiness Cluster to maintain the AI-ready data checklist. The cluster is a community-driven group focusing on developing recommendations and community standards on AI-ready open environmental data. Although the work currently focuses on environmental data, the product could be extended to data from other domains.

The goal of AI-readiness assessment:

- For data producers/providers, the purpose is to understand to what extent the data being assessed meets the common research data management practices and principles that are relevant to AI/ML application development. The assessment result can be used to justify targeted improvements for the dataset when resources become available.
- For projects generating new datasets, the purpose of the AI-readiness checklist can be used to guide the development of the dataset. For example:
 - What documentation do you want to provide accompanying the dataset?
 - Do you have a proper data quality assessment that will make the development of downstream AI/ML applications efficient?

How to use the checklist for assessment:

The current version of the checklist is available [here](#) (last updated 2023-12-20). The checklist will be maintained and updated by the community.

To assist with the assessment, we have created a fillable [Google sheet template](#). You can make a copy of the Google sheet for your assessment. Each dataset should be assessed separately as the checklist is designed for individual datasets. More effort is ongoing to address the need for linked datasets.

If you are in the early stages of developing AI/ML applications with open environmental datasets. We encourage you to assess the input data used for your applications. Although you may not have the ability to change other people's datasets, this will help you document the effort spent on preparing the dataset for your development.



What's available to help the process?

Platforms vs Libraries vs Tools

- What to know:
 - are they easy to install and use?
 - Free/open source?
 - Maintained/Support/documentation?
- What's important to ask:
 - Can it export the annotation data out in a common format?
 - Can you export the model?

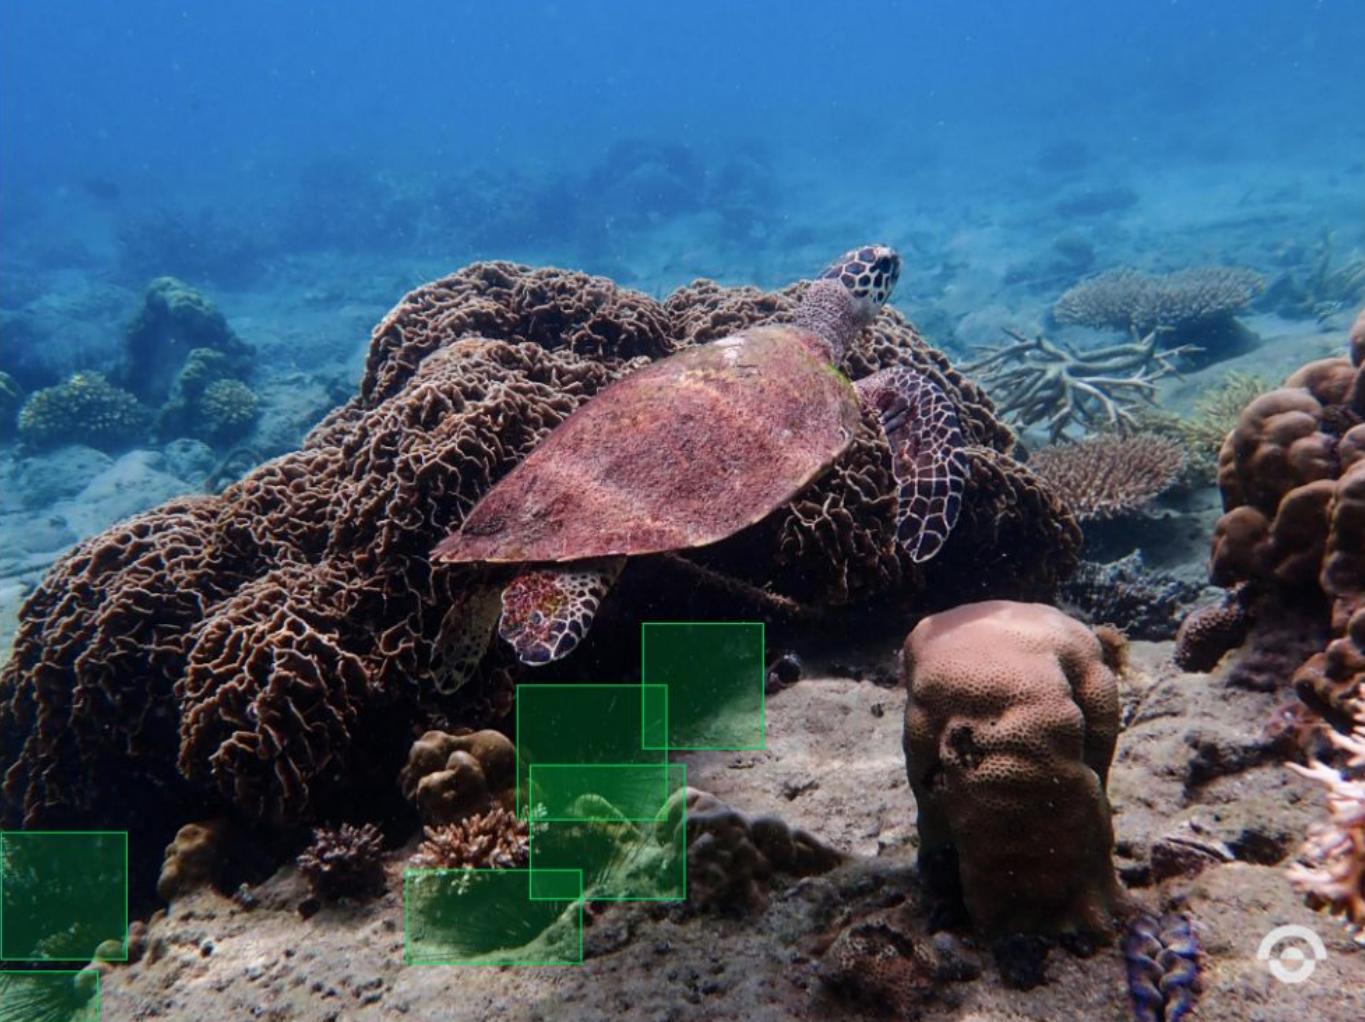
Type	AI/ML Libraries	Tools	Platforms/Software
Example	Pytorch, Tensorflow, Ultralytics	Label Studio, CVAT, TagLab	Label Studio, CVAT, Roboflow, Coralnet, Viame, CleamML
Use	for building, training, deploying, developing and more	Data Prep, Annotation, labeling of datasets (optional ML backends)	Integrate multiple functionalities

Annotation Tools

Image and Video Annotation

Name	Image	Video	Free	Open Source	AI Assisted Annotation	Notes
Label Studio	✓	✓	✓	✓	✓	Multi-type data labeling and annotation tool with standardized output format.
CVAT	✓	✓	✓	✓	✓	Open-source tool for video and image annotation.
VIAME	✓	✓	✓	✓	✓	See documentation .
Scalabel	✓	✓	✓	✓		Modern video and image annotation platform.
Tator	✓	✓	✓	✓		See Tutorials and Source Code .
BIIGLE	✓	✓	✓	✓		Check GitHub .
FishID	✓	✓	✓	→ SOON		Email to request access: fishidglow@outlook.com .
TagLab	✓	✗	✓	✓	✓	
VoTT (Archived)	✓	✗	✓	✓	✗	Officially archived, last update in 2020.
SEAMS	✓	✗	✓	✓	✗	SEafloor Annotation and Mapping Support Open-source tool for video and image annotation.

Ani



The image shows a hawksbill sea turtle resting on a coral reef. Several green rectangular boxes highlight specific areas of the turtle's shell and surrounding coral, indicating regions of interest for labeling.

Project Details:

- Project: urchin_dataset_v1
- Task: Labeling
- Owner: michael.akridge
- Last Update: #69 2 minutes ago

Selection Details:

- Regions:
- Relations:

Regions:

- 1 urchin
- 2 urchin
- 3 urchin
- 4 urchin
- 5 urchin
- 6 urchin

Relations:

- Manual
- By Time ↑

Labels:

- urchin

Annotations:

- urchin 1

Buttons:

- Update

Cloud Storage

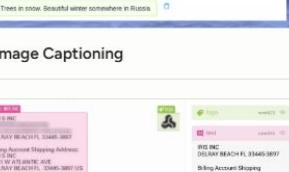
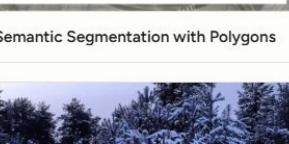
Use cloud or database storage as the source for your labeling tasks or the target of your completed annotations.

Source Cloud Storage

[Add Source Storage](#)

Create Project

- Computer Vision >
- Natural Language Processing
- Audio/Speech Processing
- Conversational AI
- Ranking & Scoring
- Structured Data Parsing
- Time Series Analysis
- Videos
- Generative AI
- Custom template



Target Cloud Storage

[Add Target Storage](#)

Project Name

Data Import

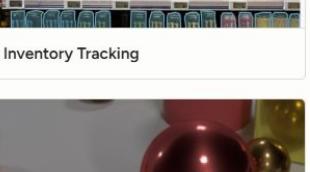
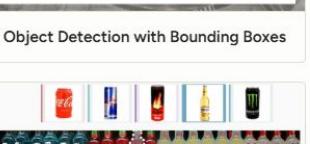
Labeling Setup

Semantic Segmentation with Polygons

Image Captioning

Image Classification

Inventory Tracking



Export data

You can export dataset in one of the following formats:

JSON

List of items in raw JSON format stored in one JSON file. Use to export both the data and the annotations for a dataset. It's Label Studio Common Format

JSON-MIN

List of items where only "from_name", "to_name" values from the raw JSON format are exported. Use to export only the annotations for a dataset.

CSV

Results are stored as comma-separated values with the column names specified by the values of the "from_name" and "to_name" fields.

TSV

Results are stored in tab-separated tabular file with column names specified by "from_name" "to_name" values

image segmentation object detection

COCO

Popular machine learning format used by the COCO dataset for object detection and image segmentation tasks with polygons and rectangles.

image segmentation object detection

Pascal VOC XML

Popular XML format used for object detection and polygon image segmentation tasks.

image segmentation object detection

YOLO

Popular TXT format is created for each image file. Each txt file contains annotations for the corresponding image file, that is object class, object coordinates, height & width.

image segmentation object detection

YOLOv8 OBB

image segmentation object detection

[See the documentation to contribute a template.](#)

Export

Dataset Best Practices

Dataset Best Practices: For training an Object Detection Model

Standard Dataset Split

- **Training Set (70-80%):** Used to train your model.
- **Validation Set (15-20%):** Used to tune and validate the model during training
 - **Optional - Test Set (15-20%):** Used to evaluate the model performance on unseen data after training



Best Practices

- **Images per class:** ≥ 1500 images per class recommended
- **Instances per class:** ≥ 10000 instances (labeled objects) per class recommended
- **Image variety.** Representative of environment.
- **Label accuracy.**
 - Labels must closely enclose each object.
 - No space should exist between an object and its bounding box.
 - No objects should be missing a label.
- **Optional - Background images (1-10%).**
Background images are images with no objects that are added to a dataset to reduce False Positives (FP).
- Have unique images in each split

Dataset Management & Balance

roboflow Projects Universe Documentation Forum James Gallagher ↗

JAMES GALLAGHER

View on Universe

Construction S... Object Detection

Sharing Page

Upload

Unassigned

Annotate

Images 665

Generate

Versions 1

Deploy

Health Check

Upgrade

Construction Safety ➜ Dataset Health Check

Generated on August 15, 2023 at 12:01 pm. ⏪ Regenerate

Images	Annotations	Average Image Size	Median Image Ratio
665	5,567	0.92 mp from 0.05 mp to 14.68 mp	1080×720 wide

Annotations: 0 missing annotations, 0 null examples

Class Balance

all train valid test

Class	Count	Status
Person	1,147	over represented
Safety Cone	599	over represented
NO-Safety Vest	582	over represented
Hardhat	574	over represented
NO-Mask	491	over represented
Safety Vest	424	over represented
NO-Hardhat	402	over represented
Gloves	296	over represented
Mask	202	over represented
Excavator	166	under represented
dump truck	157	under represented
wheel loader	126	under represented
vehicle	122	under represented
Ladder	58	under represented
sedan	54	under represented
machinery	45	under represented
truck and trailer	35	under represented
van	28	under represented
SUV	16	under represented
trailer	15	under represented
mini-van	7	under represented
semi	7	under represented

NOAA FISHERIES

Lessons Learned

- **Data augmentation** techniques can help models generalize better
 - Many common data augmentation techniques (scale, shear, rotate, perspective, flip, mosaic, and mixup) are built into many training libraries
(Link: [YOLO Data Augmentation Explained](#))

Image Scale & Mosaic Augmentation Example

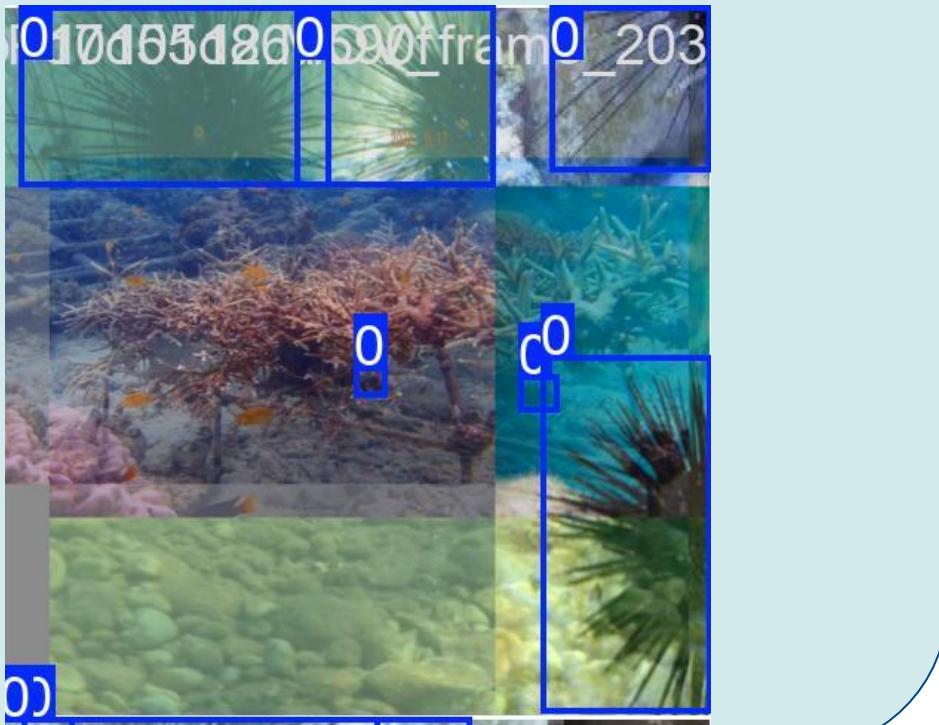
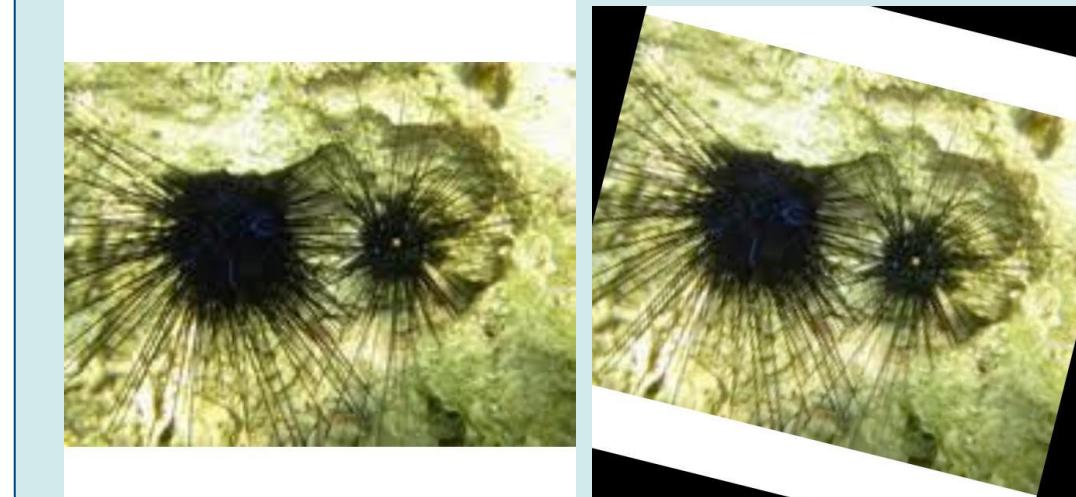
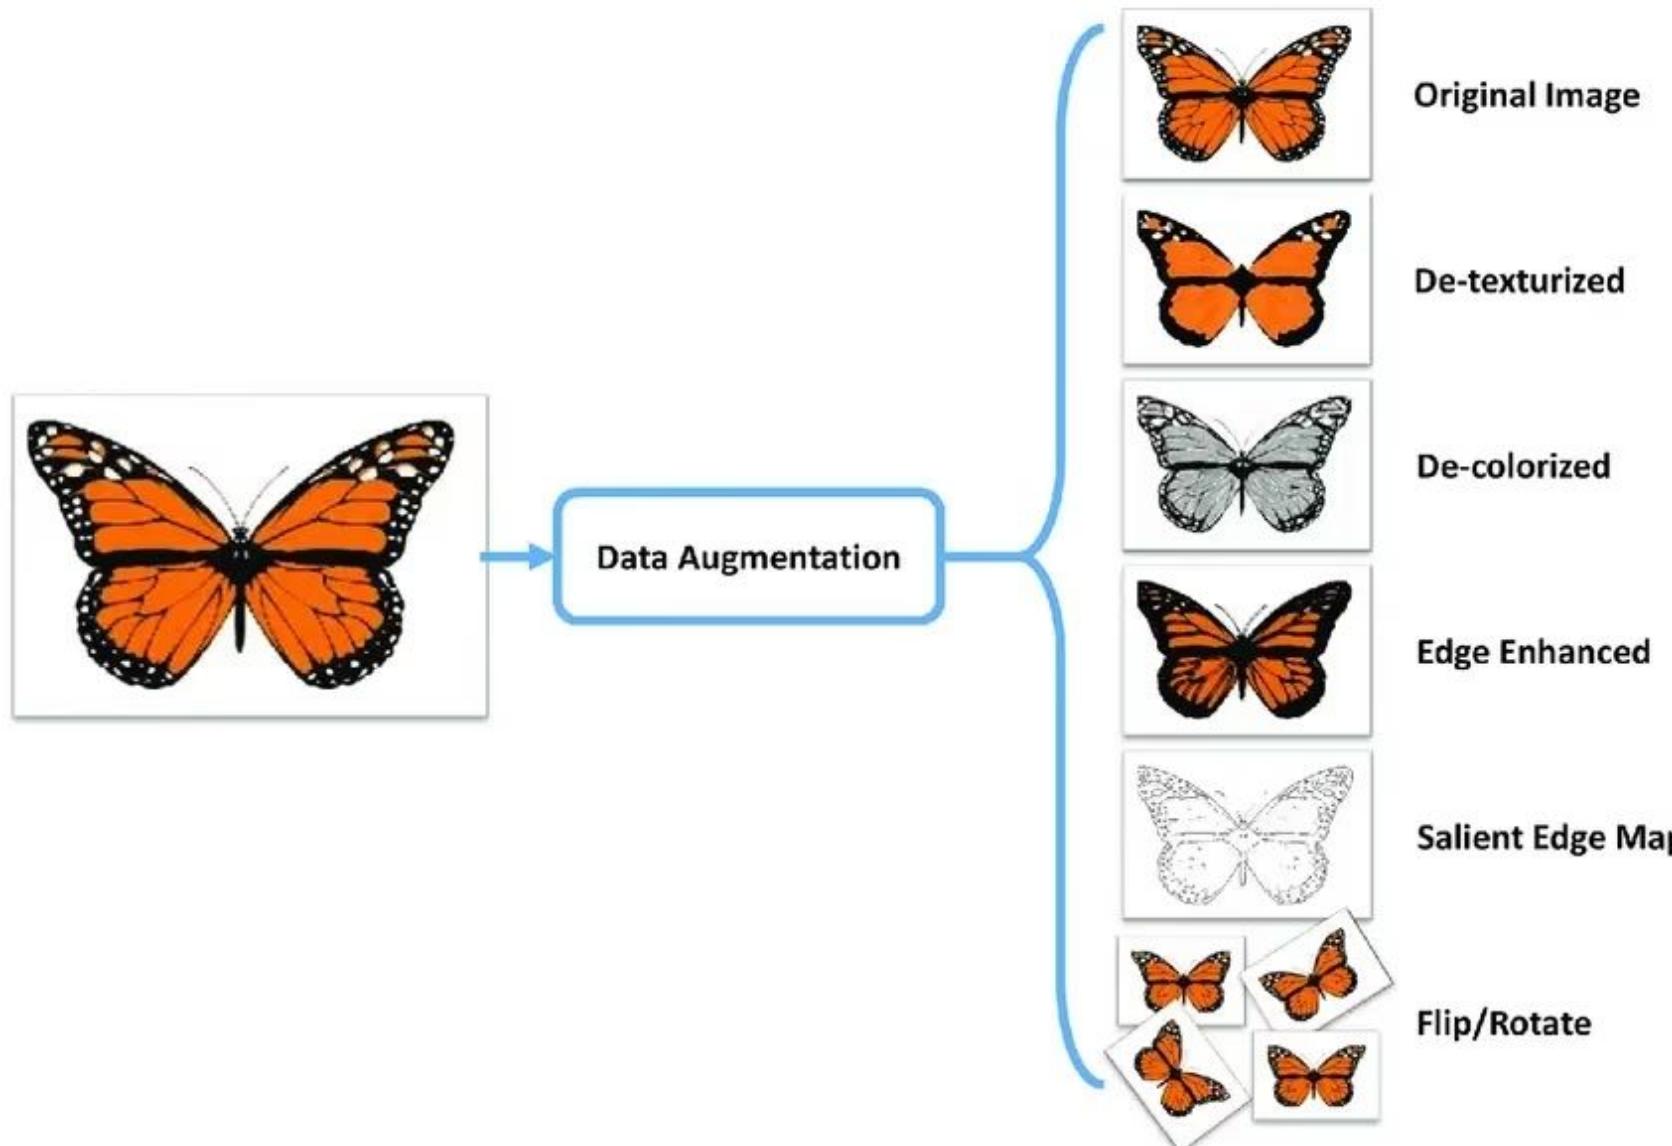


Image Angle/Degree rotation Augmentation Example

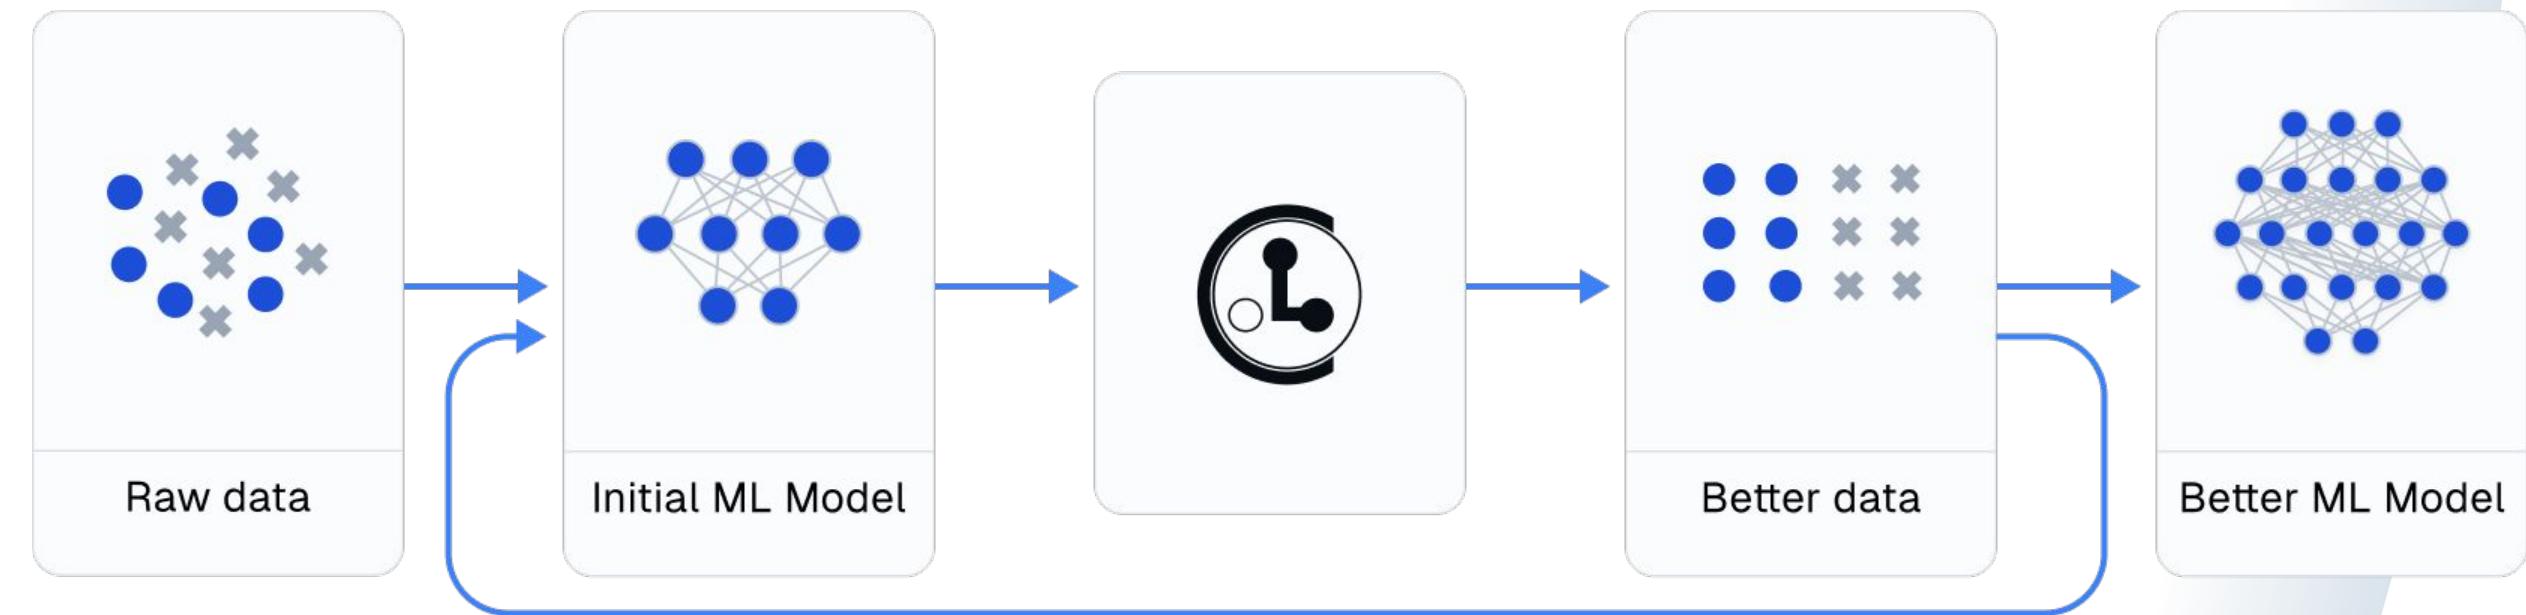




Turns 500 image dataset into a 5000 image dataset

After Data: When to Start Training a Model?

Data-Centric AI Model Flow: Start Right Now



fish_2016

7931 / 7931 ✓ 7931 - 0 0 0

28 Oct '24, 10:20 MI

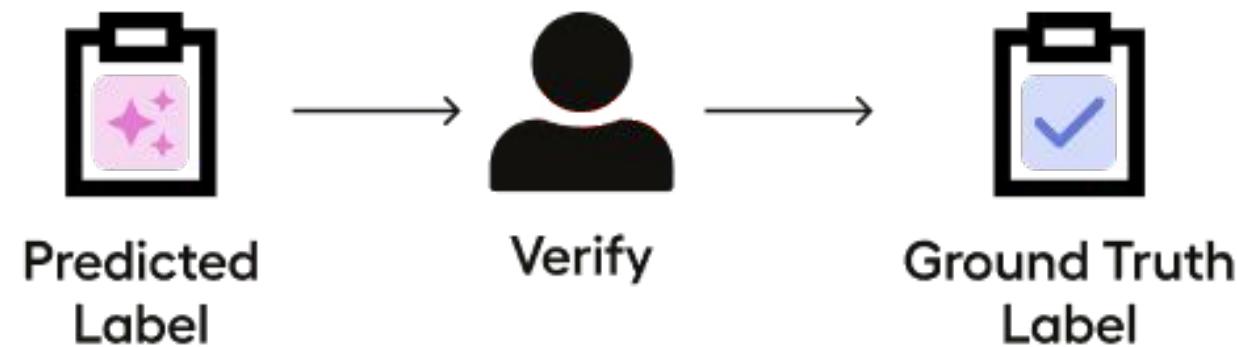
urchin_dataset_v1

356 / 356 ✓ 356 - 0 0 0

16 Oct '24, 09:58 MI

Have the Machine do the Hard work:

- Pre-annotate/autolabel data
- Interactive labeling
- Model evaluation and fine-tuning



Tasks: 7931 / 7931 Annotations: 7931 Predictions: 0

The YOLO (Ultralytics) Framework

YOLO(You Only Look Once)

Object Detection: YOLO(You Only Look Once)

YOLO: A Brief History

YOLO (You Only Look Once), a popular [object detection](#) and [image segmentation](#) model, was developed by Joseph Redmon and Ali Farhadi at the University of Washington. Launched in 2015, YOLO quickly gained popularity for its high speed and accuracy.

- [YOLOv2](#), released in 2016, improved the original model by incorporating batch normalization, anchor boxes, and dimension clusters.
- [YOLOv3](#), launched in 2018, further enhanced the model's performance using a more efficient backbone network, multiple anchors and spatial pyramid pooling.
- [YOLOv4](#) was released in 2020, introducing innovations like Mosaic [data augmentation](#), a new anchor-free detection head, and a new [loss function](#).
- [YOLOv5](#) further improved the model's performance and added new features such as hyperparameter optimization, integrated experiment tracking and automatic export to popular export formats.
- [YOLOv6](#) was open-sourced by [Meituan](#) in 2022 and is in use in many of the company's autonomous delivery robots.
- [YOLOv7](#) added additional tasks such as pose estimation on the COCO keypoints dataset.
- [YOLOv8](#) released in 2023 by Ultralytics. YOLOv8 introduced new features and improvements for enhanced performance, flexibility, and efficiency, supporting a full range of vision AI tasks,
- [YOLOv9](#) introduces innovative methods like Programmable Gradient Information (PGI) and the Generalized Efficient Layer Aggregation Network (GELAN).
- [YOLOv10](#) is created by researchers from [Tsinghua University](#) using the [Ultralytics Python package](#). This version provides real-time [object detection](#) advancements by introducing an End-to-End head that eliminates Non-Maximum Suppression (NMS) requirements.
- **YOLO11**  **NEW:** Ultralytics' latest YOLO models delivering state-of-the-art (SOTA) performance across multiple tasks, including [detection](#), [segmentation](#), [pose estimation](#), [tracking](#), and [classification](#), leverage capabilities across diverse AI applications and domains.

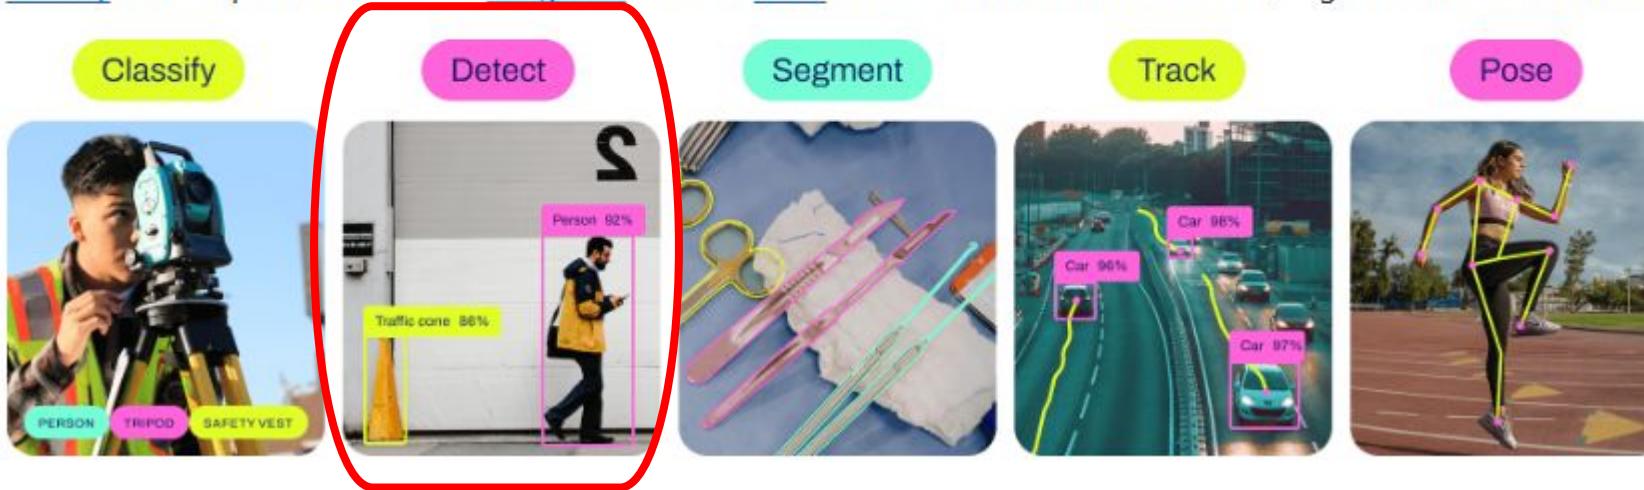
2015

2025



YOLO(You Only Look Once) - Models

YOLO11 [Detect](#), [Segment](#) and [Pose](#) models pretrained on the [COCO](#) dataset are available here, as well as YOLO11 [Classify](#) models pretrained on the [ImageNet](#) dataset. [Track](#) mode is available for all Detect, Segment and Pose models.



- Open-source (python codebase) and easy to use
- Extensive documentation , tutorials, with a large and active community(>60 million)
- Shared encoders between models (part of the CNN that learns relevant features). The encoder of one model can be used as starting point for others

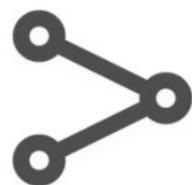


NOAA
FISHERIES

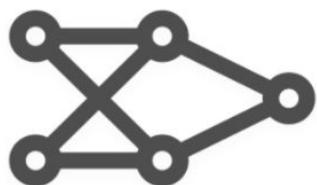
Base Foundation Model Selection

Model	size (pixels)	mAP ^{val} 50-95	Speed CPU ONNX (ms)	Speed T4 TensorRT10 (ms)	params (M)	FLOPs (B)
YOLO11n	640	39.5	56.1 ± 0.8	1.5 ± 0.0	2.6	6.5
YOLO11s	640	47.0	90.0 ± 1.2	2.5 ± 0.0	9.4	21.5
YOLO11m	640	51.5	183.2 ± 2.0	4.7 ± 0.1	20.1	68.0
YOLO11l	640	53.4	238.6 ± 1.4	6.2 ± 0.1	25.3	86.9
YOLO11x	640	54.7	462.8 ± 6.7	11.3 ± 0.2	56.9	194.9

- speed vs accuracy tradeoffs
- training time
- Resources for Deployment



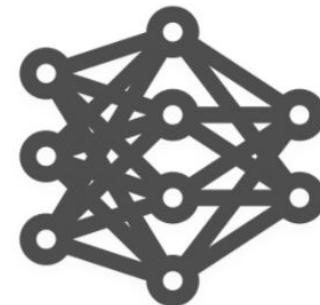
Small



Medium



Large



XLarge

YOLO Python Script to Train Model

```
# Load Library
import YOLO

# Load a foundational model
model = YOLO("yolo11n.pt")

# Train Model
model.train(data="fish_dataset.yaml", # path to dataset
            epochs=100, # number of training epochs
            imgsz=640, # training image size
            device=cpu, # device to run on - cpu, gpu(0,1)
            ...)
```



note: When training a model, an epoch refers to one complete pass through the entire training dataset.

What is the right number of Training Passes (Epochs) to Use?

It depends. **Varies based on model, dataset size and task.**

Larger datasets require more while smaller datasets will need less.

Good baselines:

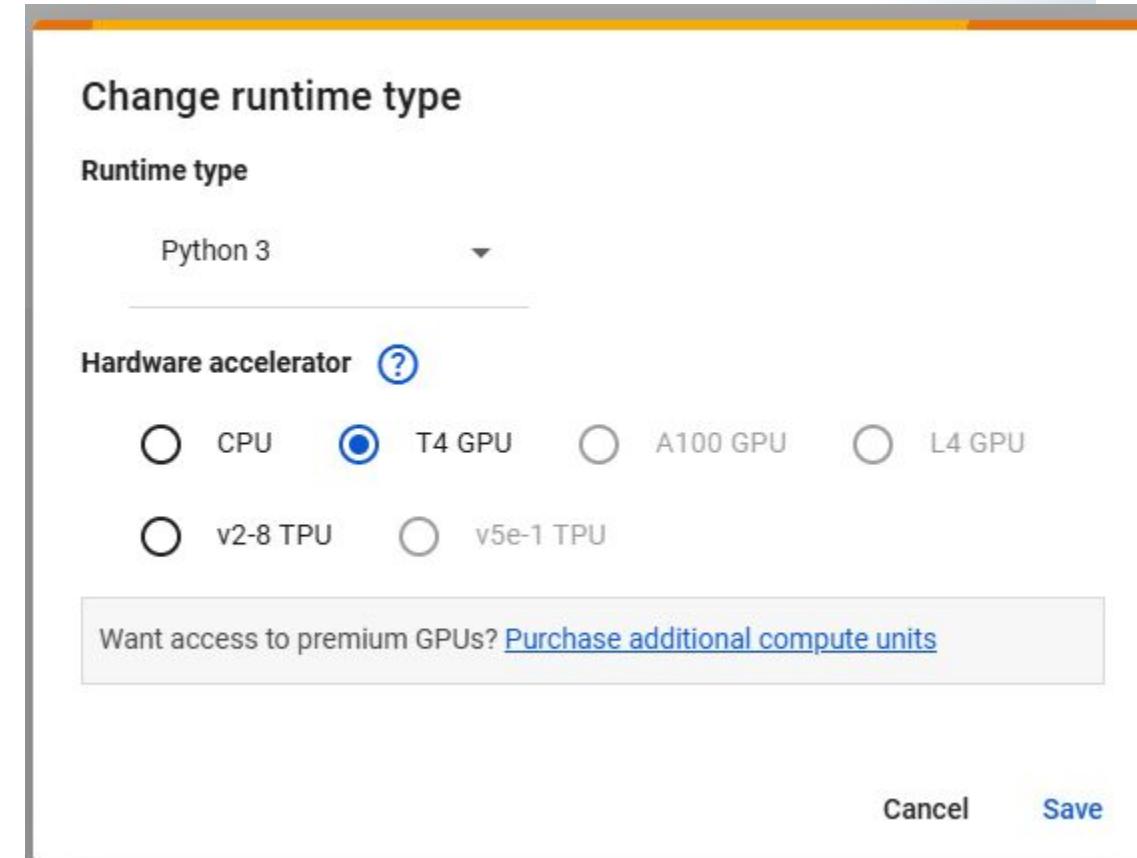
- 1-10 - test setup
- 50 - mockup model
- 100 min regular run
- 300 recommended starting point

Do I need a Supercomputer to get started?

Nope. Just a Web browser to Get Started.

Google Colaboratory

Colab is a hosted Jupyter Notebook service that requires no setup to use and provides free access to computing resources, including GPUs and TPUs. Colab is especially well suited to machine learning, data science, and education.



NOAA
FISHERIES



+ Code



```
[ ] import os
[ ] import torch
from ultralytics import YOLO

device = 'cuda' if torch.cuda.is_available() else 'cpu'
print(f"Using device: {device}")

# Load the large model (YOLOv8x) that was fine-tuned for VME detection
large_model = YOLO("/content/best.pt") # Replace with your model path
large_model = large_model.to(device) # Move the model to GPU (CUDA) if available

# Step 4: Create folders for the new dataset (images and labels)
base_path = "/content/fish_dataset"

# Step 11: Train the YOLOv11n model using the generated dataset
small_model = YOLO("yolo11n.pt") # Load the smaller YOLOv8n model

# Train the model using the generated fish-only dataset
small_model.train(data="/content/fish_dataset.yaml", epochs=50, imgsz=416, batch=16, lr0=0.001)

print("Training complete!")

# model save
small_model.save("/content/yolov11n_fish_trained.pt")

# model val stats
metrics = small_model.val(data="/content/fish_dataset.yaml") # This will evaluate precision, recall, and mAP
```



Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/50	1.45G	1.276	2.073	1.11	5	416: 100% ██████████ 93/93 [00:28<00:00, 3.25it/s]
	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 12/12 [00:03<00:00, 3.06it/s]
	all	369	473	0.466	0.431	0.426 0.255

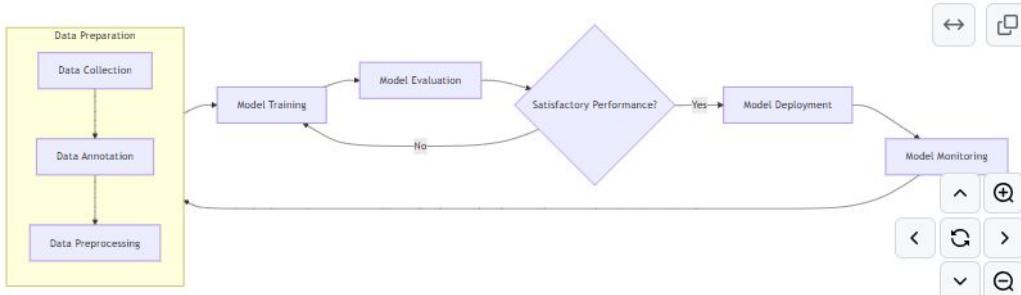
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
2/50	1.36G	1.194	1.457	1.083	11	416: 100% ██████████ 93/93 [00:26<00:00, 3.55it/s]
	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 12/12 [00:02<00:00, 5.04it/s]
	all	369	473	0.466	0.431	0.426 0.255

- <https://github.com/MichaelAkridge-NOAA/open-science-ai-toolkit>

Open Science AI Toolkit

Open source suite of tools, workflows, and processes designed to accelerate Open Science AI/ML development.

Overview



Installation

To install the entire toolkit:

```
pip install git+https://github.com/MichaelAkridge-NOAA/open-science-ai-toolkit.git
```

Features

The `noaa_ai_tools` subpackage provides functionalities like:

- **Dataset Preparation:** Easily prepare your datasets for training with functions for filtering and validating annotations.
 - Filter Images with Labels
 - Create Labels for Background Images
 - Validate YOLO Format
 - Remap Class IDs
 - Split Dataset (train, val, test)
- **Model Training:** Train and manage deep learning models, including support for YOLO models from Ultralytics(<https://github.com/ultralytics/ultralytics>).
- **Model Exporting:** Export trained models to various formats, including ONNX, TensorFlow Lite, Edge TPU, and NCNN.

```
# Training hyperparameters
small_model.train(
    data=yaml_file_path,
    epochs=100,
    imgsz=640,
    batch=32, # Adjust batch size based on GPU capacity
    lr0=0.001, # Initial learning rate
    lrf=0.0001, # Final learning rate (used for Cosine Annealing)
    optimizer='AdamW', # Use AdamW optimizer for better performance
    device=device,
    save_period=10, # Save model checkpoint every 10 epochs
    patience=10, # Early stopping if no improvement after 10 epochs
    augment=True, # Enable data augmentation
    mosaic=True, # Use mosaic augmentation
    mixup=True, # Use MixUp augmentation
    cos_lr=True, # Cosine annealing Learning rate
    project='training_logs', # TensorBoard Logging directory
)

print("Training complete!")

# Unfreeze all layers after the initial phase
for param in small_model.model.model.parameters():
    param.requires_grad = True

# Save the trained model
trained_model_path = os.path.join(base_path, "yolo1m_urchin_trainedv3.pt")
small_model.save(trained_model_path)
print(f"Trained model saved to {trained_model_path}")

# Save the model weights separately for further use
weights_path = os.path.join(base_path, "yolo1m_urchin_weightsv3.pth")
torch.save(small_model.model.state_dict(), weights_path)
print(f"Weights saved to {weights_path}")

# Evaluate model performance
```

Lessons Learned

Lessons Learned

Garbage in, garbage out

"the concept that flawed, biased or poor quality ("garbage") information or input produces a result or output of similar ("garbage") quality."

- still very relevant today
- The better your training dataset the better your model will be

$$\text{PRECISE NUMBER} + \text{PRECISE NUMBER} = \text{SLIGHTLY LESS PRECISE NUMBER}$$

$$\text{PRECISE NUMBER} \times \text{PRECISE NUMBER} = \text{SLIGHTLY LESS PRECISE NUMBER}$$

$$\text{PRECISE NUMBER} + \text{GARBAGE} = \text{GARBAGE}$$

$$\text{PRECISE NUMBER} \times \text{GARBAGE} = \text{GARBAGE}$$

$$\sqrt{\text{GARBAGE}} = \text{LESS BAD GARBAGE}$$

$$(\text{GARBAGE})^2 = \text{WORSE GARBAGE}$$

$$\frac{1}{N} \sum (\text{N PIECES OF STATISTICALLY INDEPENDENT GARBAGE}) = \text{BETTER GARBAGE}$$

$$(\text{PRECISE NUMBER})^{\text{GARBAGE}} = \text{MUCH WORSE GARBAGE}$$

$$\text{GARBAGE} - \text{GARBAGE} = \text{MUCH WORSE GARBAGE}$$

$$\frac{\text{PRECISE NUMBER}}{\text{GARBAGE} - \text{GARBAGE}} = \text{MUCH WORSE GARBAGE, POSSIBLE DIVISION BY ZERO}$$

$$\text{GARBAGE} \times 0 = \text{PRECISE NUMBER}$$

Lessons Learned: How to Create a better model

More data \neq Better Model

Better data = Better Model

Lessons Learned

Storage & Access of AI Ready Datasets is Lacking

- In general, AI ready datasets are very sparse and/or hard to access easily
- When making a dataset:
 - Choose a standard format like COCO (json)
 - Although, its very easy to convert from/to any format

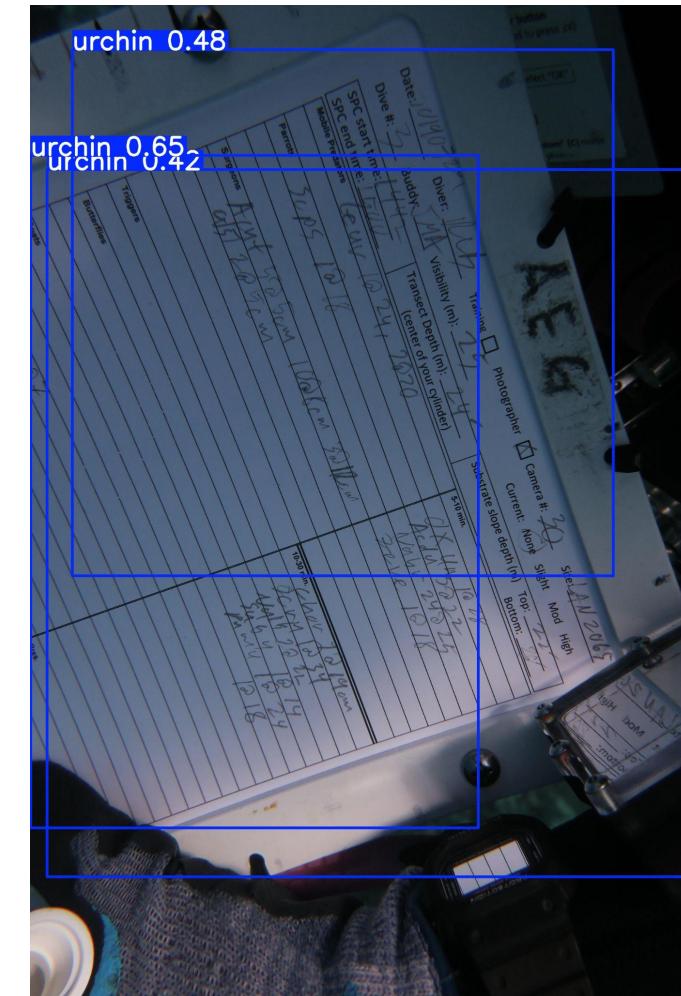
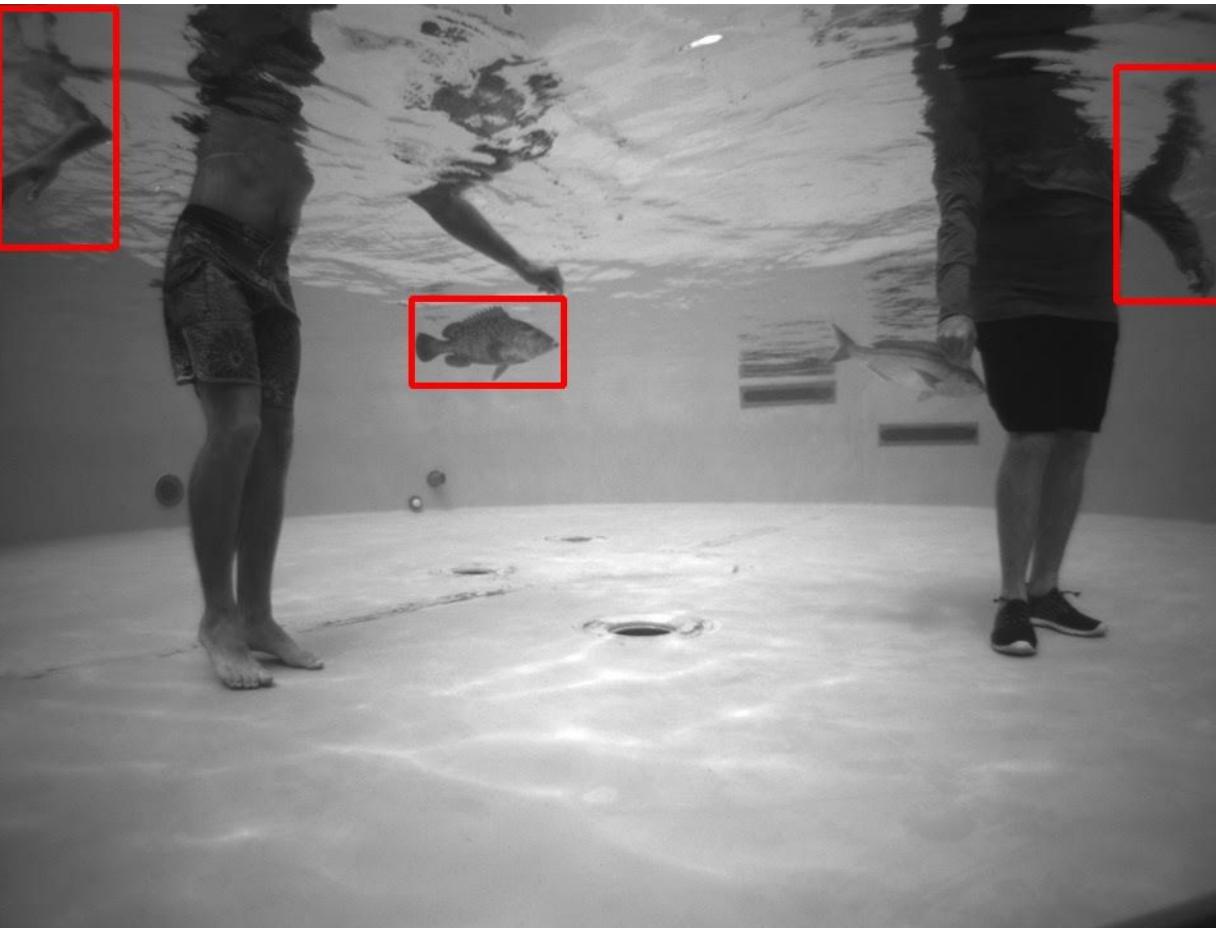
The screenshot shows the Google Cloud Storage interface. The top navigation bar includes 'Google Cloud', 'Select a project', and a search bar. The main title is 'Bucket details' for 'nmfs_odp_pifsc'. The left sidebar has tabs for 'Overview', 'Buckets' (which is selected), 'Monitoring', and 'Settings'. The main content area shows the bucket 'nmfs_odp_pifsc' with details: Location (us (multiple regions in United States)), Storage class (Standard), and Public access (Access granted to public principals). Below this are tabs for 'OBJECTS', 'CONFIGURATION', 'PERMISSIONS', 'PROTECTION', and 'LIFECYCLE'. A breadcrumb trail shows the path: Buckets > nmfs_odp_pifsc > PIFSC > ESD > ARP > pifsc-ai-data-repository. There are buttons for 'CREATE FOLDER', 'UPLOAD', 'TRANSFER DATA', and 'OTHER SERVICES'. A filter bar allows filtering by name prefix only or objects and folders. The 'OBJECTS' table lists five entries, all of which are empty folders:

Name	Size	Type	Created	Storage
coral-classification/	—	Folder	—	—
fish-detection/	—	Folder	—	—
image-segmentation/	—	Folder	—	—
object-tracking/	—	Folder	—	—

Cloud AI data repo in the works

Lessons Learned

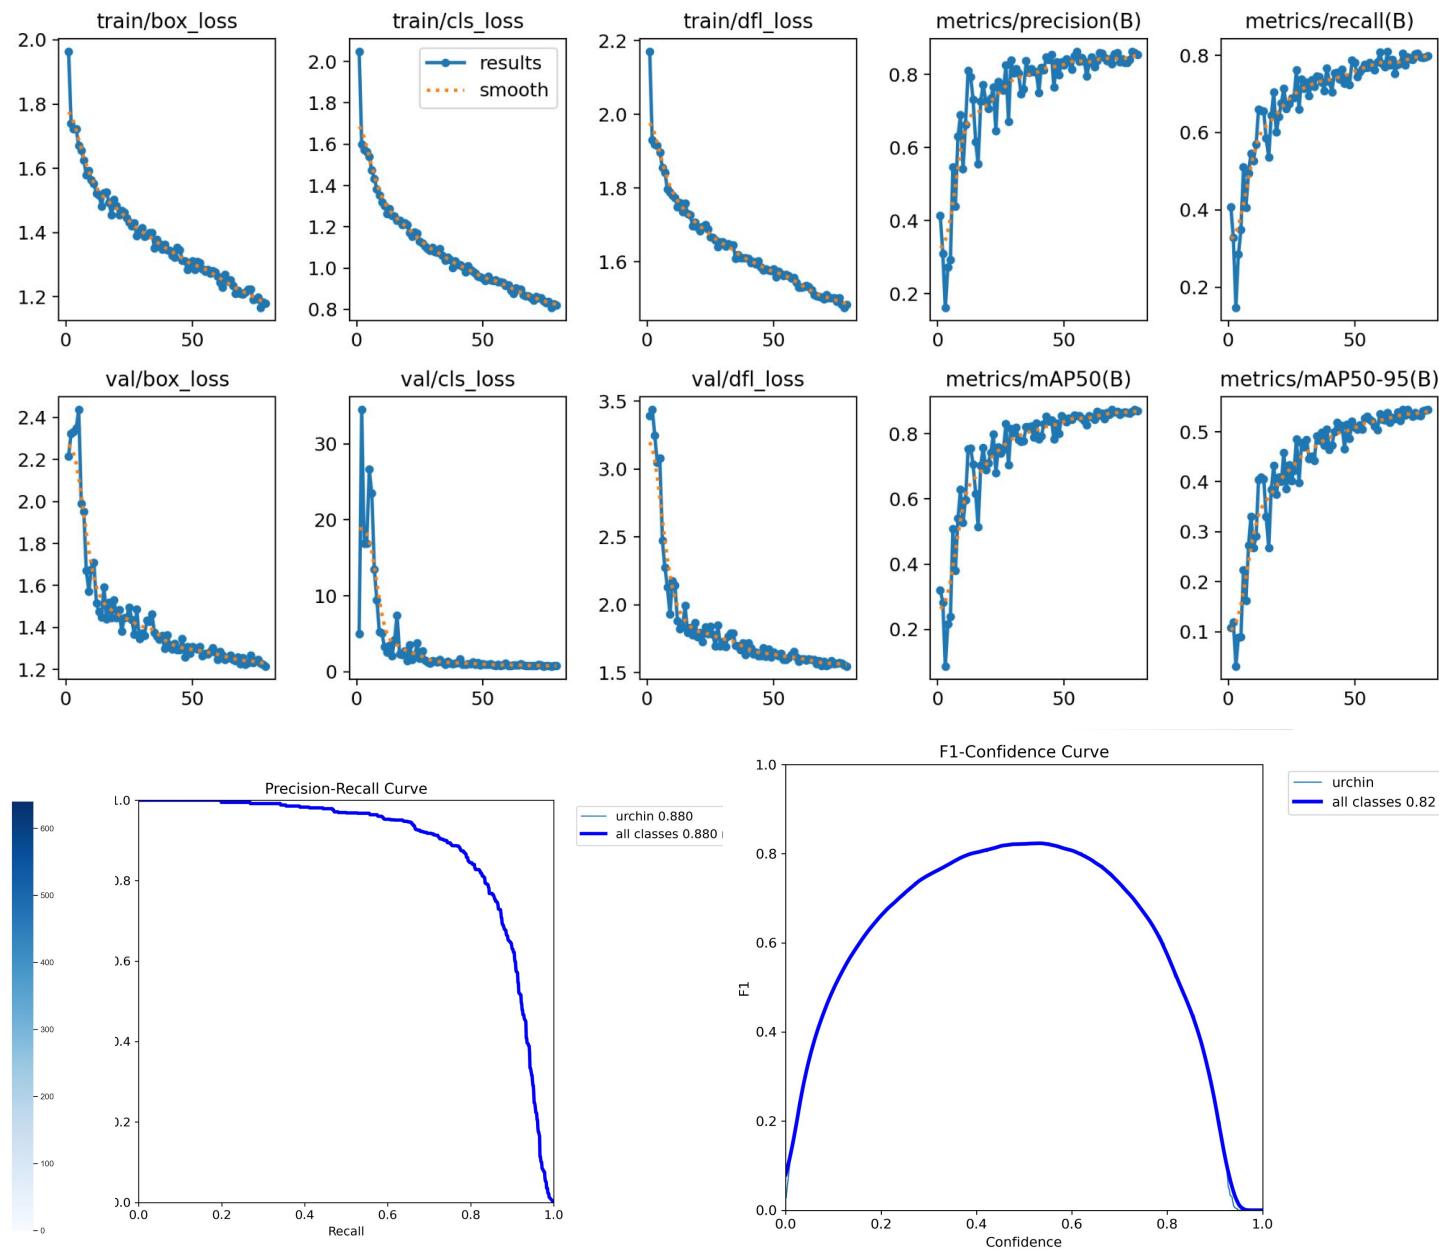
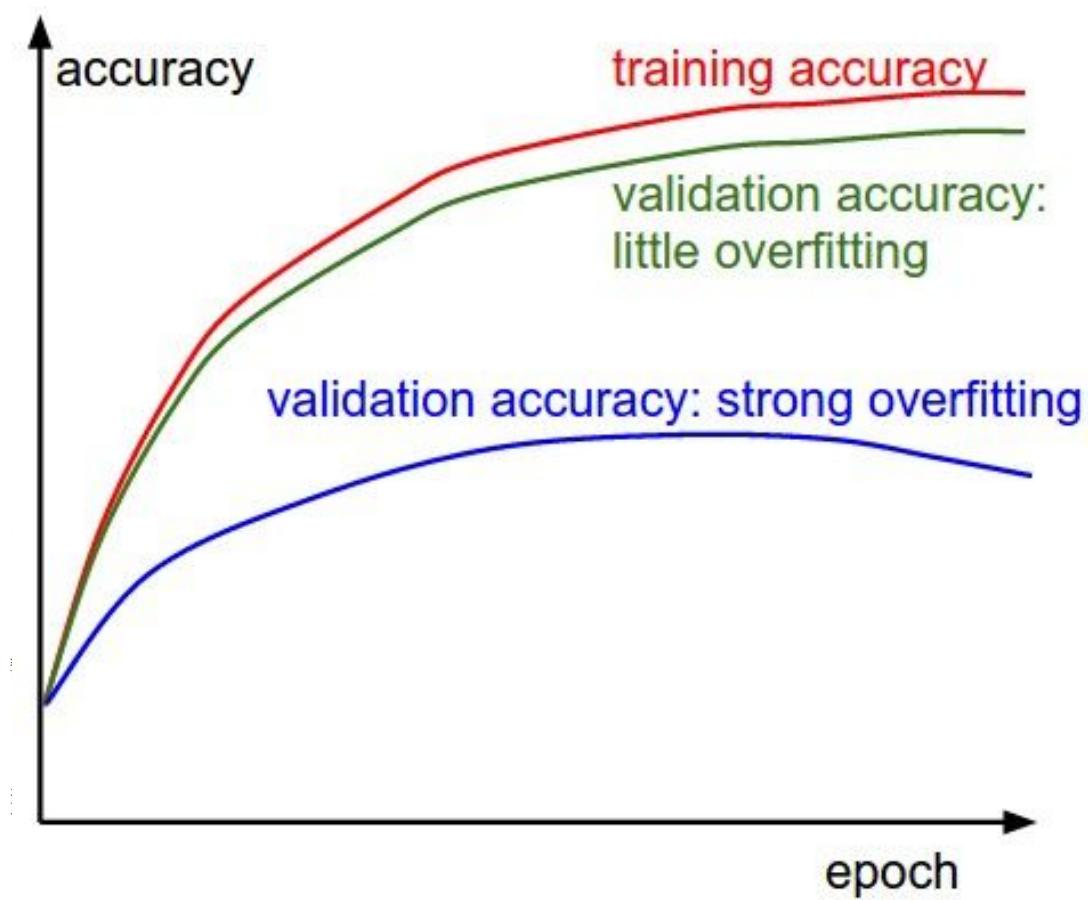
- When making a Training Dataset:
 - The model won't know what it hasn't seen



NOAA
FISHERIES

Lessons Learned

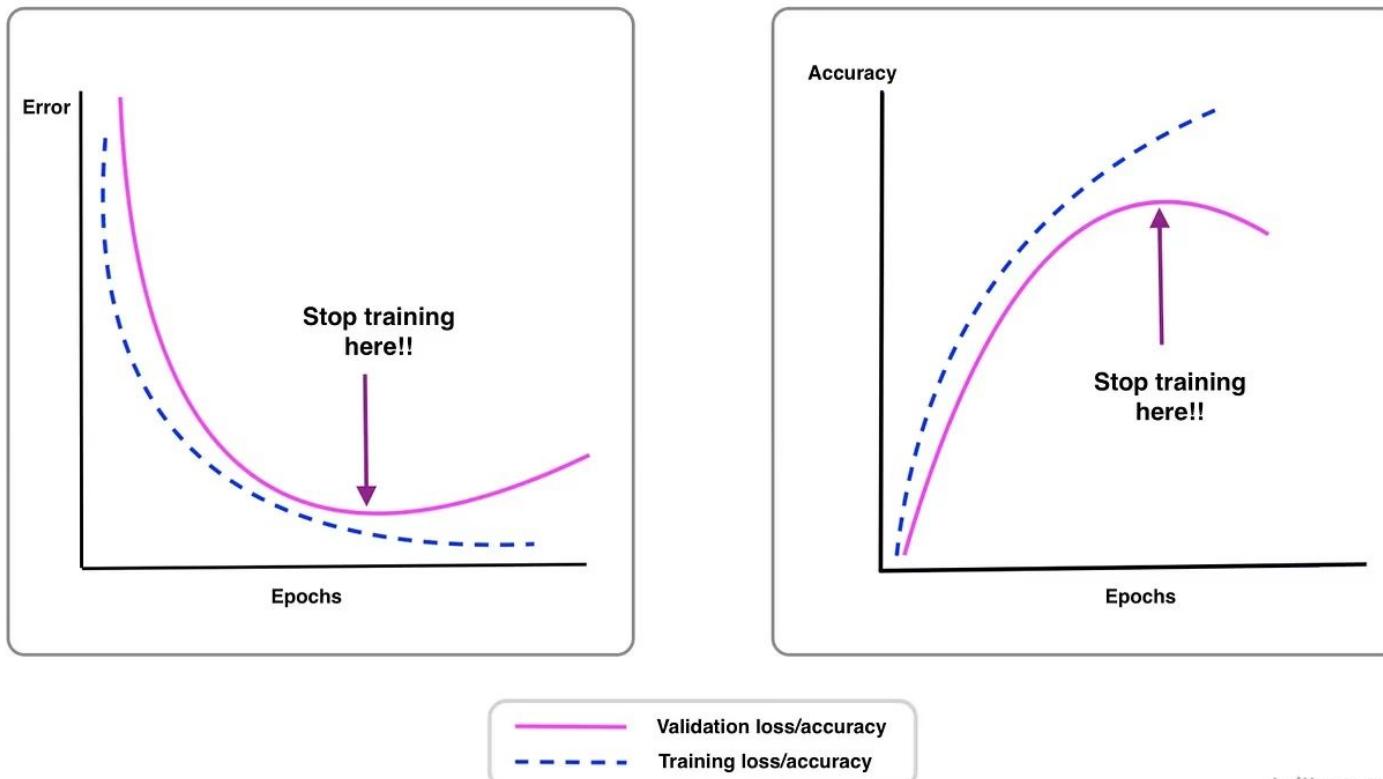
- There are built in Evaluation Tools for Metrics



Lessons Learned

- More training is not always better.
- Stopping early can help prevent “Overfitting” (memorizing of the training set) and help optimize the model.

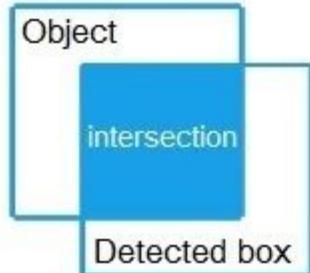
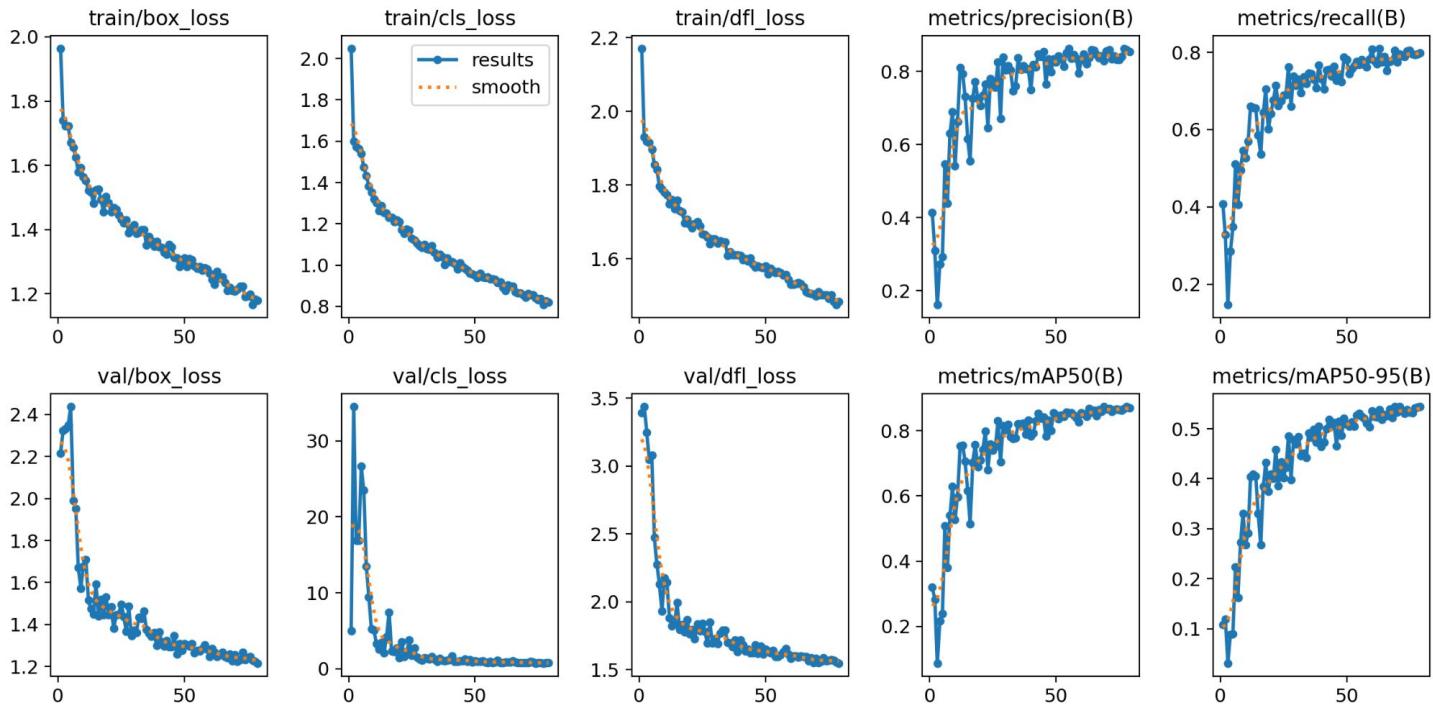
Early Stopping



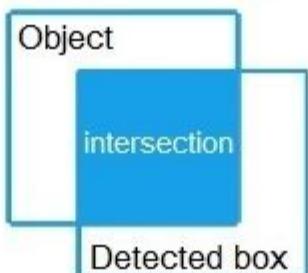
twitter.com/jeande_d

Lessons Learned

- Evaluation Tools for Metrics



$$\text{Precision} = \frac{\text{intersection}}{\text{Detected box}}$$

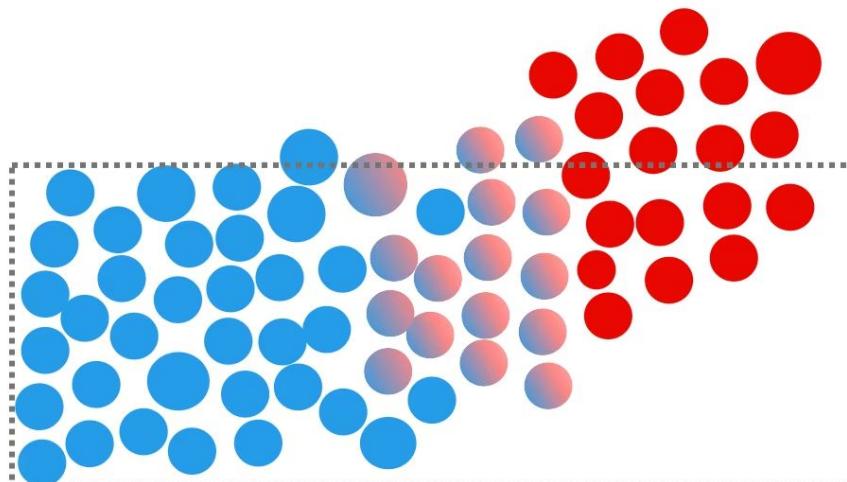


$$\text{Recall} = \frac{\text{intersection}}{\text{Object}}$$

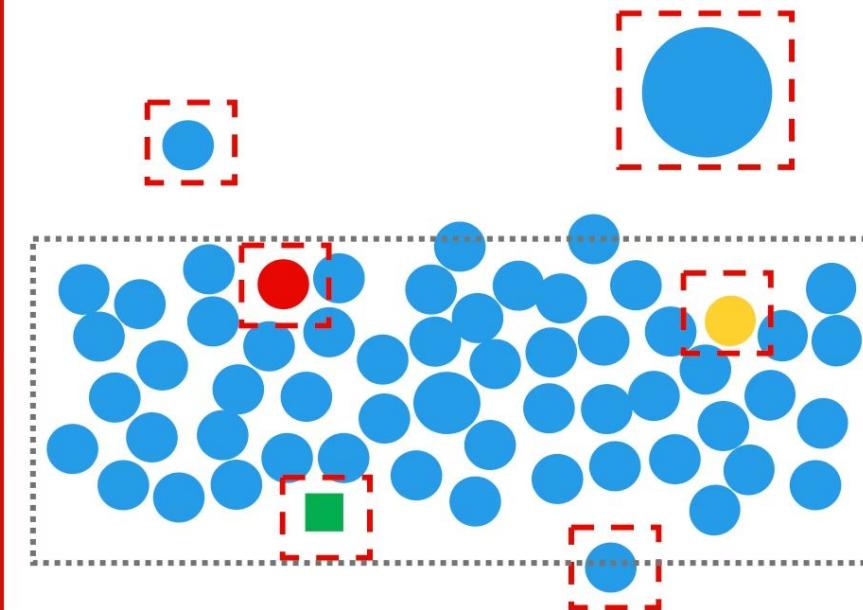
$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$

urchin
all classes 0.82 at

Data drift



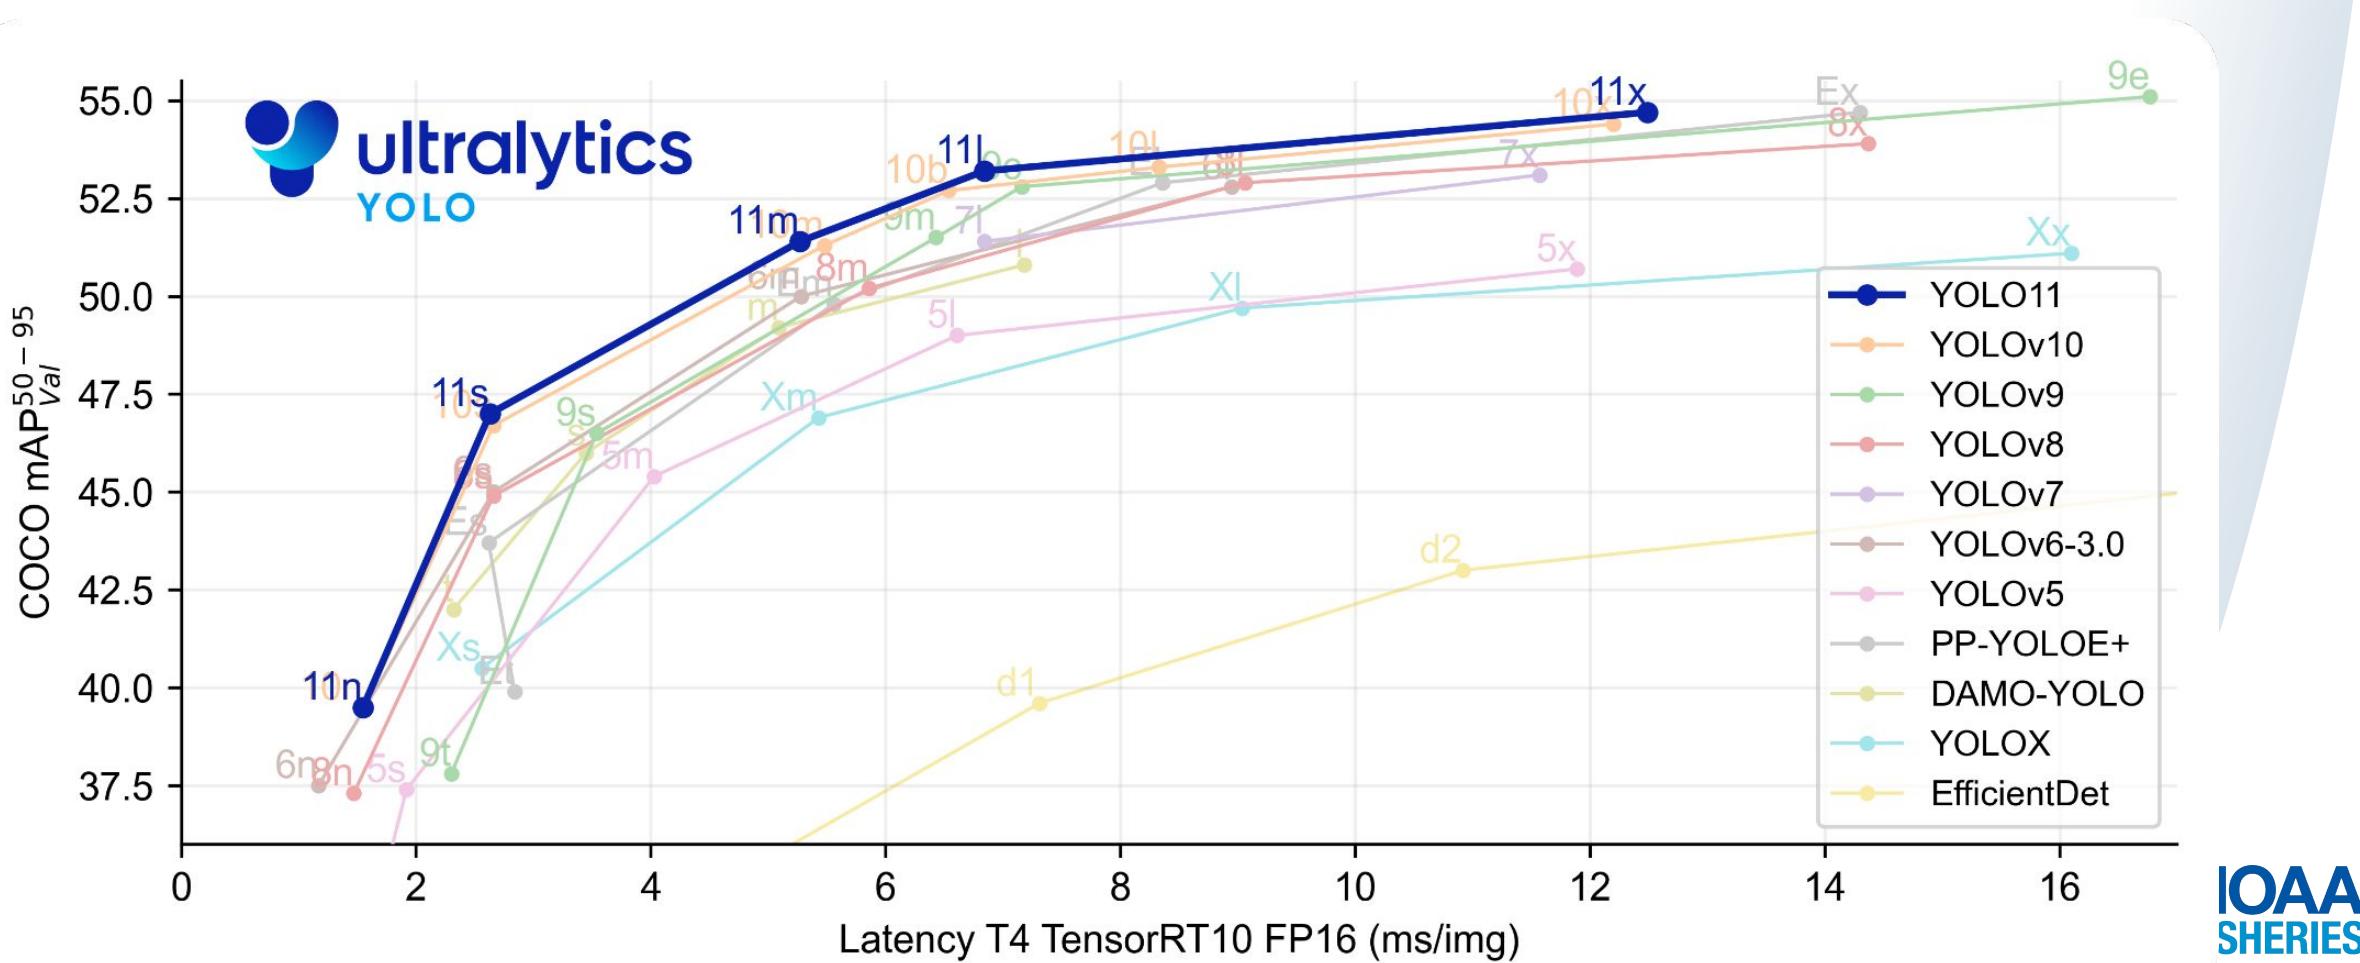
Outliers



NOAA
FISHERIES

Lessons Learned

- Changes, improvements, and updates with AI happen fast



The Fun Stuff: Test Models, Advanced Techniques & Examples



Explore our ocean

FathomNet is an open-source image database for understanding our ocean and its inhabitants



Marine Detect

This repository provides access to two YOLOv8 object detection models for identifying species of interest in underwater environments.

Datasets Details

The models utilize a combination of publicly available datasets (~ 90%) and Ténaka-based datasets (~ 10%). Some datasets were already annotated, and others undergo manual labeling.

References to the public datasets used can be found in the 'References' section of this README.

The images used with annotations (YOLO format) can be downloaded using the following links: [FishInv dataset](#), [MegaFauna dataset](#).

Datasets split details

Model	Training + Validation Sets	Test Set
FishInv	12,243 images (80%, 20%)	499 images
MegaFauna	8,130 images (80%, 20%)	253 images



Home

Datasets

Projects

Models

roboflow

Universe [Public Datasets](#) Model Zoo Blog Docs

Home >
Datasets

Dataset Type

All Datasets 40

Object Detection 36

Classification 4

The Ultra

Computer Vision Datasets

Roboflow hosts free public computer vision datasets in many po
available.

If you'd like us to host your dataset, please [get in touch](#).

Roboflow 100

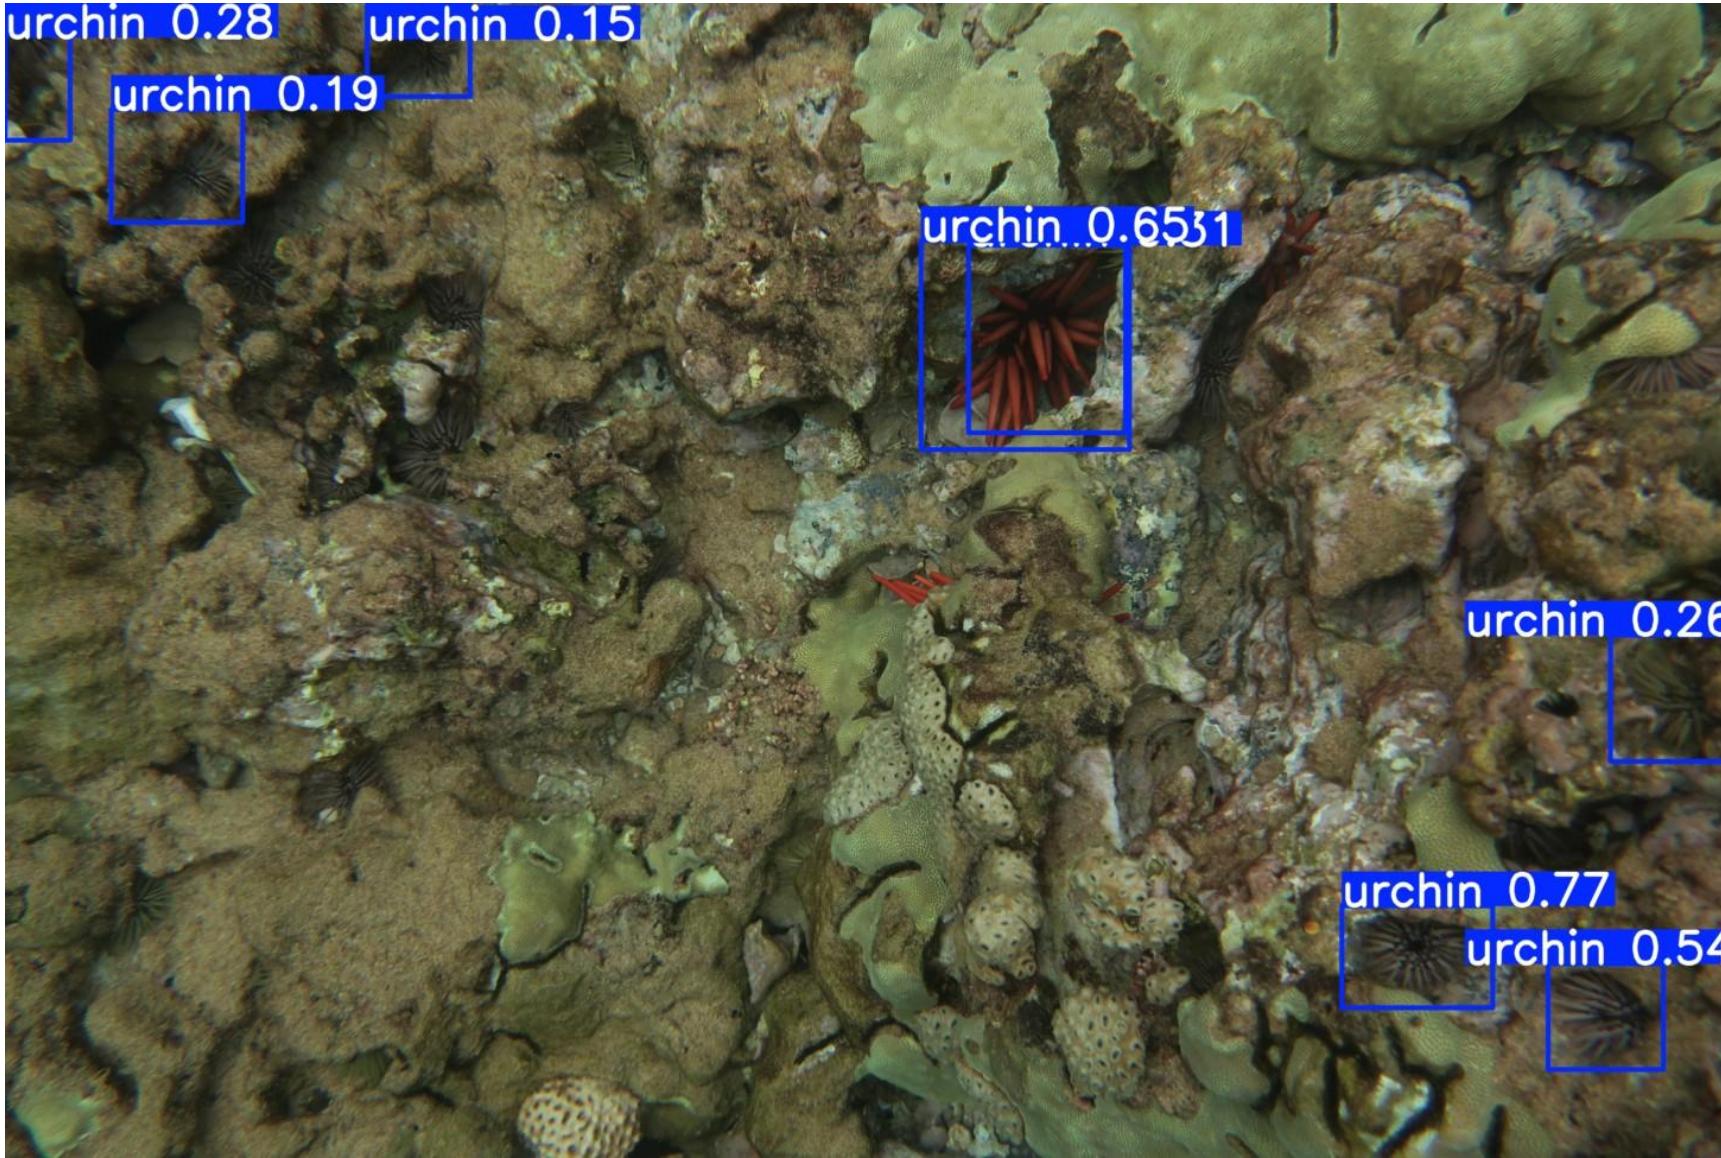
Object Detection Benchmark
100 datasets | 7 Domains | 224k Images

[Microsoft COCO 2017 Dataset](#)

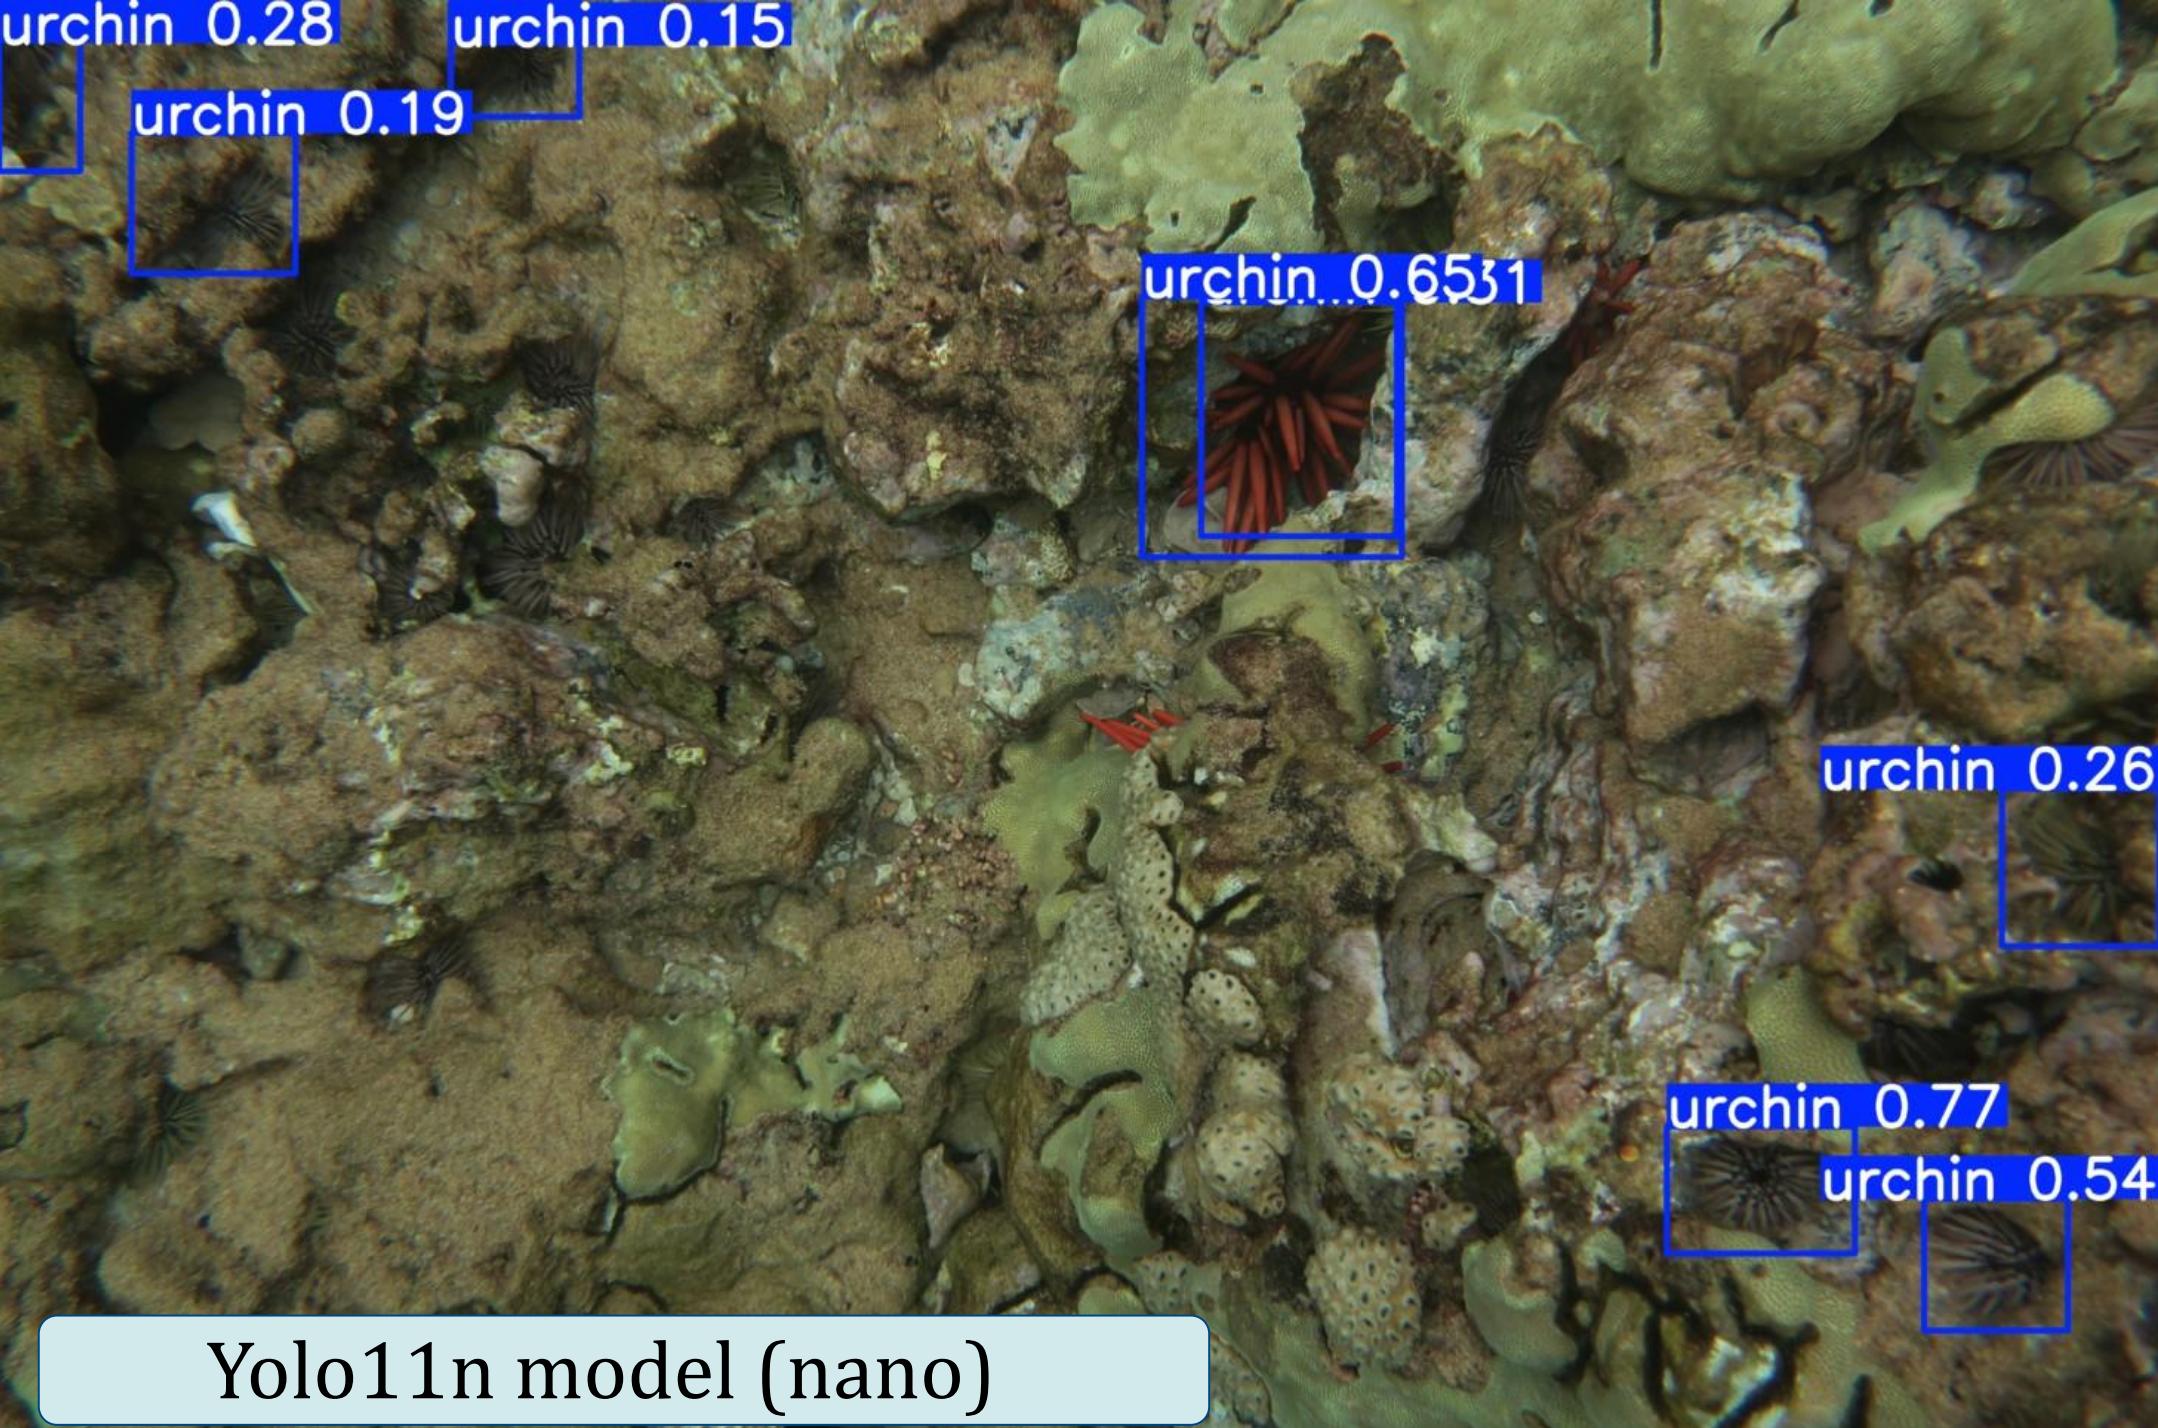
Object Detection (Bounding Box)
120358 images | 12 exports | Last updated a month ago

Object Detection Tests: Urchins

Yolo11n model (nano)



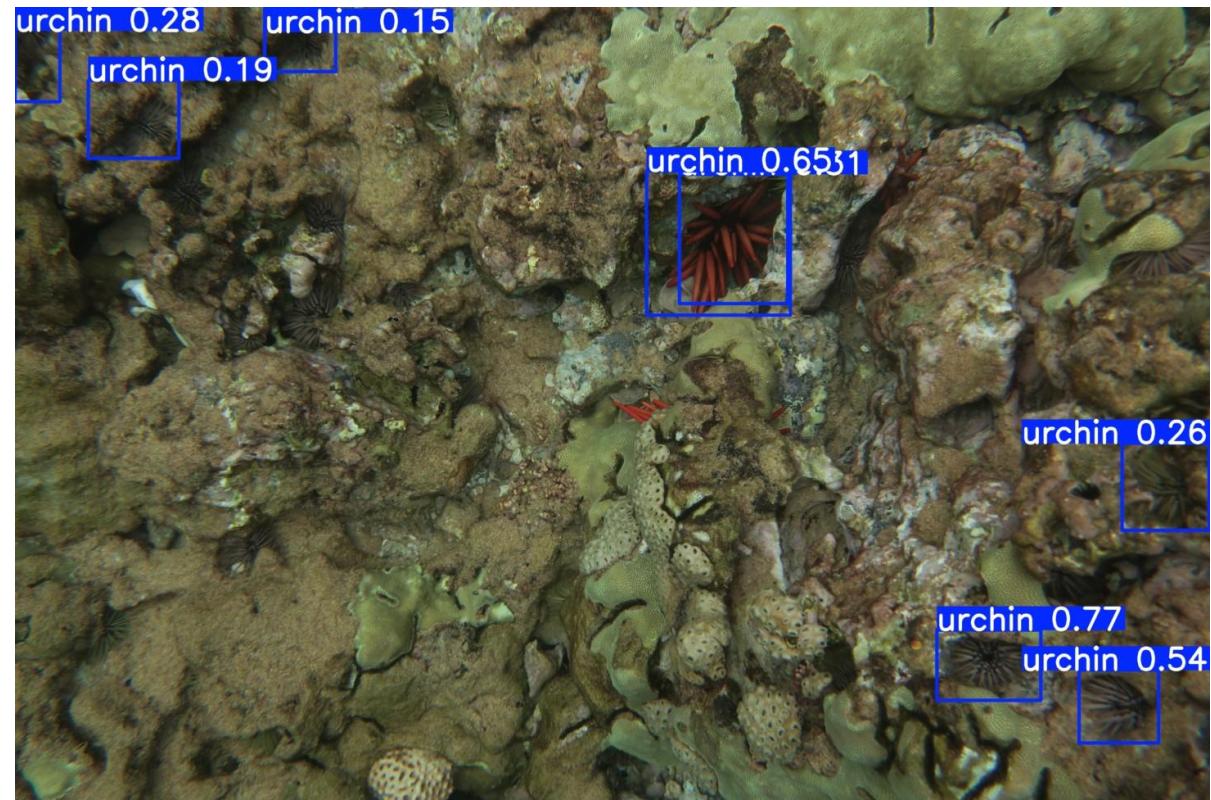
NOAA
FISHERIES



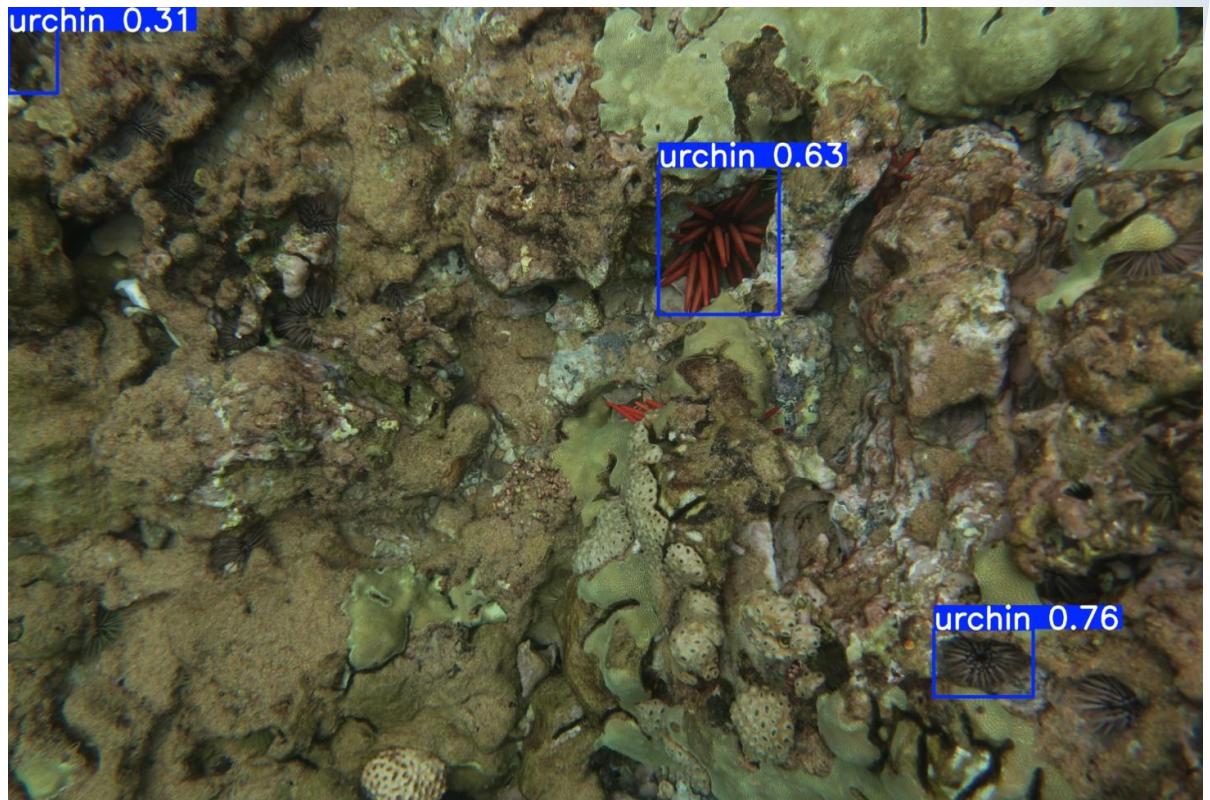
Yolo11n model (nano)

Lessons Learned

Yolo11n at 50 epochs



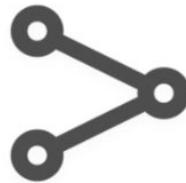
Yolo11n at 100 epochs



note: When training a model, an epoch refers to one complete pass through the entire training dataset.

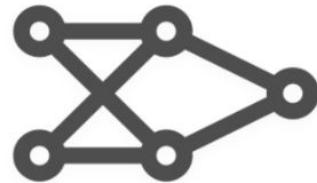
Lessons Learned

- Choose a model that makes sense for the task
 - nano for on edge and speed, and simple/large objects
 - XL for complexity
 - when unsure, try medium



Small

~45mins

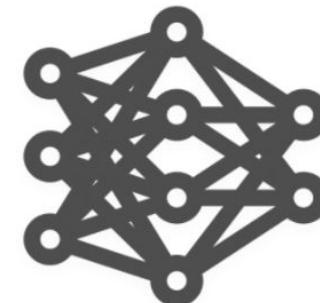


Medium

~4 hours



Large

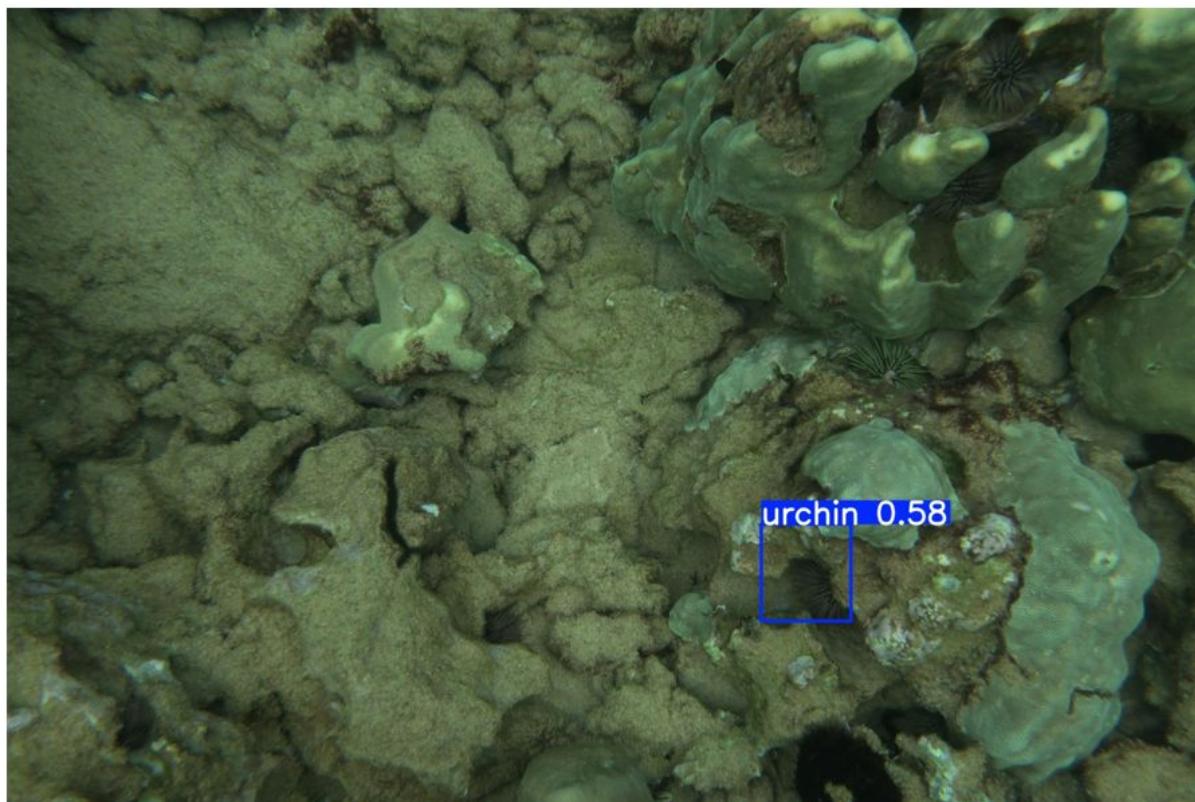


XLarge

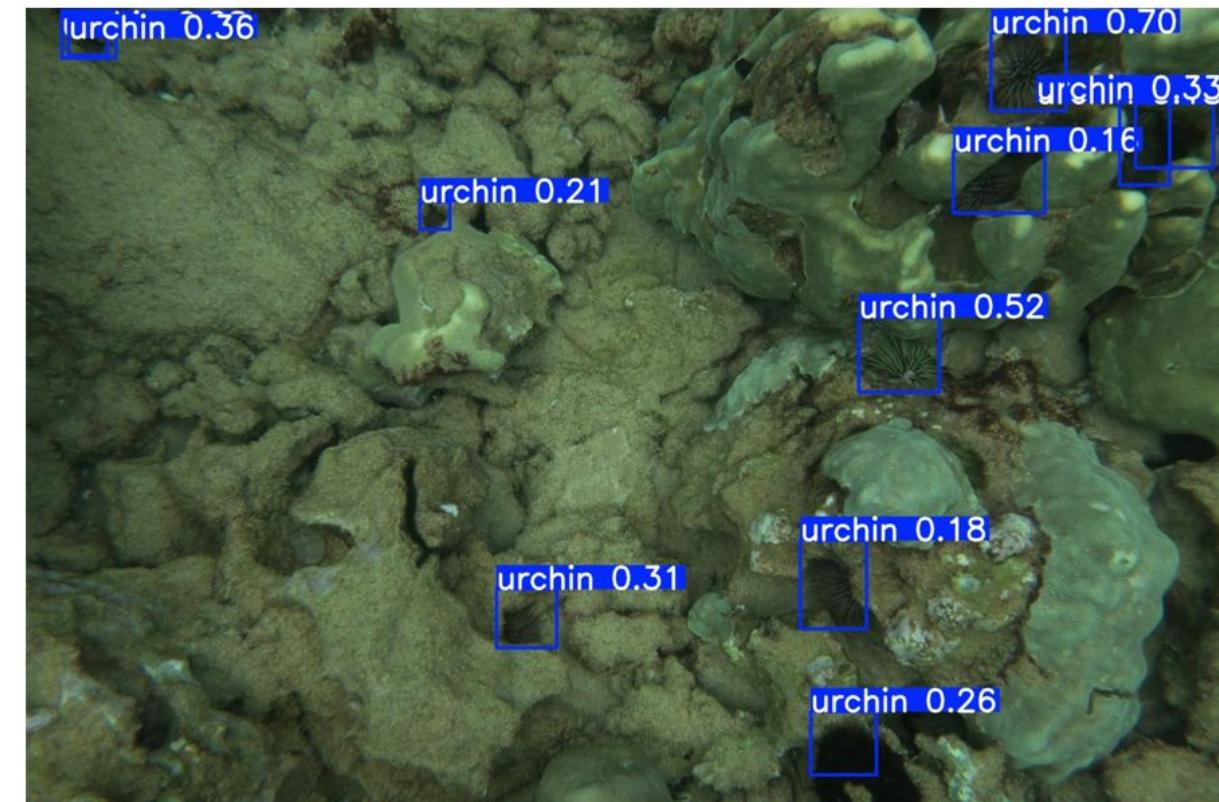
~8 hours

Yolo11 Nano Model vs Yolo11 Medium Model Performance

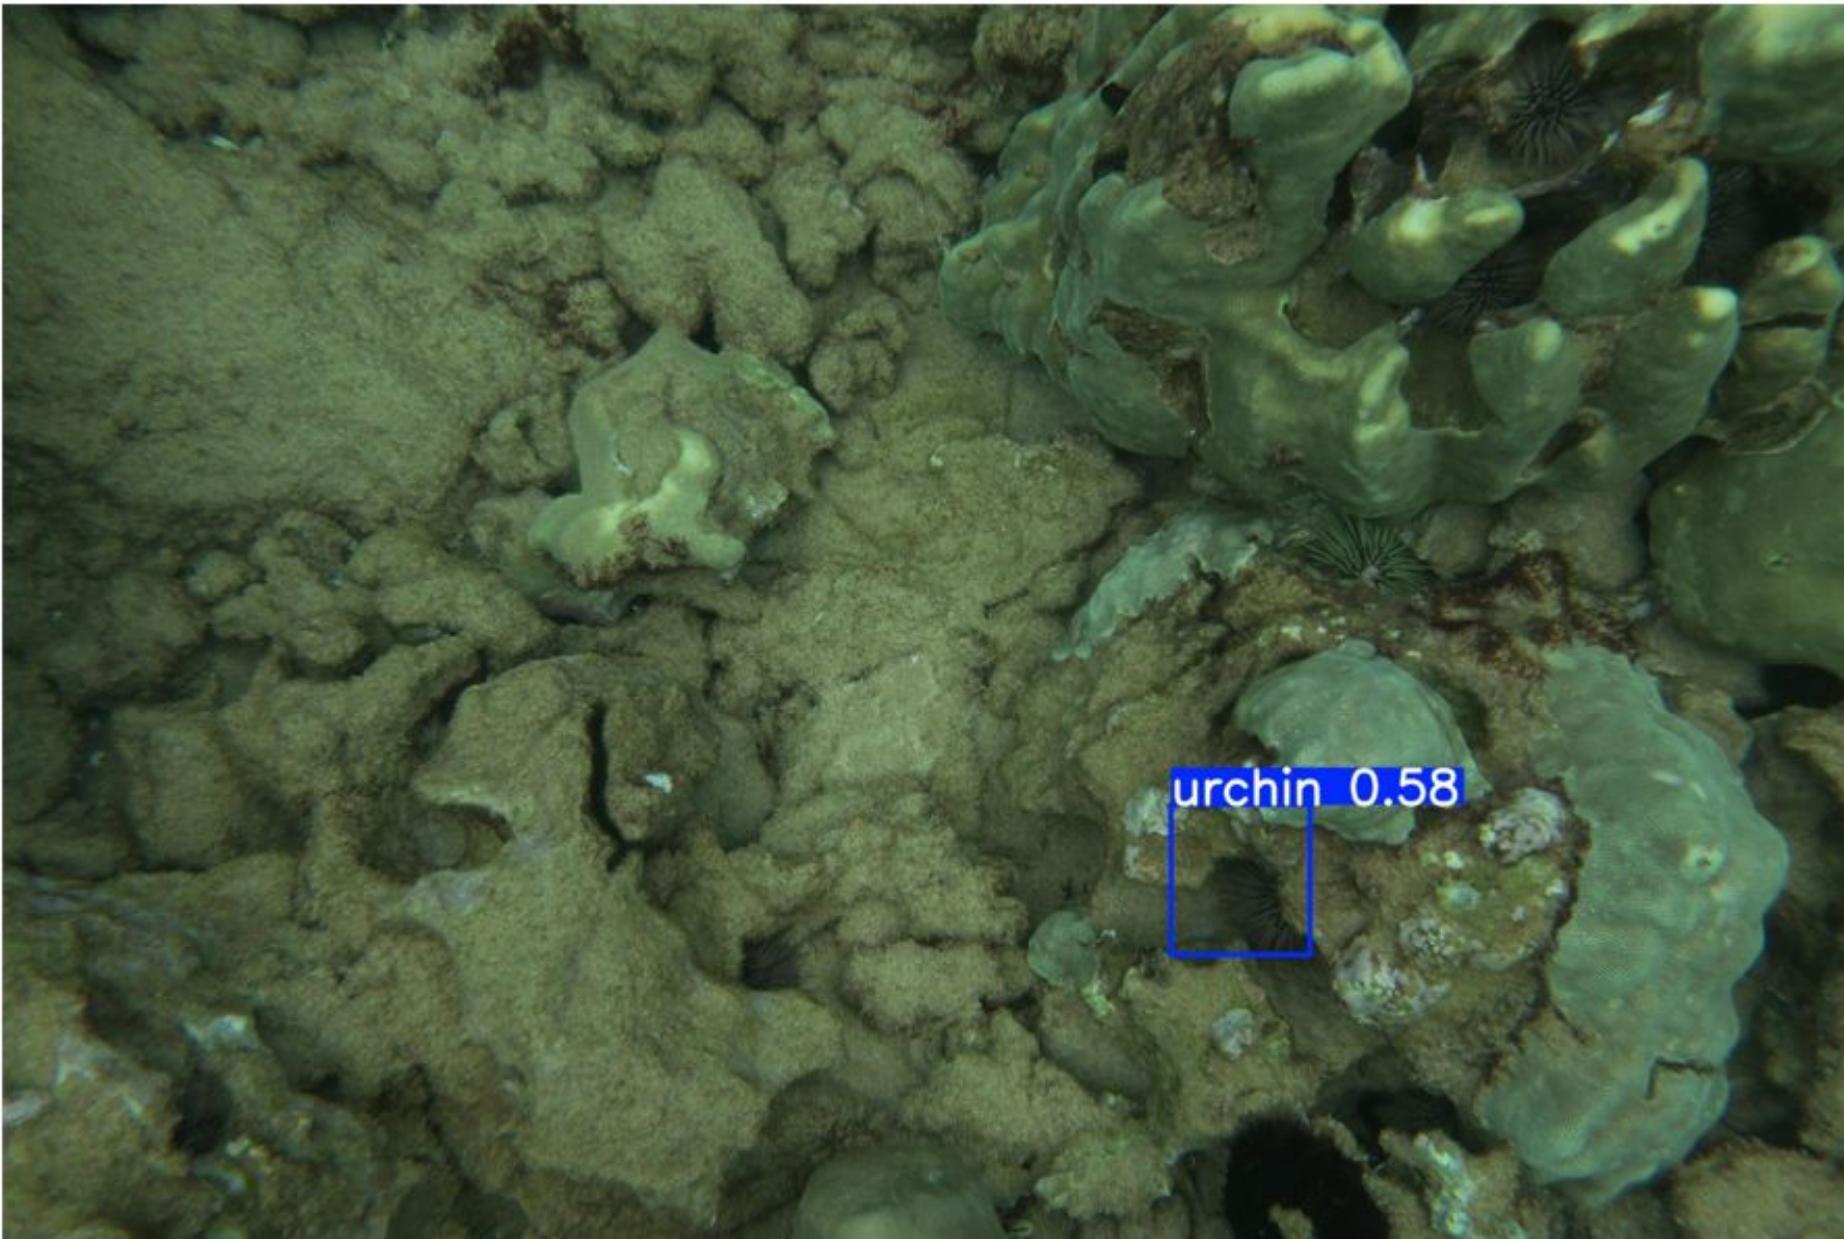
Urchin Detector yolo11n(nano)



Urchin Detector yolo11m(medium)

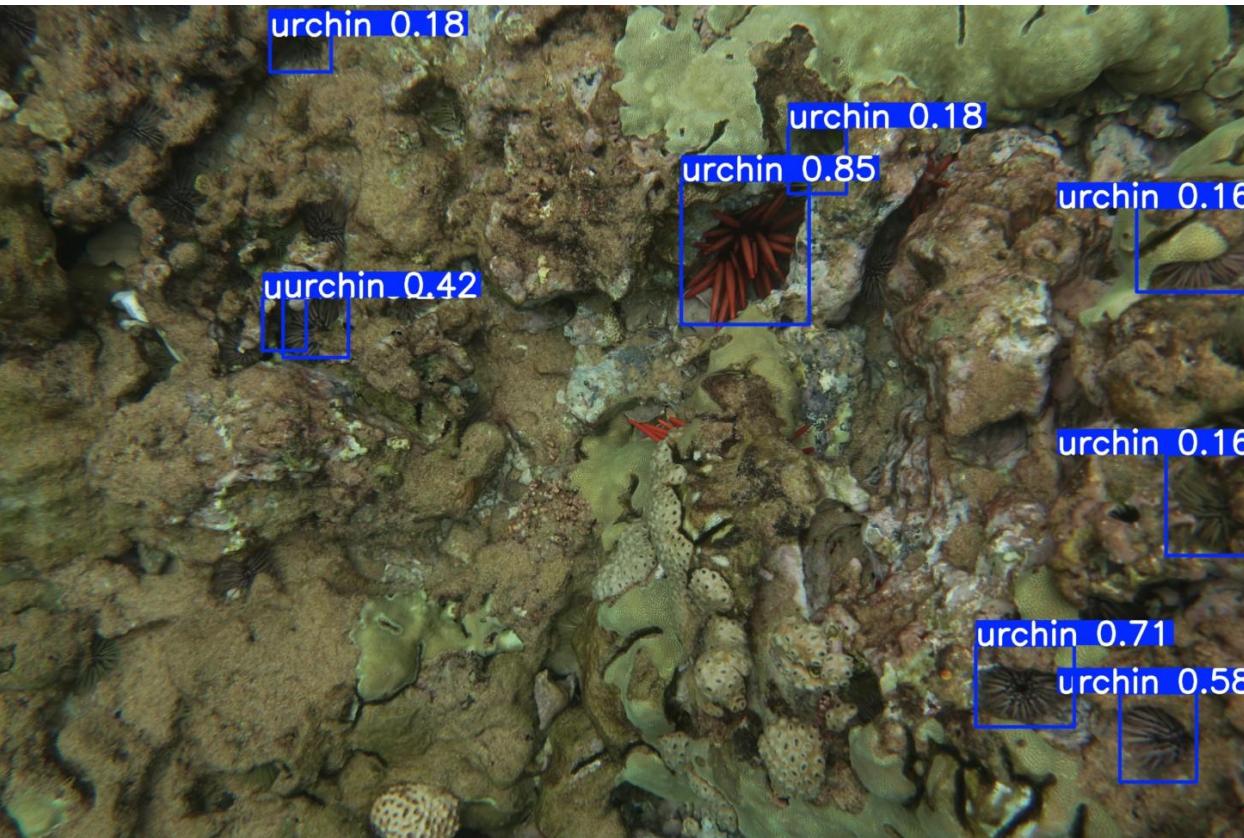


Urchin Detector yolo11n(nano)

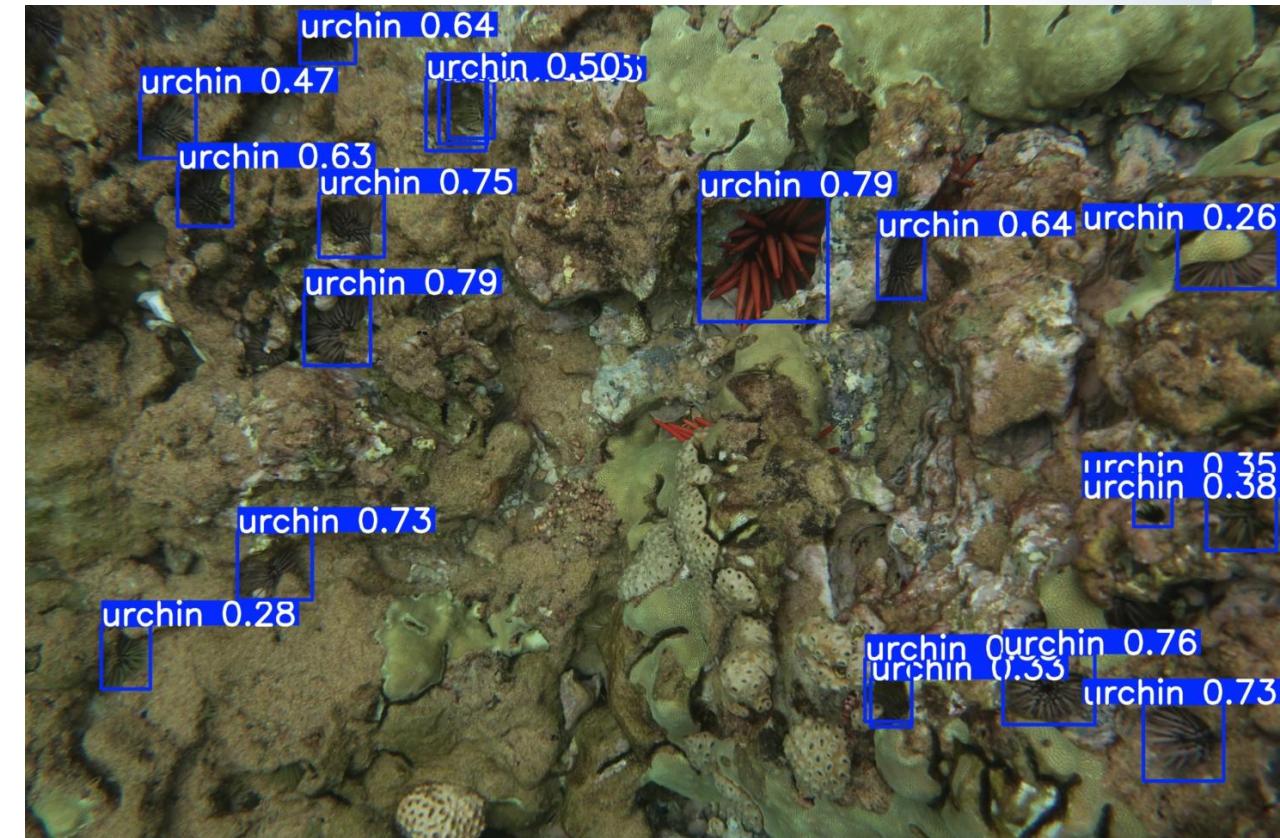


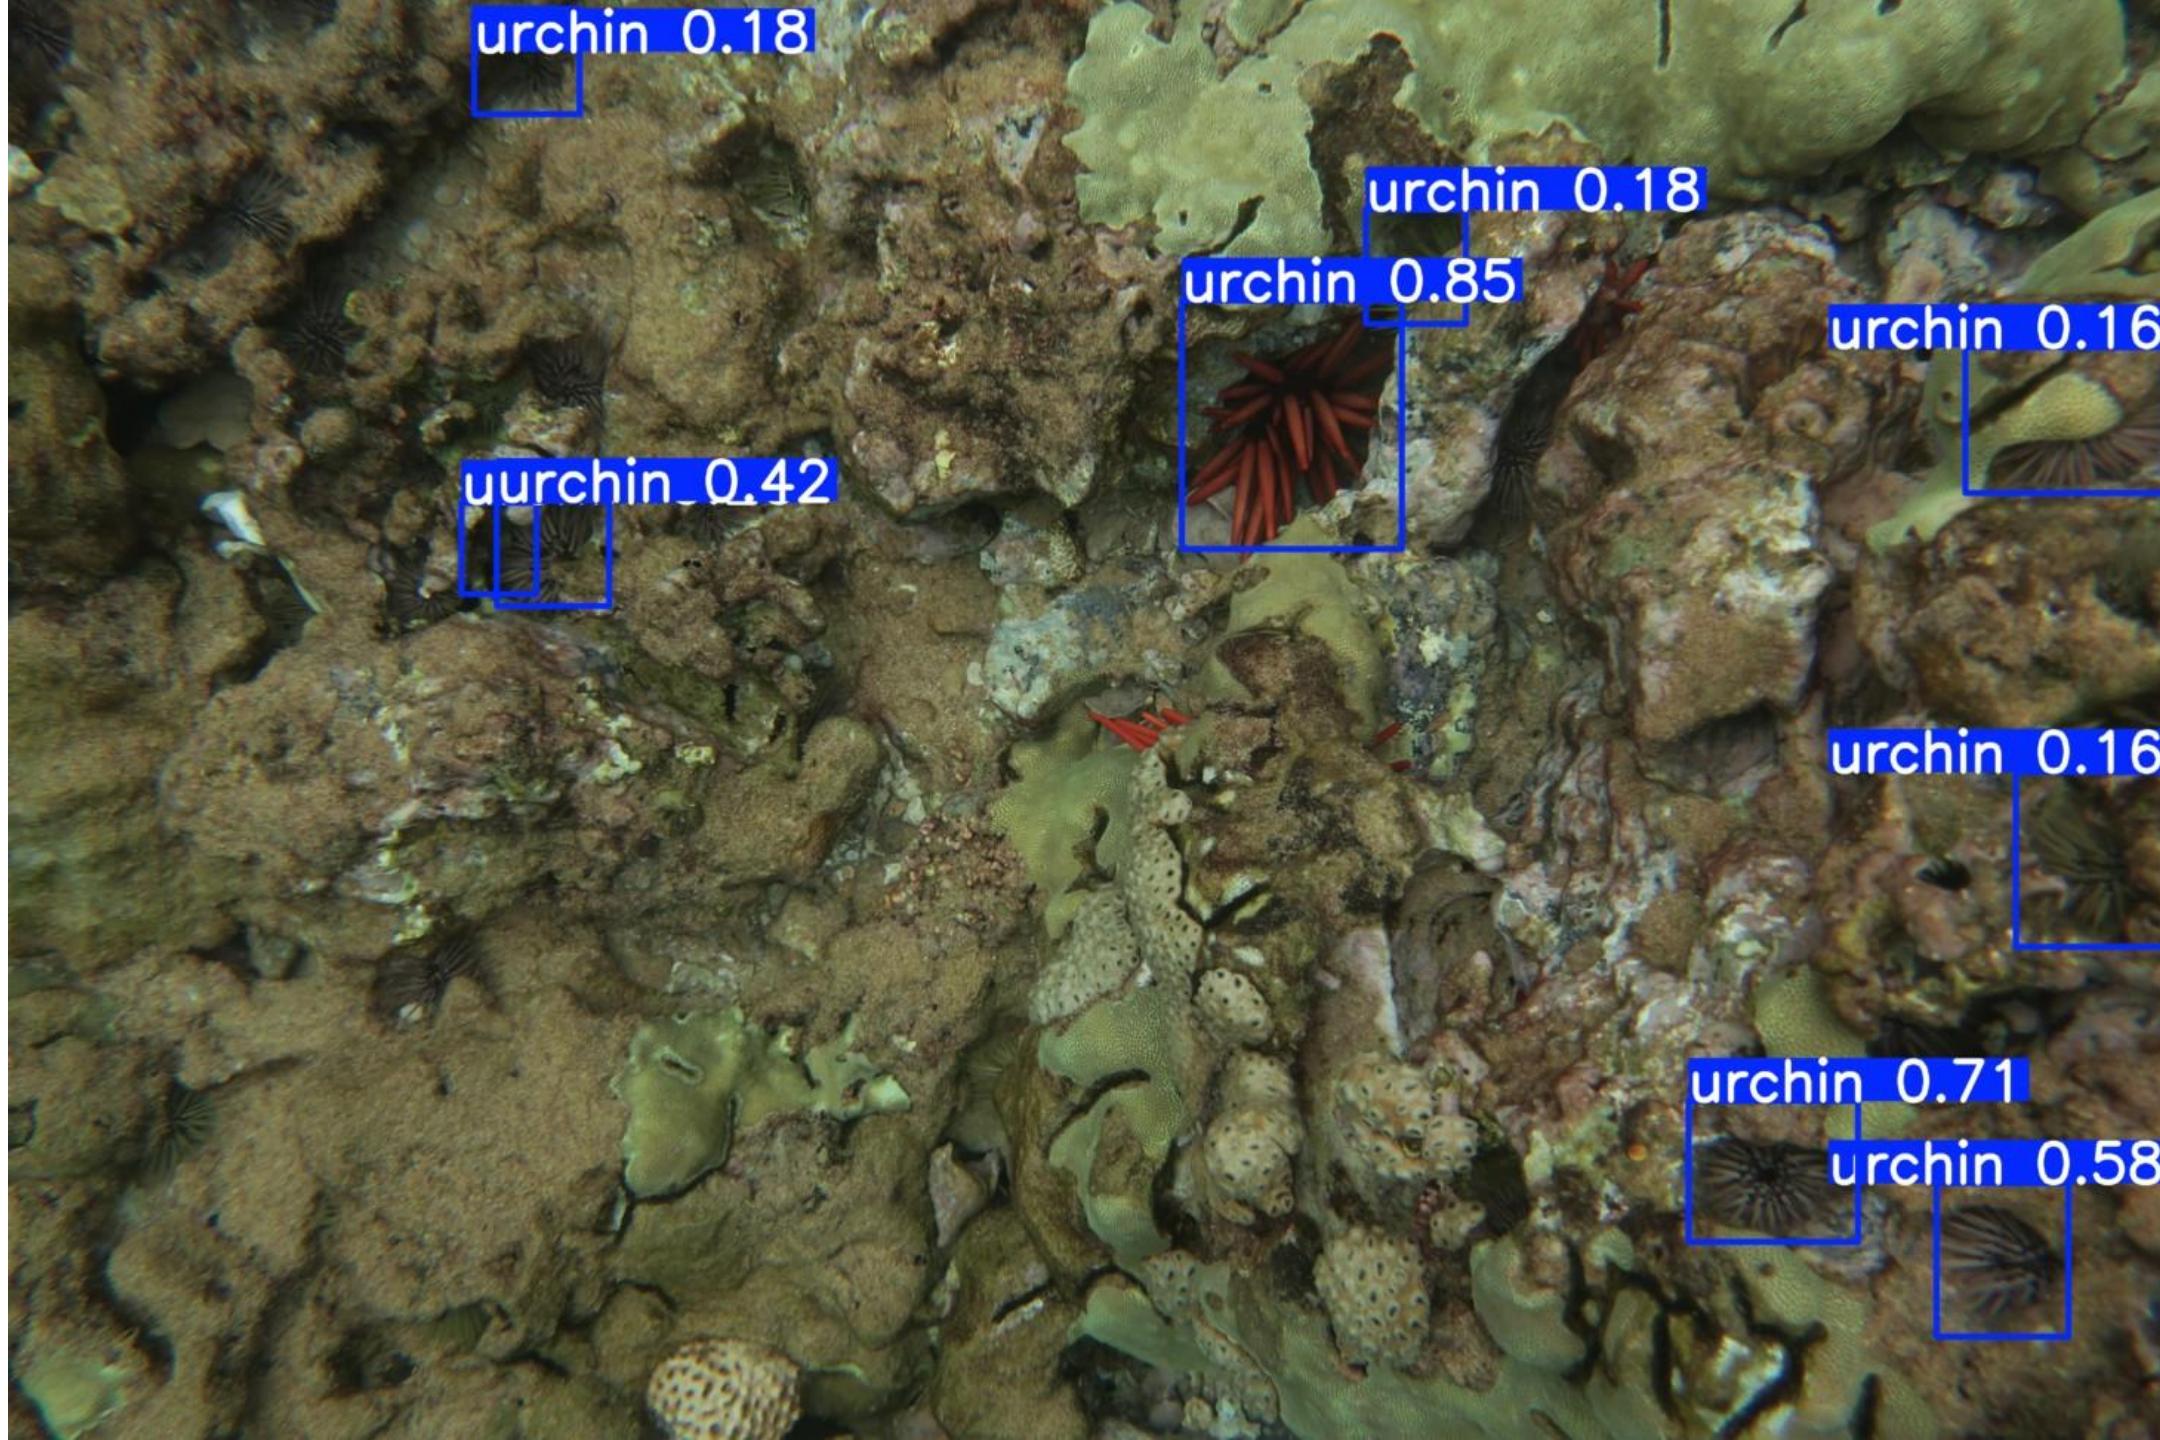
Small vs Larger Training Dataset Performance

YOLO11m | 3,330 image training dataset



YOLO11m | 10,345 image training dataset





Hybrid Object Detection: Using Patched Based Inference for small objects

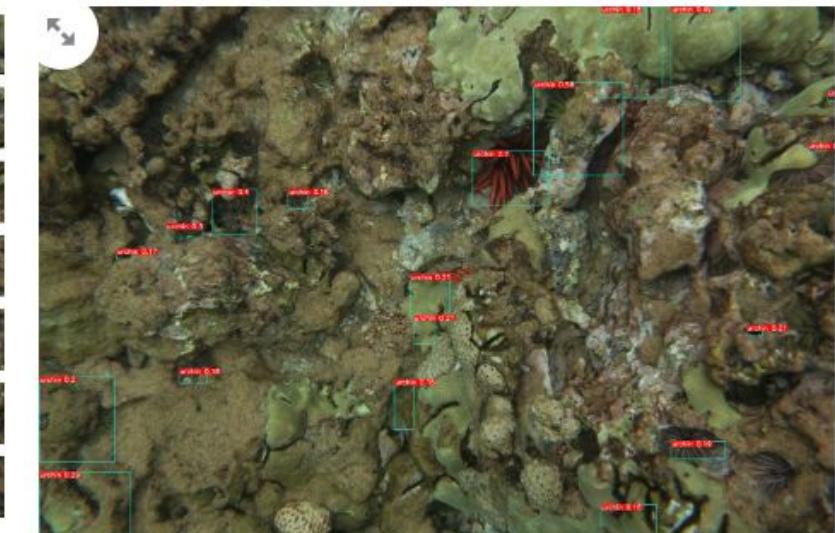
Processing MAI-B4010_2019_18.JPG...



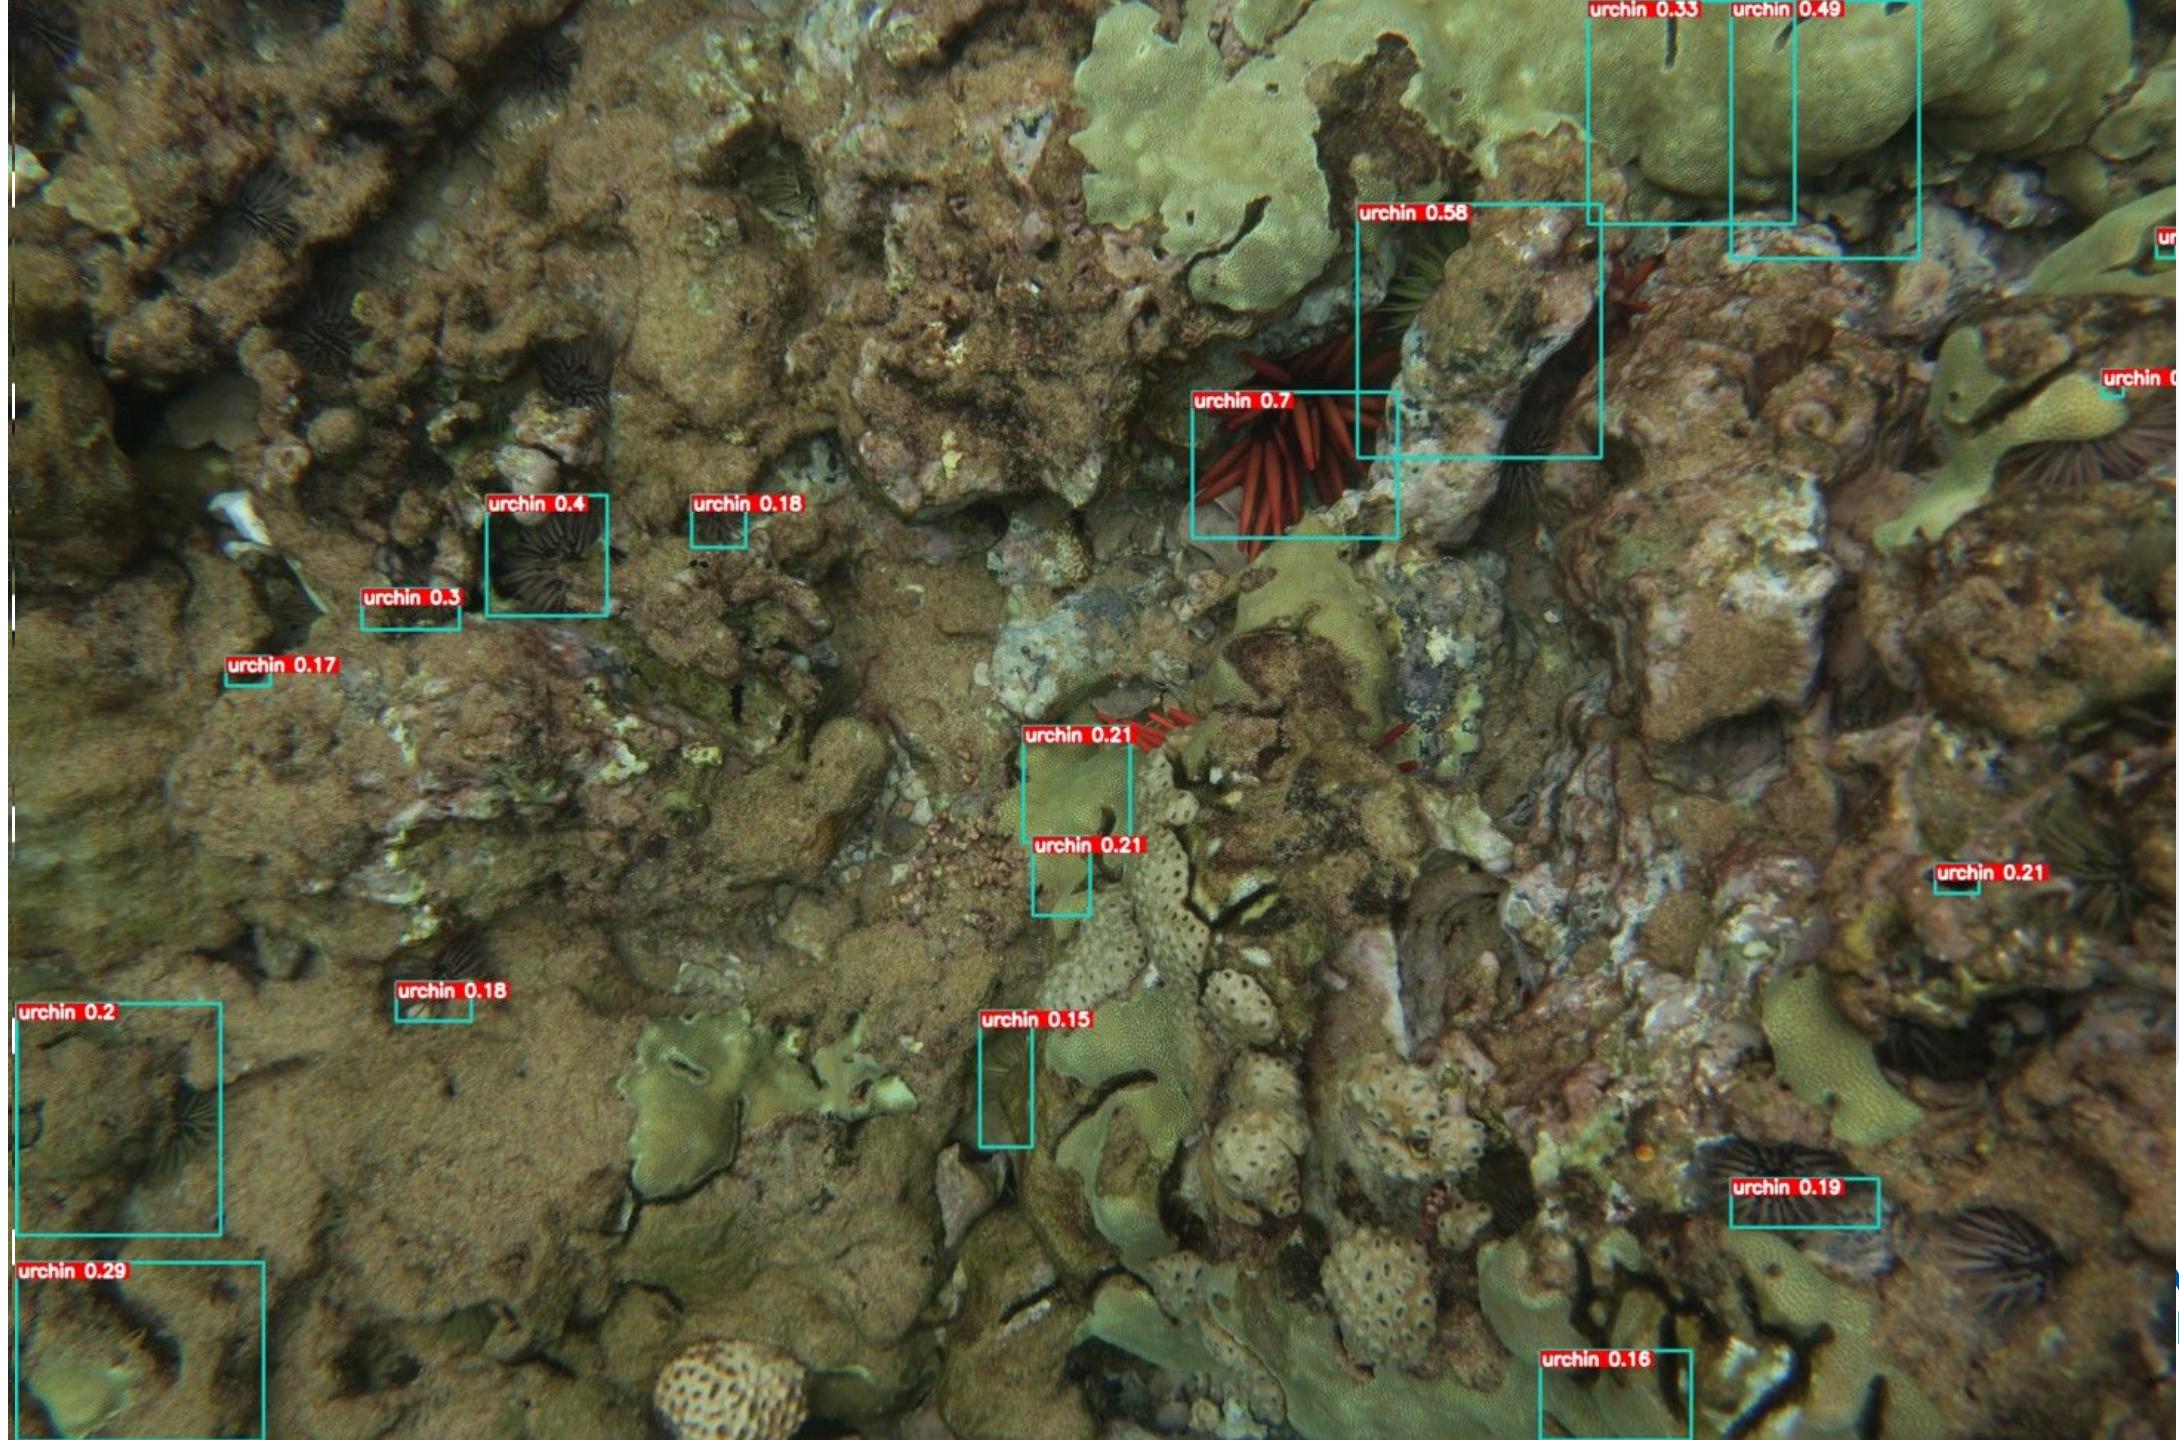
Original Image



Composite of Generated Crops

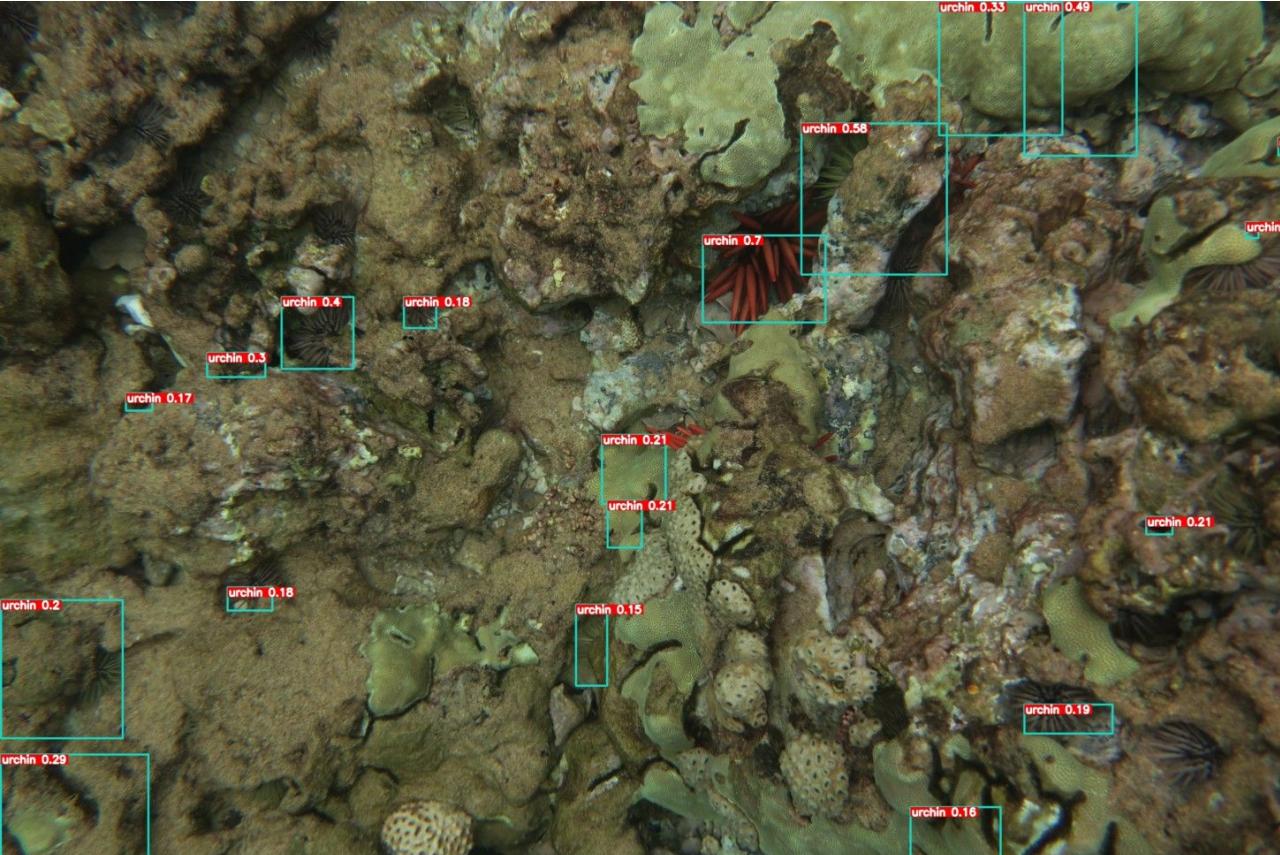


Detection Result

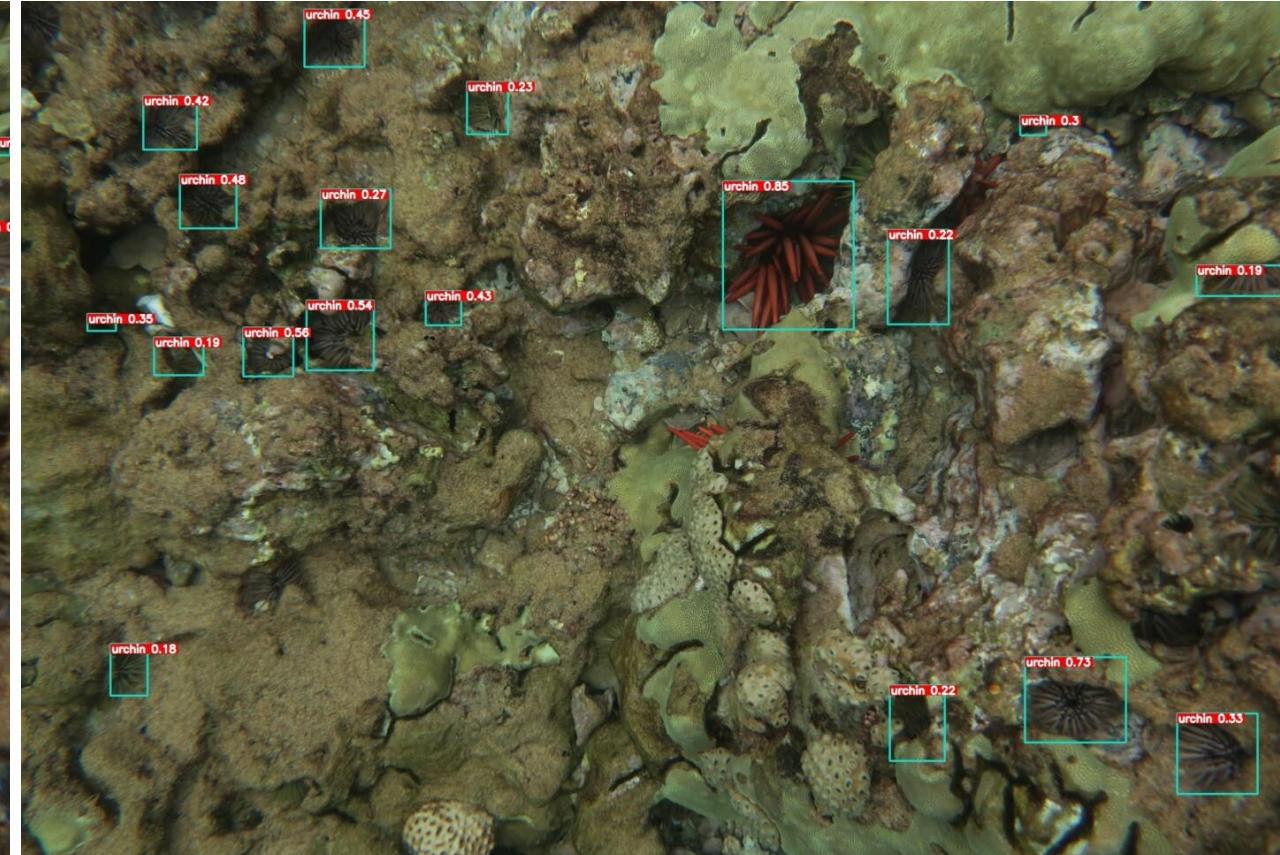


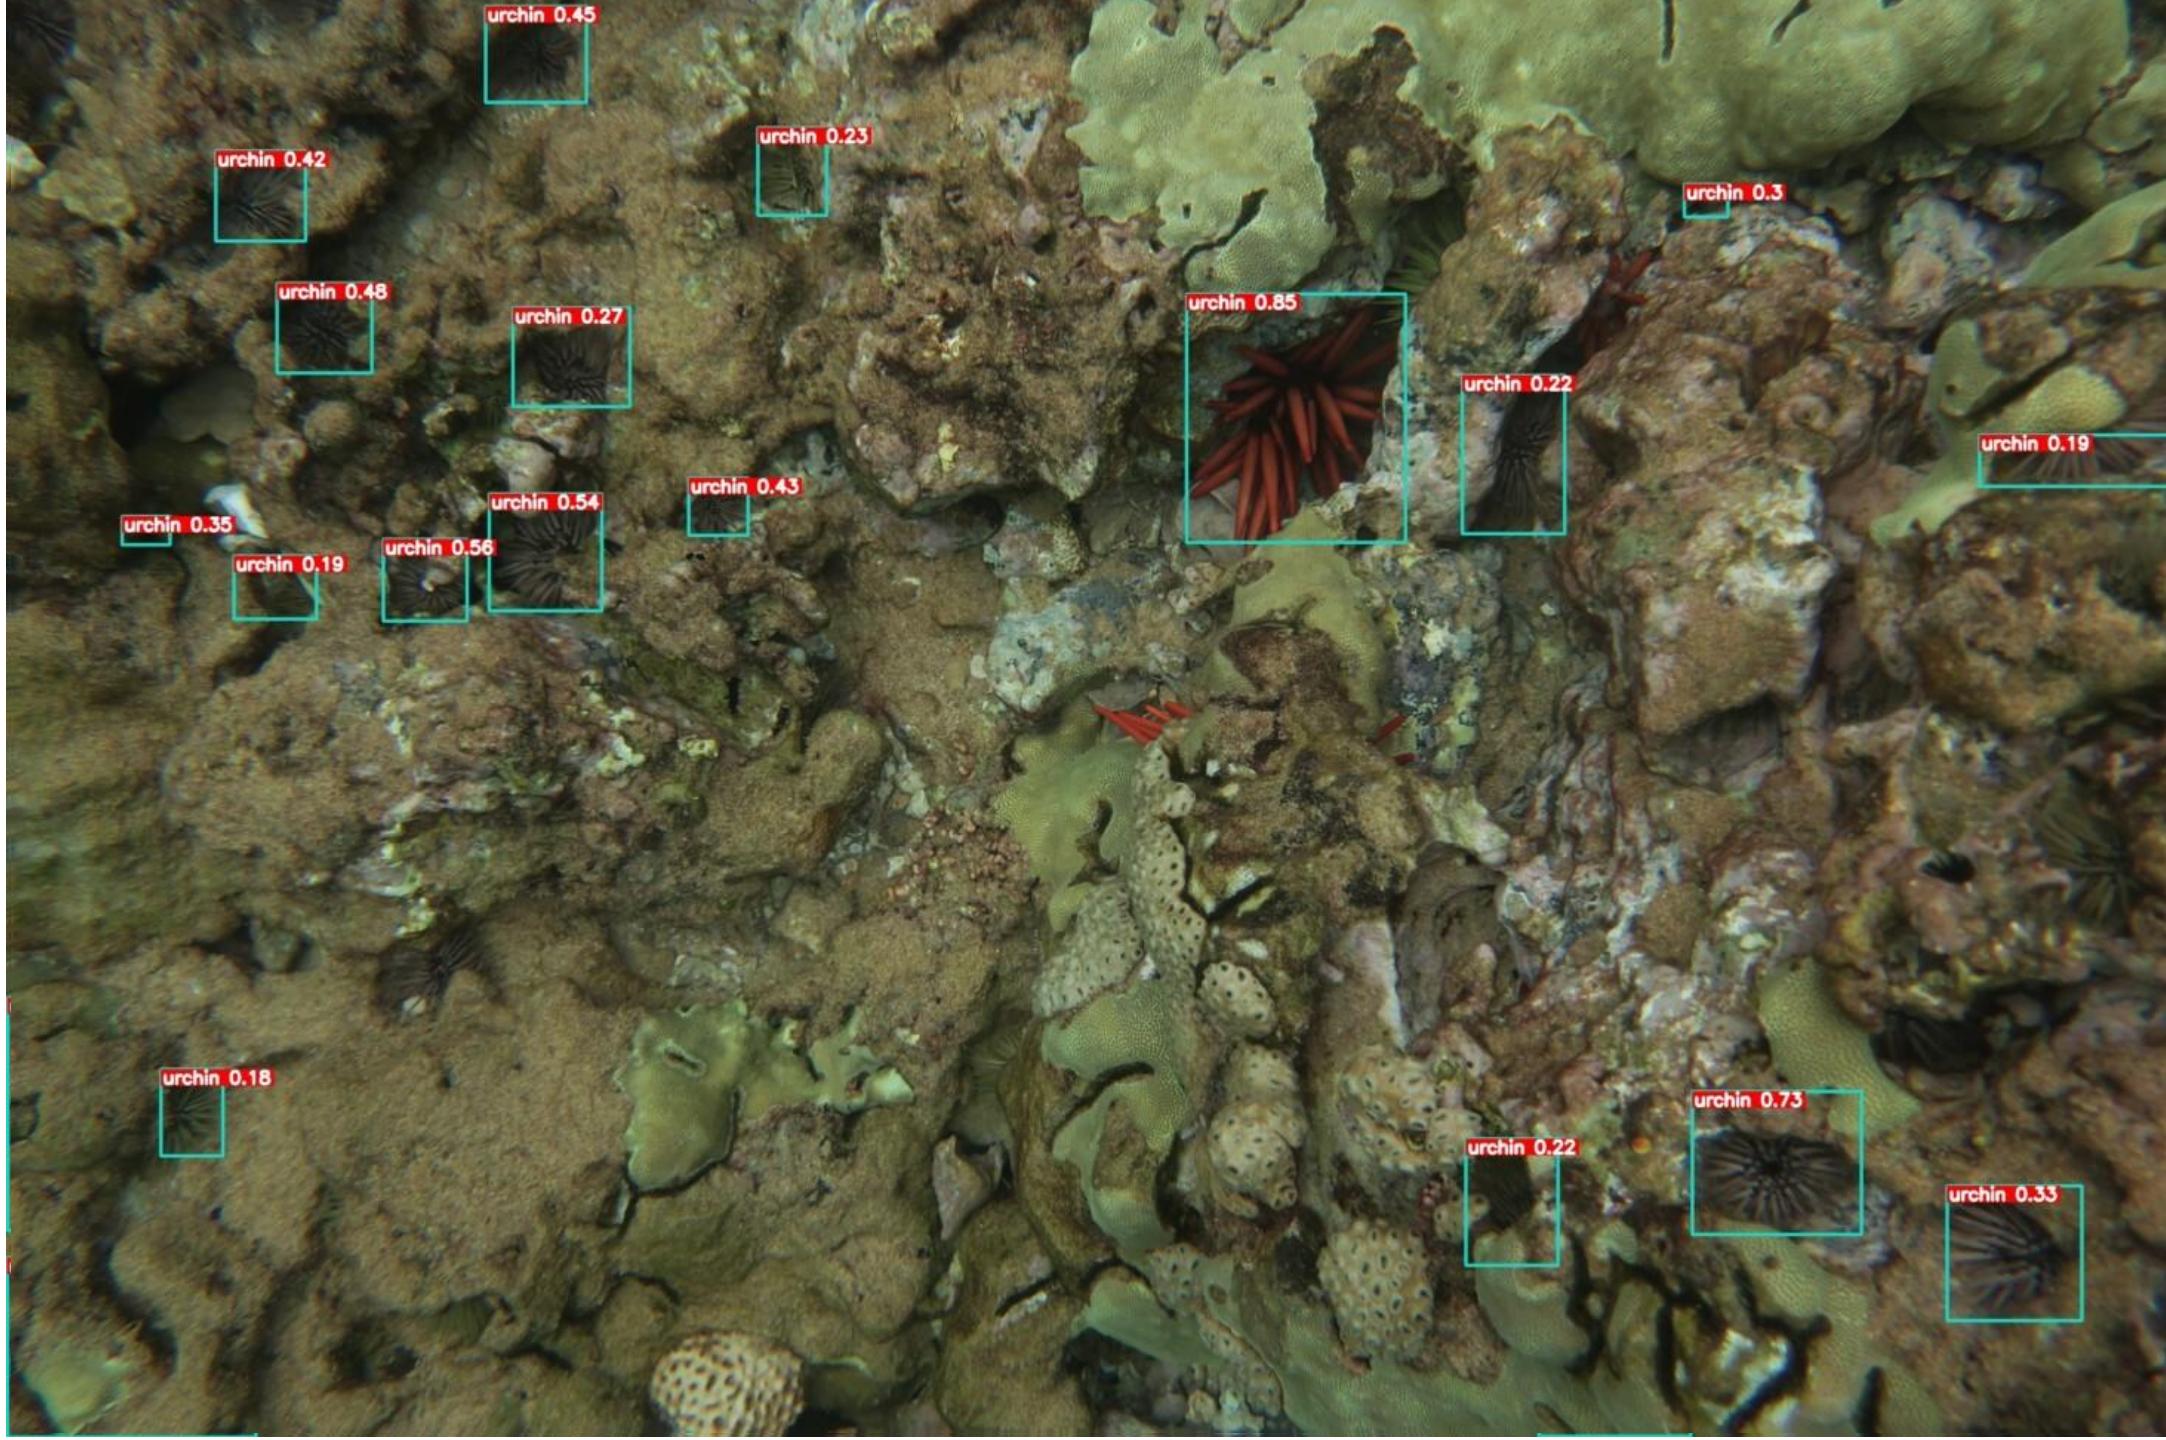
Patch Based vs Multi-Scale Patch Based

Patch Based Inference



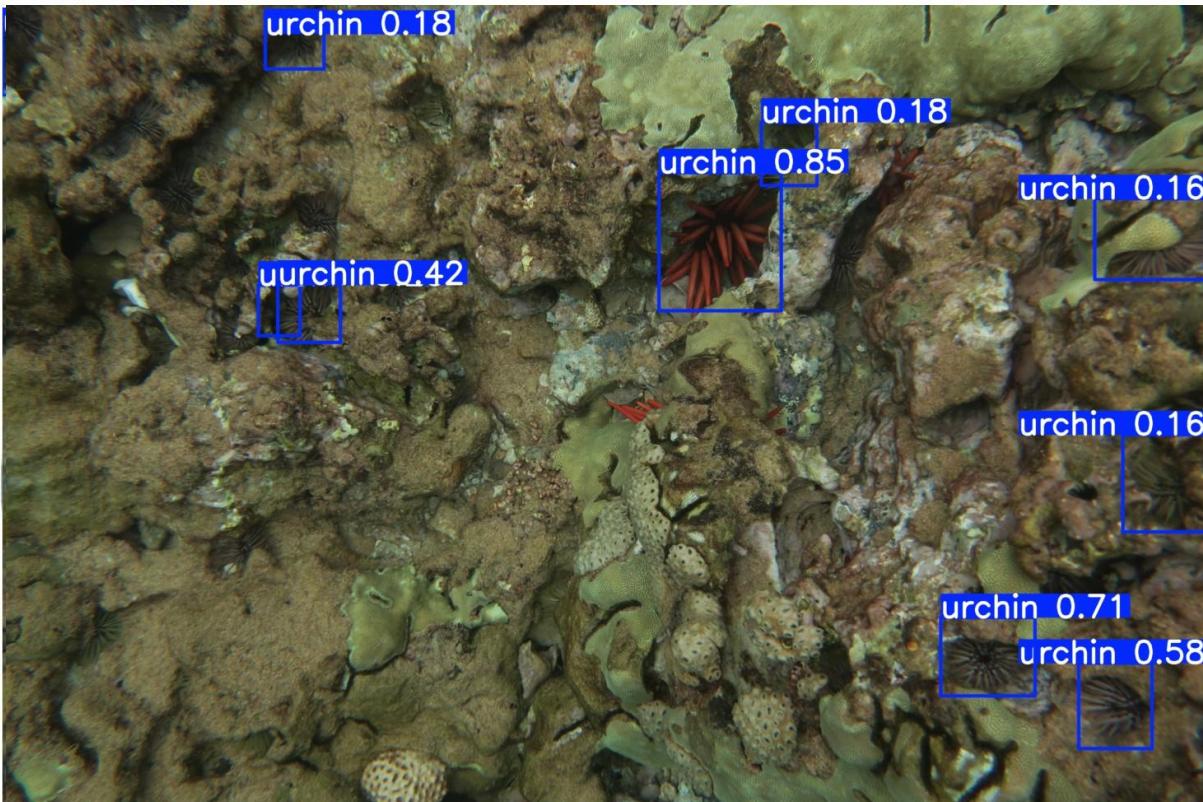
Multi-Scale Patch Based Inference



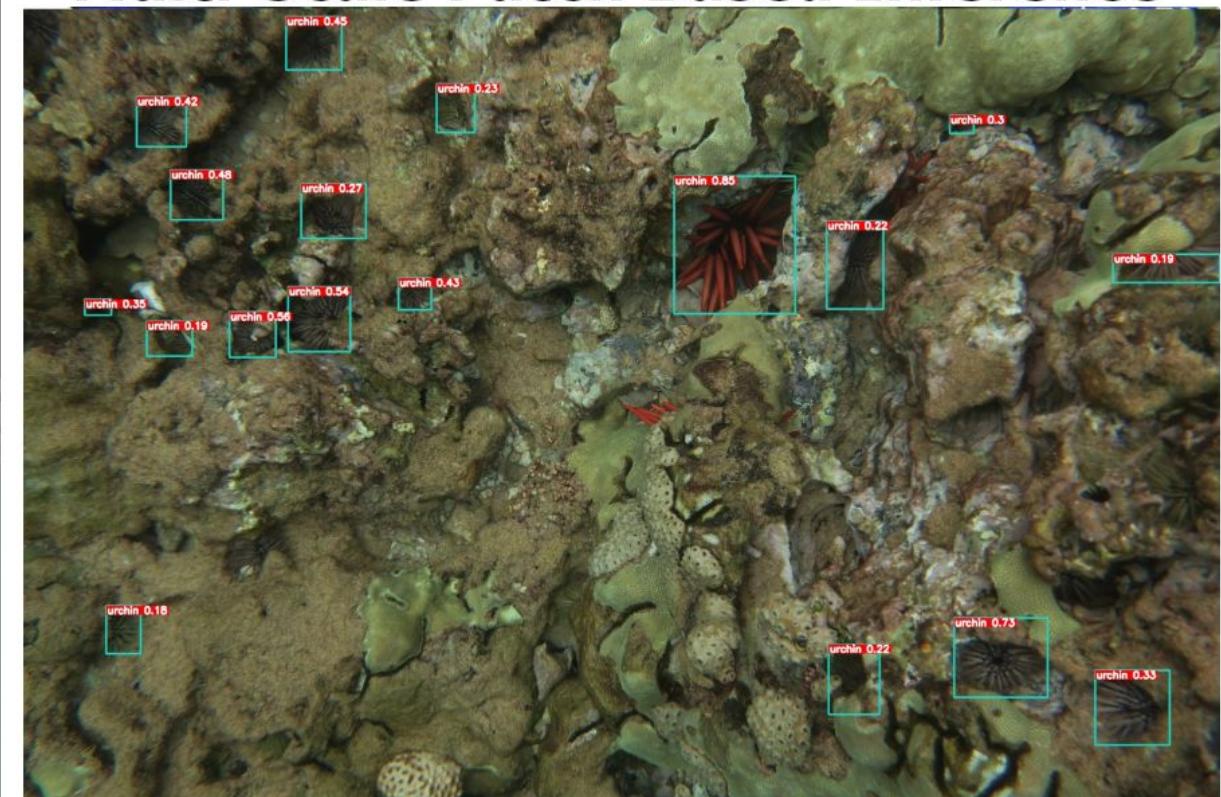


Multi-Scale Patch Based Inference Comparison

Urchin Detector yolo11m(medium)



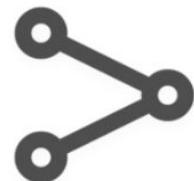
**Urchin Detector yolo11m(medium)
Multi-Scale Patch Based Inference**



Lessons Learned

- Explore a multi-model approach

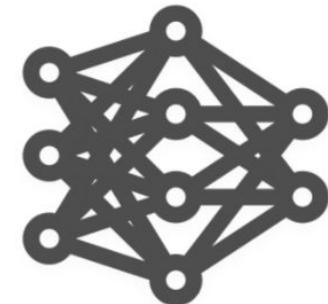
Use nano/small model for quick & real time data collection/detections



Small



Pass detections to larger model back in the office for finer detections/classifications/sizing



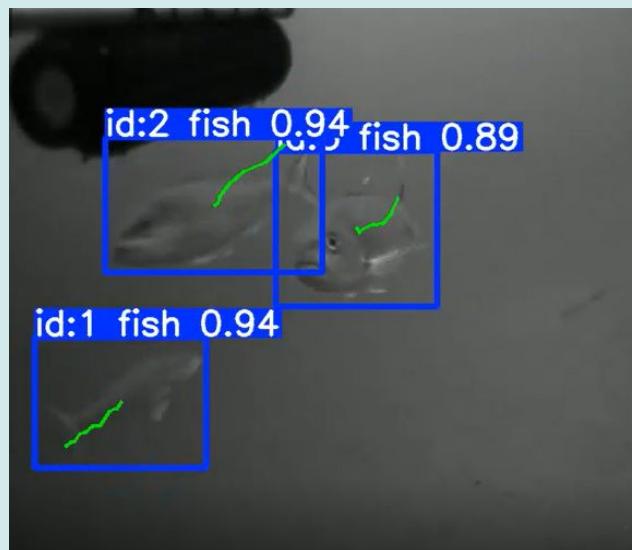
XLarge

Auto-Annotation:

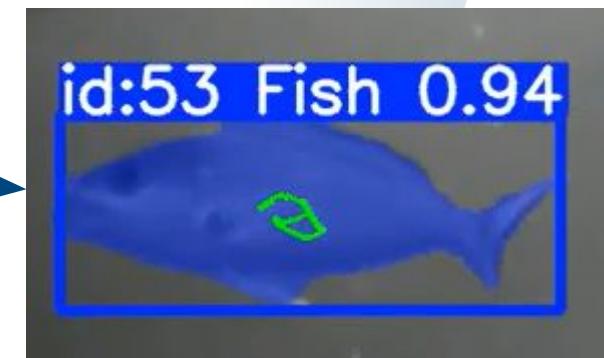
- A Quick Way to Segmentation Datasets & Models

Dataset

- + Trained Detection model
- + Segment Anything Model (SAM)
- = Segmentation dataset



Train on
segmentation
dataset for a
Segmentation Model



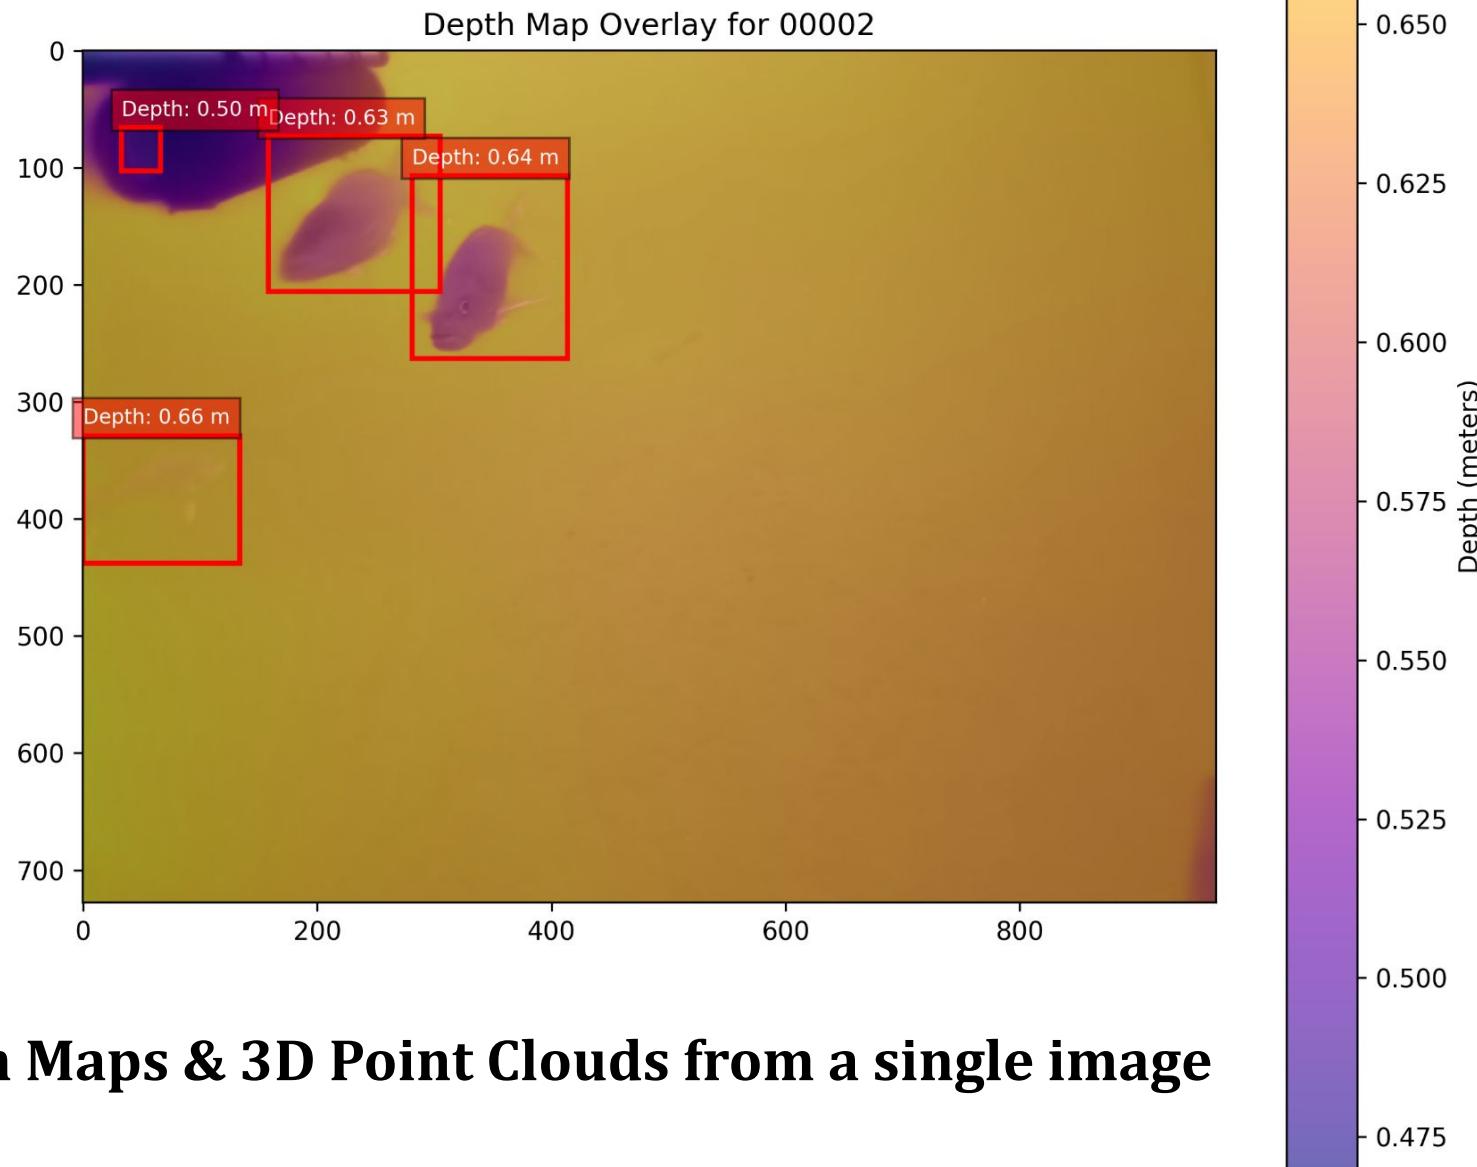
Object Detection Model



Segmentation Model



YOLO + Apple ML-Depth Pro



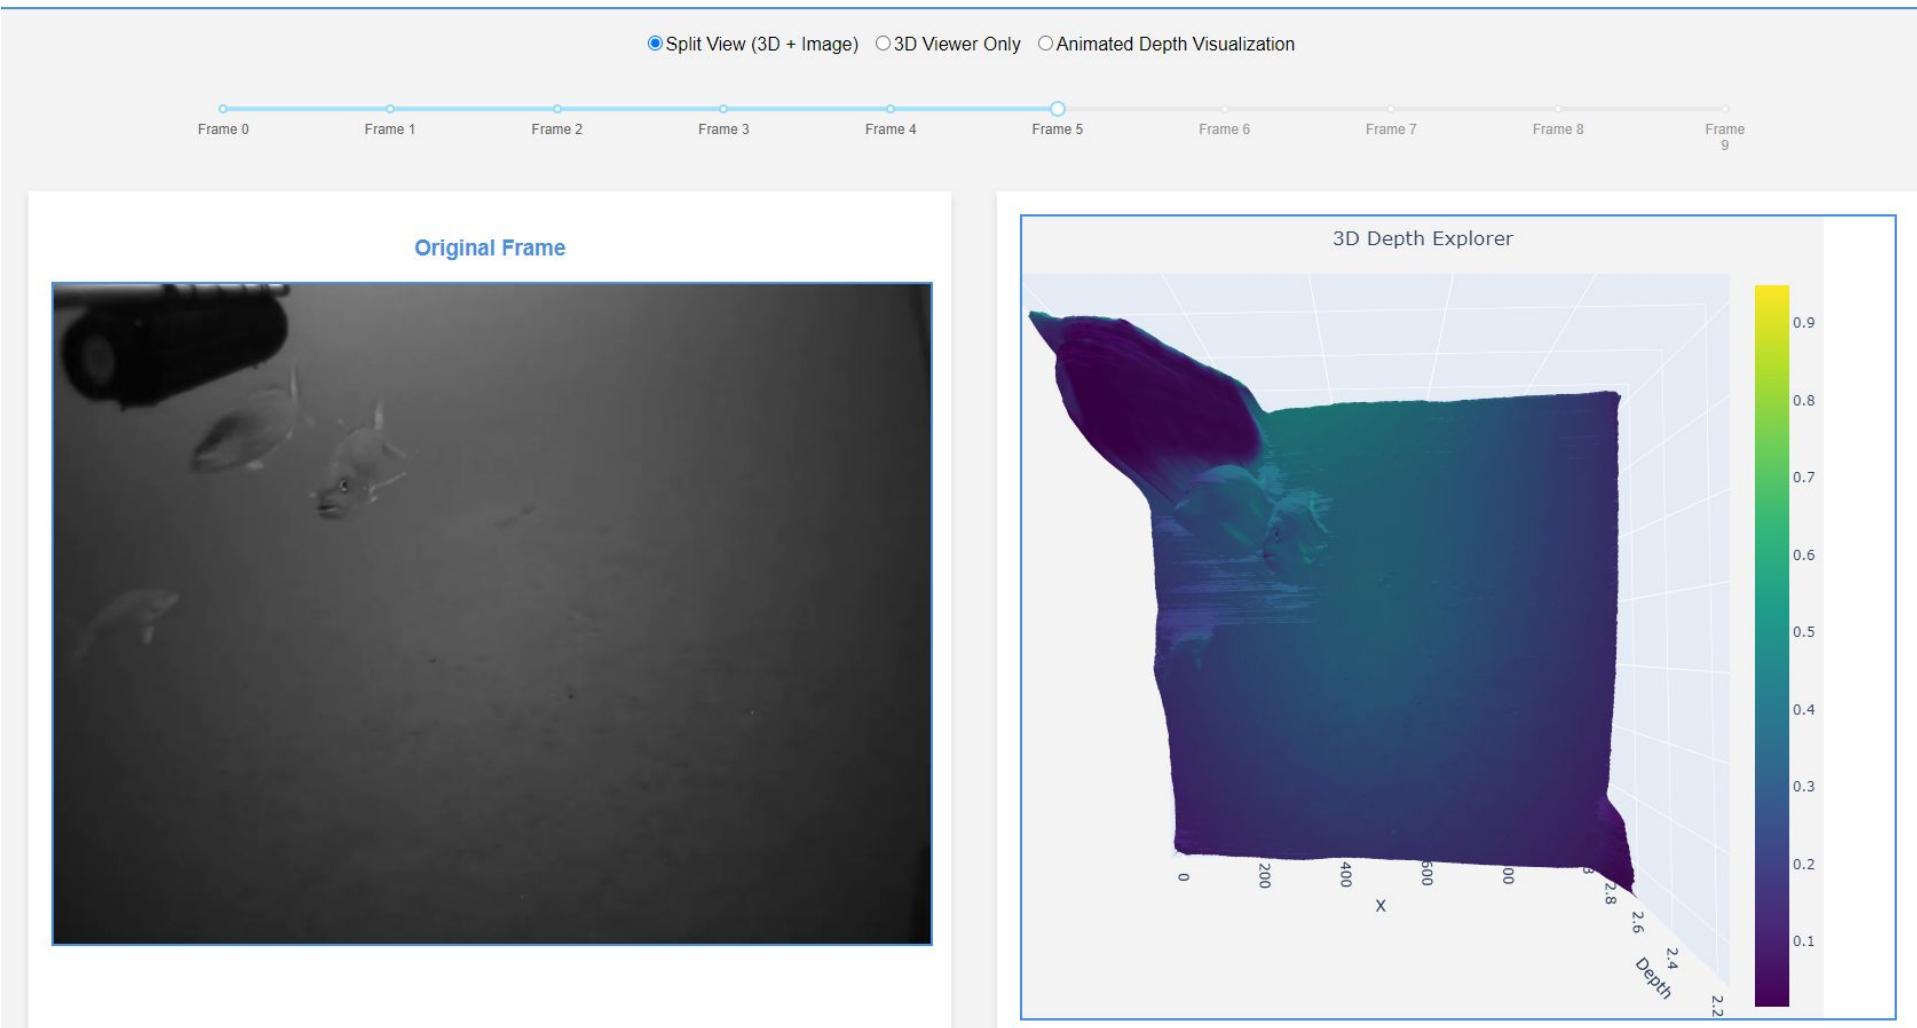
Depth Maps & 3D Point Clouds from a single image

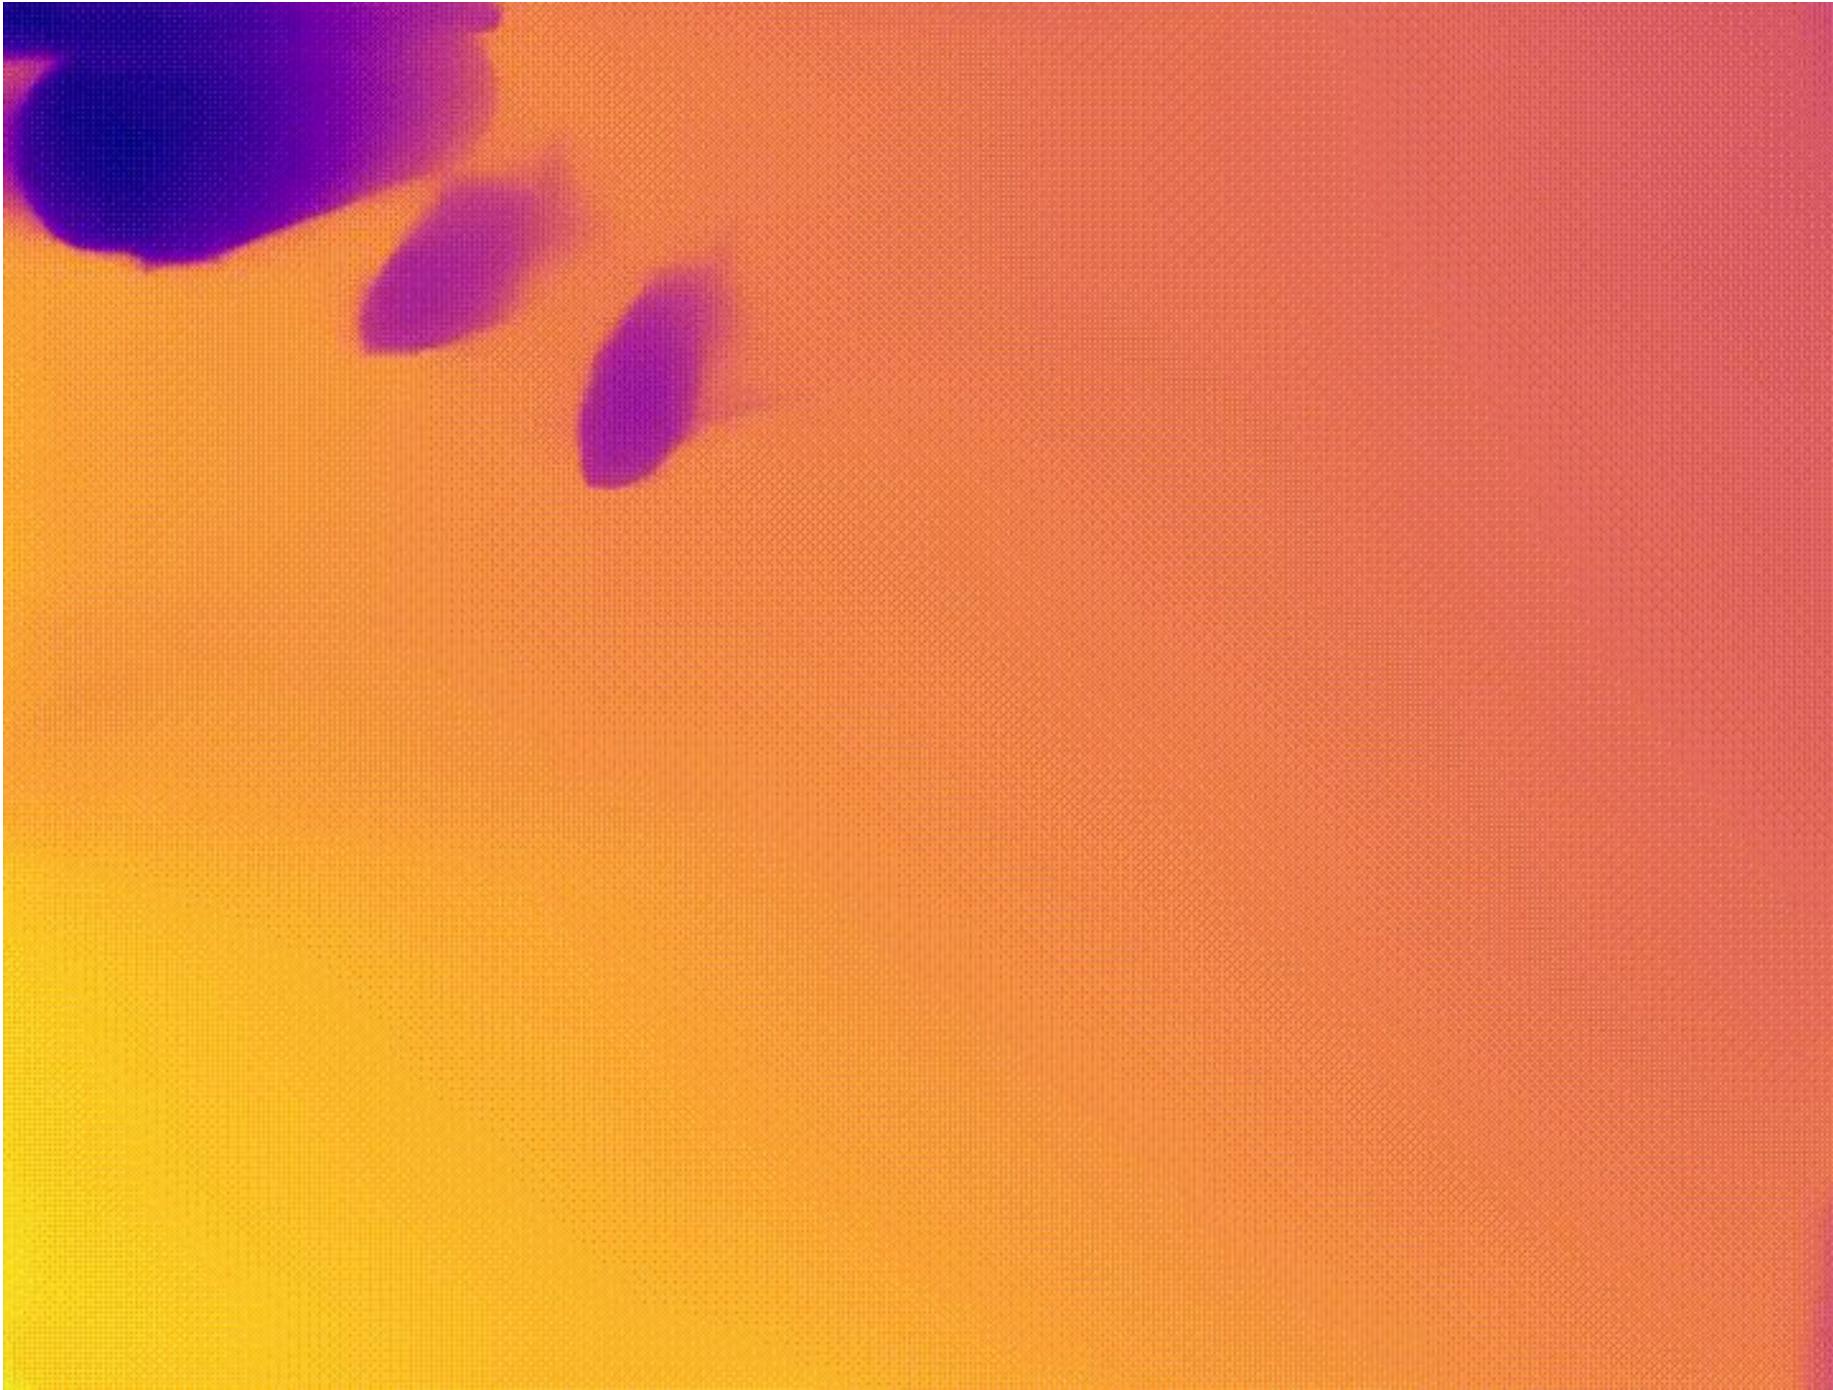
- https://connect.fisheries.noaa.gov/ml_depth_pro/

Depth Estimation using Machine Learning on Underwater Marine Imagery

Visualization of high-resolution depth maps generated from underwater marine imagery using Apple's [ML-Depth Pro](#) machine learning model. This is a test of Depth Pro's performance on underwater footage to evaluate its depth estimation in marine environments. No adjustments or fine tuning have been made at this time. Future exploration will include integrating known variables, such as camera focal length, and making adjustments for factors like water clarity.

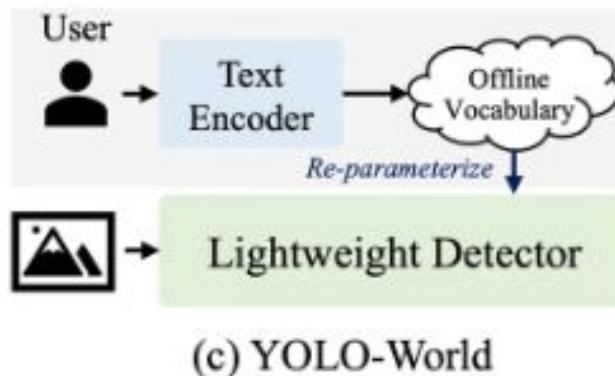
For more information, view this project's source: [GitHub Repository](#). Contact: michael.akridge@noaa.gov.





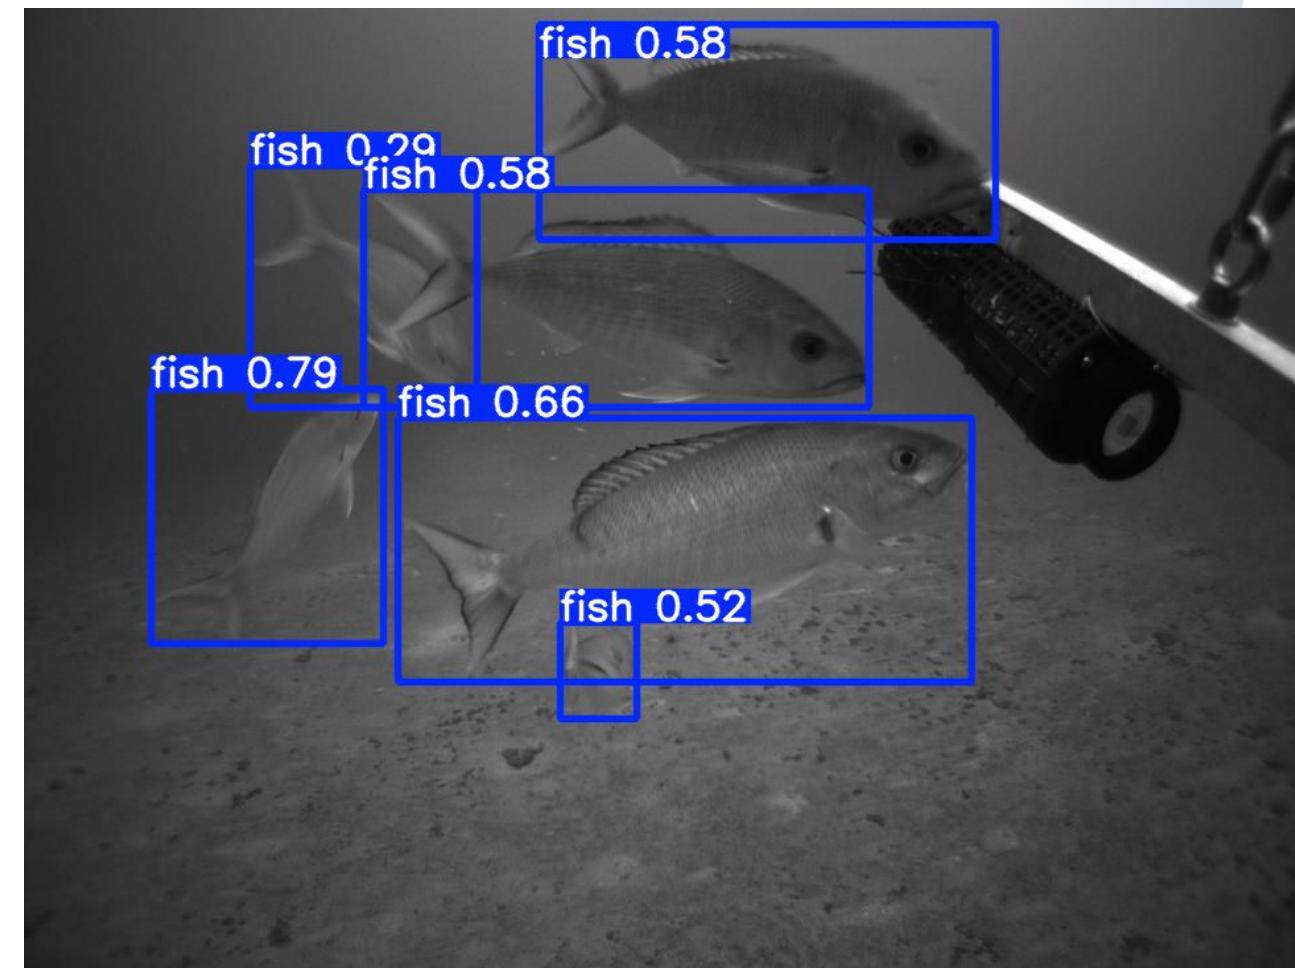
NOAA
FISHERIES

YOLO-World: "Prompt and Detect"



Zero-shot learning **enables a model to recognize object classes it has never been explicitly trained on.**

Speed: 15.2ms preprocess, 909.2ms inference, 21.4ms post process per image



Lessons Learned

- Easy to get a DOI



Connecting research and researchers

Michael Akridge

ID <https://orcid.org/0009-0007-0193-9705>

▼ Works (4)

AI Model: yolo11m-sea-urchin-detector

2024 | Research tool

DOI: [10.57967/HF/3342](https://doi.org/10.57967/HF/3342)

CONTRIBUTORS:

Source: Michael Akridge

NOAA-PIFSC Model Zoo

updated 16 days ago

Collection of pre-trained & custom models for object detection, image segmentation, classification and more.

akridge/yolo11-fish-detector-grayscale
Object Detection • Updated 15 days ago • ↓ 78 • ❤ 1

akridge/yolo11m-sea-urchin-detector
Object Detection • Updated 21 days ago • ↓ 54

akridge/yolo11-segment-fish-grayscale
Image Segmentation • Updated 15 days ago • ↓ 23 • ❤ 1



NOAA
FISHERIES

Community is Important: HuggingFace

Models

-  Urchin Model  224
-  Fish Models  272

Total: 496 Downloads



michael akridge

akridge

0 followers • 3 following

Edit profile

Settings

 [https://github.com/MichaelAkrige...](https://github.com/MichaelAkrige)

 MichaelAkridge-NOAA

AI-Ready Datasets

-  Object Detection  1,480
-  Segmentation  1,681

Total: 3,161 Downloads

Models 7

 akridge/yolo11x-sea-urchin-detector

Object Detection • Updated 13 days ago • ↓ 19

 akridge/yolo11-segment-fish-grayscale

Image Segmentation • Updated Oct 29, 2024 • ↓ 13 • ❤ 1

 akridge/yolo11n-sea-urchin-detector

Object Detection • Updated Oct 23, 2024 • ↓ 2

 akridge/Llama-2-7b-chat-hf-GGUF

Updated Sep 26, 2023 • ❤ 1

Datasets 2

↑ Sort: Recently updated

 akridge/MOUSS_fish_segment_dataset_grayscale

Viewer • Updated Oct 7, 2024 • ↓ 1.85k • ↓ 83

 akridge/MOUSS_fish_imagery_dataset_grayscale

Viewer • Updated Oct 7, 2024 • ↓ 1.85k • ↓ 31

Thanks!

Links:

- YOLO Library:
 - <https://github.com/ultralytics/ultralytics>
- Label Studio
 - <https://github.com/HumanSignal/label-studio>
- Toolkit:
 - <https://github.com/MichaelAkridge-NOAA/open-science-ai-toolkit>
- Models:
 - <https://huggingface.co/akridge>
- **Welcome to the Community Computer Vision Course**
- Cloud Data:
 - [List of NOAA Open Data](#)
 - PIFSC Cloud Console URL: [Google Cloud Storage Browser](#)

MAKER Lab

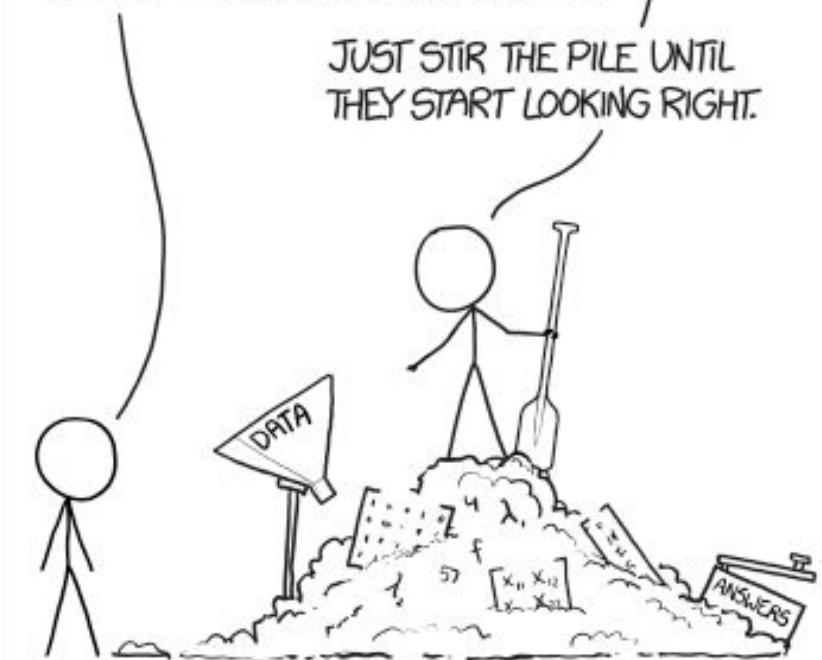


THIS IS YOUR MACHINE LEARNING SYSTEM?

YUP! YOU POUR THE DATA INTO THIS BIG PILE OF LINEAR ALGEBRA, THEN COLLECT THE ANSWERS ON THE OTHER SIDE.

WHAT IF THE ANSWERS ARE WRONG?

JUST STIR THE PILE UNTIL THEY START LOOKING RIGHT.



NOAA
FISHERIES

<https://xkcd.com/1838/>